




US 20220232739A1

(19) **United States**(12) **Patent Application Publication**
Heydari(10) **Pub. No.: US 2022/0232739 A1**(43) **Pub. Date: Jul. 21, 2022**(54) **INTELLIGENT COLD PLATE SYSTEM
WITH ACTIVE AND PASSIVE FEATURES
FOR A DATACENTER COOLING SYSTEM**(71) Applicant: **Nvidia Corporation**, Santa Clara, CA
(US)(72) Inventor: **Ali Heydari**, Albany, CA (US)(21) Appl. No.: **17/154,708**(22) Filed: **Jan. 21, 2021****Publication Classification**(51) **Int. Cl.**
H05K 7/20 (2006.01)(52) **U.S. Cl.**CPC **H05K 7/20509** (2013.01); **H05K 7/20318**
(2013.01); **H05K 7/20309** (2013.01); **H05K**
7/20836 (2013.01); **H05K 7/20381** (2013.01);
H05K 7/20827 (2013.01); **H05K 7/20336**
(2013.01)

(57)

ABSTRACT

Systems and methods for cooling a datacenter are disclosed. In at least one embodiment, a cold plate has microchannels and a heat pipe to support a first fluid in an active mode of operation of a cold plate that uses microchannels, and to support a second fluid in a passive mode of operation of a cold plate that uses a heat pipe.







FIGURE 1






FIGURE 2B






FIGURE 3






FIGURE 6

DATA CENTER
700




FIGURE 7




FIGURE 8




FIGURE 9




FIGURE 10A




FIGURE 10B

1000C →




FIGURE 10C




FIGURE 11




FIGURE 12




FIGURE 13




FIGURE 14




FIGURE 15




FIGURE 16




FIGURE 17




FIGURE 18A




FIGURE 18B




FIGURE 19




FIGURE 20

2100




FIGURE 21




FIGURE 22




FIGURE 23

2400




FIGURE 24

2500




FIGURE 25




FIGURE 26




FIGURE 27




FIGURE 28




FIGURE 29




FIGURE 30




FIGURE 31




FIGURE 32




FIGURE 33




FIGURE 34A




FIGURE 34B




FIGURE 35A



FIGURE 35B




FIGURE 36A




FIGURE 36B




FIGURE 36C




FIGURE 37




FIGURE 38




FIGURE 39




FIGURE 40




FIGURE 41




FIGURE 42

INTELLIGENT COLD PLATE SYSTEM WITH ACTIVE AND PASSIVE FEATURES FOR A DATACENTER COOLING SYSTEM

FIELD

[0001] At least one embodiment pertains to cooling systems, including systems and methods for operating those cooling systems. In at least one embodiment, such a cooling system can be utilized in a datacenter containing one or more racks or computing servers.

BACKGROUND

[0002] Datacenter cooling systems use fans to circulate air through server components.

[0003] Certain supercomputers or other high capacity computers may use water or other cooling systems instead of air-cooling systems to draw heat away from the server components or racks of the datacenter to an area external to the datacenter. The cooling systems may include a chiller within the datacenter area, which may include area external to the datacenter itself. Further, the area external to the datacenter may include a cooling tower or other external heat exchanger that receives heated coolant from the datacenter and that disperses the heat by forced air or other means to the environment (or an external cooling medium). The cooled coolant is recirculated back into the datacenter. The chiller and the cooling tower together form a chilling facility.

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 illustrates an exemplary datacenter cooling system subject to improvements described in at least one embodiment;

[0005] FIG. 2A illustrates server-level features associated with an intelligent cold plate system having active and passive features for a datacenter cooling system, according to at least one embodiment;

[0006] FIG. 2B illustrates cold plate-level features associated with an intelligent cold plate system having active and passive features for a datacenter cooling system, according to at least one embodiment;

[0007] FIG. 3 illustrates rack-level features associated with an intelligent cold plate system having active and passive features for a datacenter cooling system, according to at least one embodiment;

[0008] FIG. 4 illustrates datacenter-level features associated with an intelligent cold plate system having active and passive features for a datacenter cooling system, according to at least one embodiment;

[0009] FIG. 5 illustrates a method associated with a datacenter cooling system of FIG. 2A—4, according to at least one embodiment;

[0010] FIG. 6 illustrates a distributed system, in accordance with at least one embodiment;

[0011] FIG. 7 illustrates an exemplary datacenter, in accordance with at least one embodiment;

[0012] FIG. 8 illustrates a client-server network, in accordance with at least one embodiment;

[0013] FIG. 9 illustrates a computer network, in accordance with at least one embodiment;

[0014] FIG. 10A illustrates a networked computer system, in accordance with at least one embodiment;

[0015] FIG. 10B illustrates a networked computer system, in accordance with at least one embodiment;

[0016] FIG. 10C illustrates a networked computer system, in accordance with at least one embodiment;

[0017] FIG. 11 illustrates one or more components of a system environment in which services may be offered as third-party network services, in accordance with at least one embodiment;

[0018] FIG. 12 illustrates a cloud computing environment, in accordance with at least one embodiment;

[0019] FIG. 13 illustrates a set of functional abstraction layers provided by a cloud computing environment, in accordance with at least one embodiment;

[0020] FIG. 14 illustrates a supercomputer at a chip level, in accordance with at least one embodiment;

[0021] FIG. 15 illustrates a supercomputer at a rack module level, in accordance with at least one embodiment;

[0022] FIG. 16 illustrates a supercomputer at a rack level, in accordance with at least one embodiment;

[0023] FIG. 17 illustrates a supercomputer at a whole system level, in accordance with at least one embodiment;

[0024] FIG. 18A illustrates inference and/or training logic, in accordance with at least one embodiment;

[0025] FIG. 18B illustrates inference and/or training logic, in accordance with at least one embodiment;

[0026] FIG. 19 illustrates training and deployment of a neural network, in accordance with at least one embodiment;

[0027] FIG. 20 illustrates an architecture of a system of a network, in accordance with at least one embodiment;

[0028] FIG. 21 illustrates an architecture of a system of a network, in accordance with at least one embodiment;

[0029] FIG. 22 illustrates a control plane protocol stack, in accordance with at least one embodiment;

[0030] FIG. 23 illustrates a user plane protocol stack, in accordance with at least one embodiment;

[0031] FIG. 24 illustrates components of a core network, in accordance with at least one embodiment;

[0032] FIG. 25 illustrates components of a system to support network function virtualization (NFV), in accordance with at least one embodiment;

[0033] FIG. 26 illustrates a processing system, in accordance with at least one embodiment;

[0034] FIG. 27 illustrates a computer system, in accordance with at least one embodiment;

[0035] FIG. 28 illustrates a system, in accordance with at least one embodiment;

[0036] FIG. 29 illustrates an exemplary integrated circuit, in accordance with at least one embodiment;

[0037] FIG. 30 illustrates a computing system, according to at least one embodiment;

[0038] FIG. 31 illustrates an APU, in accordance with at least one embodiment;

[0039] FIG. 32 illustrates a CPU, in accordance with at least one embodiment;

[0040] FIG. 33 illustrates an exemplary accelerator integration slice, in accordance with at least one embodiment;

[0041] FIGS. 34A-34B illustrate exemplary graphics processors, in accordance with at least one embodiment;

[0042] FIG. 35A illustrates a graphics core, in accordance with at least one embodiment;

[0043] FIG. 35B illustrates a GPGPU, in accordance with at least one embodiment;

[0044] FIG. 36A illustrates a parallel processor, in accordance with at least one embodiment;

[0045] FIG. 36B illustrates a processing cluster, in accordance with at least one embodiment;

[0046] FIG. 36C illustrates a graphics multiprocessor, in accordance with at least one embodiment;

[0047] FIG. 37 illustrates a software stack of a programming platform, in accordance with at least one embodiment;

[0048] FIG. 38 illustrates a CUDA implementation of a software stack of FIG. 37, in accordance with at least one embodiment;

[0049] FIG. 39 illustrates a ROCm implementation of a software stack of FIG. 37, in accordance with at least one embodiment;

[0050] FIG. 40 illustrates an OpenCL implementation of a software stack of FIG. 37, in accordance with at least one embodiment;

[0051] FIG. 41 illustrates software that is supported by a programming platform, in accordance with at least one embodiment; and

[0052] FIG. 42 illustrates compiling code to execute on programming platforms of FIGS. 37-40, in accordance with at least one embodiment.

DETAILED DESCRIPTION

[0053] In at least one embodiment, an exemplary datacenter 100 can be utilized as illustrated in FIG. 1, which has a cooling system subject to improvements described herein. In at least one embodiment, numerous specific details are set forth to provide a thorough understanding, but concepts herein may be practiced without one or more of these specific details. In at least one embodiment, datacenter cooling systems can respond to sudden high heat requirements caused by changing computing-loads in present day computing components. In at least one embodiment, as these requirements are subject to change or tend to range from a minimum to a maximum of different cooling requirements, these requirements must be met in an economical manner, using an appropriate cooling system. In at least one embodiment, for moderate to high cooling requirements, liquid cooling system may be used. In at least one embodiment, high cooling requirements are economically satisfied by localized immersion cooling. In at least one embodiment, these different cooling requirements also reflect different heat features of a datacenter. In at least one embodiment, heat generated from these components, servers, and racks are cumulatively referred to as a heat feature or a cooling requirement as cooling requirement must address a heat feature entirely.

[0054] In at least one embodiment, a datacenter liquid cooling system is disclosed. In at least one embodiment, this datacenter cooling system addresses heat features in associated computing or datacenter devices, such as in graphics processing units (GPUs), in switches, in dual inline memory module (DIMMs), or central processing units (CPUs). In at least one embodiment, these components may be referred to herein as high heat density computing components. Furthermore, in at least one embodiment, an associated computing or datacenter device may be a processing card having one or more GPUs, switches, or CPUs thereon. In at least one embodiment, each of GPUs, switches, and CPUs may be a heat generating feature of a computing device. In at least one embodiment, a GPU, a CPU, or a switch may have one or more cores, and each core may be a heat generating feature.

[0055] In at least one embodiment, a cold plate system is provided that has both active and passive features for a

datacenter cooling system. In at least one embodiment, a cold plate enables liquid cooling for datacenter cooling systems. In at least one embodiment, a cold plate may have microchannels and an associated flow controller that actively drives fluid, such as coolant, through microchannels of a cold plate. In at least one embodiment, a cold plate may be associated with a computing device and fluid through provided microchannels enable extracting heat from an associated computing device. In at least one embodiment, a heat pipe may be part of a cold plate system. In at least one embodiment, a heat pipe may be embedded in a cavity of a cold plate. In at least one embodiment, a cold plate system may be referred to as a cold plate even if a passive feature, such as a heat pipe, extends out of a cold plate.

[0056] In at least one embodiment, a heat pipe may be in a separate section of a cold plate that a section provided for microchannels of a cold plate. In at least one embodiment, a heat pipe extends out of a cold plate to a location remote from a cold plate to dissipate heat from liquid therein to an environment away from at least one computing device. In at least one embodiment, a section of a heat pipe located externally from a cold plate is a condenser section. In at least one embodiment, a condenser section of a heat pipe may be supported by a fan providing forced air cooling. In at least one embodiment, an intelligent control system selects to engage, activate, or deactivate an active feature of a cold plate (having microchannels) to provide additional cooling to a passive feature (having a heat pipe). In at least one embodiment, an intelligent control system may therefore combine active and passive features (representing active and passive modes) for addressing cooling requirements and for efficient power utilization in a datacenter.

[0057] In at least one embodiment, an intelligent cold plate system with active and passive features is able to address issues of a cold plate that may be related to a thermal margin for managing high heat density servers in datacenters. In at least one embodiment, coolant may be a forced fluid provided to flow within microchannels for cooling imparted to at least one computing device by removing heat from a high heat density component. In at least one embodiment, an intelligent cold plate system with active and passive features can balance power consumption requirements associated with multiple component cooling as well as with singular cooling requirements to promote efficiency of a cold plate.

[0058] In at least one embodiment, an intelligent cold plate system with active and passive features provides different thermal margins to address different cooling requirements. In at least one embodiment, an intelligent cold plate system with active and passive features can be used with its passive feature alone or with an active feature combined with the passive feature, depending on temperature sensed from, as well as power consumed by, a computing component. In at least one embodiment, power consumed may reflect a future high heat generation from at least one computing component and, subsequently, a future cooling requirement for at least one computing component.

[0059] In at least one embodiment, an exemplary datacenter 100 can be utilized as illustrated in FIG. 1, which has a cooling system subject to improvements described herein. In at least one embodiment, a datacenter 100 may be one or more rooms 102 having racks 110 and auxiliary equipment to house one or more servers on one or more server trays. In at least one embodiment, a datacenter 100 is supported by a

cooling tower **104** located external to a datacenter **100**. In at least one embodiment, a cooling tower **104** dissipates heat from within a datacenter **100** by acting on a primary cooling loop **106**. In at least one embodiment, a cooling distribution unit (CDU) **112** is used between a primary cooling loop **106** and a second or secondary cooling loop **108** to enable extraction of heat from a second or secondary cooling loop **108** to a primary cooling loop **106**. In at least one embodiment, a secondary cooling loop **108** can access various plumbing into a server tray as required, in an aspect. In at least one embodiment, loops **106**, **108** are illustrated as line drawings, but a person of ordinary skill would recognize that one or more plumbing features may be used. In at least one embodiment, flexible polyvinyl chloride (PVC) pipes may be used along with associated plumbing to move fluid along in each provided loop **106**; **108**. In at least one embodiment, one or more coolant pumps may be used to maintain pressure differences within coolant loops **106**, **108** to enable movement of coolant according to temperature sensors in various locations, including in a room, in one or more racks **110**, and/or in server boxes or server trays within one or more racks **110**.

[0060] In at least one embodiment, coolant in a primary cooling loop **106** and in a secondary cooling loop **108** may be at least water and an additive. In at least one embodiment, an additive may be glycol or propylene glycol. In operation, in at least one embodiment, each of a primary and a secondary cooling loops may have their own coolant. In at least one embodiment, coolant in secondary cooling loops may be proprietary to requirements of components in a server tray or in associated racks **110**. In at least one embodiment, a CDU **112** is capable of sophisticated control of coolants, independently or concurrently, within provided coolant loops **106**, **108**. In at least one embodiment, a CDU may be adapted to control flow rate of coolant so that coolant is appropriately distributed to extract heat generated within associated racks **110**. In at least one embodiment, more flexible tubing **114** is provided from a secondary cooling loop **108** to enter each server tray to provide coolant to electrical and/or computing components therein.

[0061] In at least one embodiment, tubing **118** that forms part of a secondary cooling loop **108** may be referred to as room manifolds. Separately, in at least one embodiment, further tubing **116** may extend from row manifold tubing **118** and may also be part of a secondary cooling loop **108** but may be referred to as row manifolds. In at least one embodiment, coolant tubing **114** enters racks as part of a secondary cooling loop **108** but may be referred to as rack cooling manifold within one or more racks. In at least one embodiment, row manifolds **116** extend to all racks along a row in a datacenter **100**. In at least one embodiment, plumbing of a secondary cooling loop **108**, including coolant manifolds **118**, **116**, and **114** may be improved by at least one embodiment herein. In at least one embodiment, a chiller **120** may be provided in a primary cooling loop within datacenter **102** to support cooling before a cooling tower. In at least one embodiment, additional cooling loops that may exist in a primary control loop and that provide cooling external to a rack and external to a secondary cooling loop, may be taken together with a primary cooling loop and is distinct from a secondary cooling loop, for this disclosure.

[0062] In at least one embodiment, in operation, heat generated within server trays of provided racks **110** may be transferred to a coolant exiting one or more racks **110** via

flexible tubing of a row manifold **114** of a second cooling loop **108**. In at least one embodiment, second coolant (in a secondary cooling loop **108**) from a CDU **112**, for cooling provided racks **110**, moves towards one or more racks **110** via provided tubing. In at least one embodiment, second coolant from a CDU **112** passes from on one side of a room manifold having tubing **118**, to one side of a rack **110** via a row manifold **116**, and through one side of a server tray via different tubing **114**. In at least one embodiment, spent or returned second coolant (or exiting second coolant carrying heat from computing components) exits out of another side of a server tray (such as enter left side of a rack and exits right side of a rack for a server tray after looping through a server tray or through components on a server tray). In at least one embodiment, spent second coolant that exits a server tray or a rack **110** comes out of different side (such as exiting side) of tubing **114** and moves to a parallel, but also exiting side of a row manifold **116**. In at least one embodiment, from a row manifold **116**, spent second coolant moves in a parallel portion of a room manifold **118** and is going in an opposite direction than incoming second coolant (which may also be renewed second coolant), and towards a CDU **112**.

[0063] In at least one embodiment, spent second coolant exchanges its heat with a primary coolant in a primary cooling loop **106** via a CDU **112**. In at least one embodiment, spent second coolant may be renewed (such as relatively cooled when compared to a temperature at a spent second coolant stage) and ready to be cycled back to through a second cooling loop **108** to one or more computing components. In at least one embodiment, various flow and temperature control features in a CDU **112** enable control of heat exchanged from spent second coolant or flow of second coolant in and out of a CDU **112**. In at least one embodiment, a CDU **112** may be also able to control a flow of primary coolant in primary cooling loop **106**.

[0064] In at least one embodiment, server-level features **200** as illustrated in FIG. 2A can be associated with an intelligent cold plate system having active and passive features for a datacenter cooling system. In at least one embodiment, server-level features **200** include a server tray or box **202**. In at least one embodiment, a server tray or box **202** includes a server manifold **204** to be intermediately coupled between provided cold plates **210A-D** of a server tray or box **202** and rack manifolds of a rack hosting a server tray or box **202**. In at least one embodiment, a server tray or box **202** includes one or more cold plates **210A-D** associated with one or more computing or datacenter components or devices **220A-D**. In at least one embodiment, one or more server-level cooling loops **214A, B** may be provided between a server manifold **204** and one or more colds plates **210A-D**.

[0065] In at least one embodiment, each server-level cooling loop **214A; B** includes an inlet line **210** and an outlet line **212**. In at least one embodiment, when there are series configured cold plates **210A, B**, an intermediate line **216** may be provided. In at least one embodiment, one or more cold plates **210A-D** may support distinct ports and channels for a secondary coolant of a secondary cooling loop or a different local coolant, such as a fluid circulated from other sources associated with a local cooling loop. In at least one embodiment, fluid for cooling may be provided to a server manifold **204** via provided inlet and outlets **206A, 206B**.

[0066] In at least one embodiment, a server tray 202 is an immersive-cooled server tray that may be flooded by fluid. In at least one embodiment, a fluid for an immersive-cooled server tray may be a dielectric engineered fluid capable of being used in an immersive-cooled server. In at least one embodiment, a secondary coolant or local coolant may be used to cool engineered fluid. In at least one embodiment, a local coolant may be used to cool engineered fluid when a primary cooling loop associated with a secondary cooling loop circulating a secondary coolant has failed or is failing. In at least one embodiment, at least one cold plate therefore has ports for a secondary cooling loop and for a local cooling loop and can support a local cooling loop that is activated in an event of a failure in a primary cooling loop. In at least one embodiment, an intelligent cold plate system with active and passive features may be used without a secondary cooling loop by at least passive features alone and an associated local cooling loop.

[0067] In at least one embodiment, at least one dual-cooling cold plate 210B; 250 may be configured to work alongside regular cold plates 210A, C, D. In at least one embodiment, both a dual-cooling cold plate and a regular cold plate may have active and passive features. In at least one embodiment, a three-dimensional (3D) blow-up illustration (cold plate 250) provides internal detail of at least some features that may be included in a dual-cooling cold plate 210B of an intelligent cold plate system with active and passive features. In at least one embodiment, a blow-up illustrated in FIG. 2A illustrates active features within section or cavity 282A of a cold plate 250. In at least one embodiment, a dual-cooling or a regular cold plate may have sections or cavities 282A, 282B for microchannels and a heat pipe. In at least one embodiment, sections or cavities 282A, 282B of a cold plate system maintain separation between a first fluid (for microchannels) and a second fluid (for a heat pipe).

[0068] In at least one embodiment, a regular cold plate may have one set of microchannels 264; 270 instead of two sets illustrated. In at least one embodiment, a dual-cooling cold plate 250 has distinct paths 264, 270 (each path also referred to as microchannels) for secondary coolant of a secondary cooling loop and for local coolant of a local cooling loop. In at least one embodiment, secondary or local coolant may not be dielectric in property. In at least one embodiment, in a use case of an immersive-cooled server, fluid that may be a dielectric engineered fluid may be adapted for both, a cold plate application and an immersive-cooled server tray application.

[0069] In at least one embodiment, reference to cold plate, along with its dual-cooling features, implies a reference to a cold plate that can support at least two types of cooling loops, unless otherwise stated. In at least one embodiment, both types of colds plates receive fluid for cooling from a same secondary cooling loop and can both support a local cooling loop. In at least one embodiment, a standard coolant, such as facility water may be used in both a secondary cooling loop and a local cooling loop. In at least one embodiment, secondary coolant already within a cold plate is diverted to a local cooling loop. In at least one embodiment, a heat pipe in a section 282B of a cold plate provides distinct cooling than a secondary coolant or a local cooling.

[0070] In at least one embodiment, local coolant may be same or similar to a secondary coolant to avoid issues regarding chemistry differences and manufacturer require-

ments of cold plates used in a datacenter cooling system. In at least one embodiment, a fluid may only support cold plate usage and may not be available for immersive cooling. In at least one embodiment, each type of cold plate receives different fluid from respective secondary or other cooling loops interfacing with a primary cooling loop. In at least one embodiment, in situations where different fluids are used with different coolant distribution units (CDUs) of different secondary loops, then a local cooling loop may be suited for a dual-cooling cold plate so that different channels may be used for each of a local coolant and different secondary coolants.

[0071] In at least one embodiment, a dual-cooling cold plate 250 is adapted to receive a two types of fluids (such as a secondary coolant and a local coolant) and to keep two types of fluids distinct from each other via their distinct ports 252, 272; 268, 262 and their distinct paths 264, 270. In at least one embodiment, each distinct path is a fluid path. In at least one embodiment, fluid (such as local coolant) from a fluid source and a secondary coolant may be of a same or similar composition and may be restocked from a same source in a datacenter cooling system.

[0072] In at least one embodiment, a dual-cooling cold plate 250 includes ports 252, 272 to receive fluid into a cold plate 250 and to pass fluid out of a cold plate 250. In at least one embodiment, a dual-cooling cold plate 250 includes ports 268, 262 to receive a secondary coolant into a cold plate 250 and to pass a secondary coolant out of a cold plate 250. In at least one embodiment, ports 252, 272 may have valve covers 254, 260 that may be directional, and pressure controlled. In at least one embodiment, valve covers may be associated with all provided ports. In at least one embodiment, provided valve covers 254, 260 are mechanical features of associated flow controllers that also have corresponding electronic features (such as at least one processor to execute instructions stored in associated memory and to control mechanical features for associated flow controllers).

[0073] In at least one embodiment, each valve may be actuated by an electronic feature of an associated flow controller. In at least one embodiment, electronic and mechanical features of provided flow controllers are integrated. In at least one embodiment, electronic and mechanical features of provided flow controllers are physically distinct. In at least one embodiment, reference to flow controllers may be to one or more of provided electronic and mechanical features or to their union but is at least in reference to features enabling control of flow of coolant or fluid through each cold plate or an immersion-cooled server tray or box.

[0074] In at least one embodiment, electronic features of provided flow controllers receive control signals and assert control over mechanical features. In at least one embodiment, electronic features of provided flow controllers may be actuators or other electronic parts of other similar electromechanical features. In at least one embodiment, flow pumps may be used as flow controllers. In at least one embodiment, impellers, pistons, or bellows may be mechanical features, and an electronic motor and circuitry form electronic features of provided flow controllers.

[0075] In at least one embodiment, circuitry of provided flow controllers may include processors, memories, switches, sensors, and other components, altogether forming electronic features of provided flow controllers. In at least one embodiment, provided ports 252, 262, 272, 268 of

provided flow controllers are adapted to either allow entry or to allow egress of an immersive fluid. In at least one embodiment, flow controllers **280** may be associated with fluid lines **276** (also **256**, **274**) that enable entry and egress of fluid (such as a local coolant) to a cold plate **210B**. In at least one embodiment, other flow controllers may be similarly associated with coolant lines **210**, **216**, **212** (also **266**, **258**) to enable entry and egress of a secondary coolant to a cold plate **210B**.

[0076] In at least one embodiment, fluid (such as a local coolant) enters provided fluid lines **276** via dedicated fluid inlet and outlet lines **208A**, **B**. In at least one embodiment, a server manifold **204** is adapted with channels therein (illustrated by dotted lines) to support distinct paths to distinct fluid lines **276** (also **256**, **274**) and to any remaining loops **214A**, **B** that are associated with secondary coolant inlet and outlet lines **206A**, **B**. In at least one embodiment, there may be multiple manifolds to support fluid (a local coolant) and secondary coolant distinctly. In at least one embodiment, there may be multiple manifolds to support entry and egress, distinctly, for each of a fluid and a secondary coolant. In at least one embodiment, if a fluid is same or similar as a secondary coolant, then at least two different flows via a same fluid path (at least within a cold plate or a server tray) to a fluid source and to a secondary coolant row manifold (such as row manifold **350** in FIG. 3) are enabled.

[0077] In at least one embodiment, a first flow may be to enable fluid (such as local coolant) to flow through one or more provided ports **252**, **272** and an associated path **270**. In at least one embodiment, a dual-cooling cold plate **250** may have isolated plate sections that are flooded with a fluid and/or a secondary coolant, while being kept distinct from each other by gaskets or seals. In at least one embodiment, a second flow may be to enable secondary coolant to flow through provided ports **268**, **262**, and an associated path **264**.

[0078] In at least one embodiment, flow controllers **278** may be associated with a fluid inlet **276** and outlet portions at a server manifold **204** instead of provided flow controllers **280** at respective cold plates. In at least one embodiment, a first flow uses only local coolant and may be enabled when a failure is determined in a secondary cooling loop or a primary cooling loop, so that a secondary coolant is unable to effectively extract heat from at least one computing device. In at least one embodiment, a failure may be that a secondary coolant is not sufficiently cooled via a CDU and so it may be unable to extract sufficient heat of at least one computing device via its associated cold plate.

[0079] In at least one embodiment, cold plate-level features **290** as illustrated in FIG. 2B can be associated with an intelligent cold plate system having active and passive features for a datacenter cooling system. In at least one embodiment, cold plate-level features **290** illustrate a heat pipe within a section or cavity **282B** of a cold plate **250**. In at least one embodiment, a section or cavity **282A** may be overlapping, under, above, or adjacent (such as, on a side) of a section or cavity **282B**. In at least one embodiment, a cold plate, such as dual-cooling or a regular cold plate having sections or cavities **282A**, **282B**, is part of a datacenter cooling system.

[0080] In at least one embodiment, a cold plate having sections or cavities **282A**, **282B** therein may include microchannels, as described in reference to FIG. 2A, and a heat pipe **296**. In at least one embodiment, provided microchan-

nels of a cold plate support a first fluid, such as a secondary coolant or a local coolant, in an active mode of operation of a cold plate **250**. In at least one embodiment, a first fluid may be forced through microchannels of a cold plate **250** by at least one flow controller **280**. In at least one embodiment, at least one flow controllers may provide suction or pumping action to cause force on a first fluid that enables its flow through provided microchannels **264**, **270**.

[0081] In at least one embodiment, a heat pipe **296** of a cold plate **250** supports a second fluid, such as a refrigerant. In at least one embodiment, a refrigerant may be R134A. In at least one embodiment, a heat pipe **296** supports a second fluid in a passive mode of operation of a cold plate **250**. In at least one embodiment, a second fluid may be forced to circulate in a heat pipe **296** by differential temperature and capillary actions between an evaporator section **294A** and a condenser section **294B** of a heat pipe **250**. In at least one embodiment, capillary action may be referred to as wicking action or effect, capillary effect, or capillarity, and is in reference to fluid flow through narrow channels that may occur without external actions (such as forces). In at least one embodiment, change in a differential temperature may cause faster circulation of a second fluid in a heat pipe. In at least one embodiment, change in a differential temperature may be enabled by change in air flow from an associated fan **298**. In at least one embodiment, a faster speed of a fan **298** may cause further cooling of a condenser section **294B** relative to an evaporator section **294A**, which in turn causes faster movement of second fluid from an evaporator section to a condenser section of a heat pipe. In at least one embodiment, second fluid is in reference to different phases that may be part of properties for a fluid, including a vapor phase and a liquid phase.

[0082] In at least one embodiment, a heat pipe **296** may be formed of concentric tubes having an outer shell **284**, an inner shell **286**, a space having a wicking mesh or capillary tubes **292**, and a hollow pipe section **288**. In at least one embodiment, a heat pipe's evaporator section **294A** is within a cold plate **250** and its condenser section **294B** extends out of (or is located out of) a cold plate **250**. In at least one embodiment, a condenser section **294B** of a heat pipe is located in an area that is remote from a cold plate **250**. In at least one embodiment, an area remote from a cold plate **250** may be unoccupied portions of a server tray or box, and is at least one portion that is not disturbed by heat dissipation. In at least one embodiment, an area remote of a cold plate **250** is external to a server tray or box but within a rack. In at least one embodiment, a second fluid used within a heat pipe **296** may be a phase-changeable fluid capable of transitioning from vapor to liquid phases as part of heat extraction and dissipation features of a heat sink. In at least one embodiment, a second fluid is a refrigerant that is in vapor form within an evaporator section because of a higher temperature in an environment around an evaporator section.

[0083] In at least one embodiment, a second fluid extracts heat from at least one associated computing device. In at least one embodiment, a second fluid retains heat from an associated computing device and transforms from a fluid state, within a wicking mesh **292**, to become a vapor. In at least one embodiment, a vapor-form second fluid moves from an evaporator section **294A** to a condenser section **294B** because of, in part, a temperature difference between these sections. In at least one embodiment, a condenser

section **294B** is at a cooler temperature and allows dissipation of extracted heat retained within a second fluid. In at least one embodiment, dissipation of heat cools a second fluid and causes a phase change from vapor to liquid state. In at least one embodiment, liquid form or state of a second fluid falls on a wicking mesh or capillary tubes **292** within a shell area formed by shells **284**, **286** of a condenser section **294B**. In at least one embodiment, capillary action of a wicking mesh or of capillary tubes causes movement of a liquid form second liquid from a condenser section **294B** to an evaporator section **294A** of a heat pipe **296**. In at least one embodiment, this cycle is repeated by differential temperature and by capillary action between an evaporator section **294A** and a condenser section **294B** of a heat pipe **296**. In at least one embodiment, a fan **298** provides forced air cooling of a condenser section **294B** of a heat pipe **296**.

[0084] In at least one embodiment, at least one processor may determine a temperature associated with a computing device or a cold plate **250**. In at least one embodiment, at least one processor may cause at least one flow controller **280** to pass first fluid through microchannels **264**; **270** to enable an active mode for a cold plate **250**. In at least one embodiment, at least one processor may be adapted to determine a temperature or power associated with a computing device. In at least one embodiment, as power consumption increases for a computing device, heat generated from a computing device increases. In at least one embodiment, as heat generated from a computing device increases, its cooling requirement also increases. In at least one embodiment, therefore, it is possible to determine (or predict) cooling requirements for a computing device by its power consumption.

[0085] In at least one embodiment, at least one processor may cause one or more of an active mode or a combination of an active mode and a passive mode for a cold plate. In at least one embodiment, at least one processor controls at least one flow controller to initiate an active mode, but a passive mode may be uncontrolled and available by default in a cold plate having a heat pipe with microchannels. In at least one embodiment, a passive mode may be partly controllable by changing a temperature at a condenser section of a heat pipe. In at least one embodiment, a passive mode is always enabled but may be also be enabled to extract further heat by a further differential temperature for a heat pipe. In at least one embodiment, as a differential temperature enables circulation of second fluid in a heat pipe, a higher differential temperature, such as by more cooling at a condenser section, may cause faster circulation of second fluid. In at least one embodiment, a level of circulation of second fluid is always available in a heat pipe and so a heat pipe is always enabled. In at least one embodiment, a cold plate **250** may be associated with a computing device and may include ports as described with respect to FIG. 2A. In at least one embodiment, provided ports of a cold plate may support entry and egress of a first fluid, such as a secondary coolant or a local coolant, through provided microchannels. In at least one embodiment, a first fluid may be maintained distinctly from a second fluid of a heat pipe by at least different cavities or sections of a cold plate **250**.

[0086] In at least one embodiment, at least one processor may be adapted to receive sensor inputs from sensors associated with a computing device or a cold plate **250**. In at least one embodiment, at least one processor may be adapted to determine a first cooling requirement associated

with a passive mode, a second cooling requirement associated with both active mode and passive mode for a cold plate **250**, based in part on sensor inputs. In at least one embodiment, further circulation (by an enforced differential temperature between sections of a heat pipe by a fan) may enable some form of control for a passive mode of a cold plate. In at least one embodiment, one or more neural networks may be adapted to receive sensor inputs and to infer a first cooling requirement and a second cooling requirement associated with a computing device or with a second fluid. In at least one embodiment, a first cooling requirement includes a temperature which may be addressed by a heat pipe having higher differential temperature and no further dissipation of heat may be possible, thereby classifying further cooling requirements as a second cooling requirement.

[0087] In at least one embodiment, a second fluid provides passive cooling in a heat sink and when a determination is made that a second fluid is retaining too much heat after higher differential temperature is enabled for a heat pipe, then an active mode may be enabled for a cold plate. In at least one embodiment, an active mode provides secondary coolant or local coolant for supplemental cooling in a cold plate **250** (and for an associated computing device). In at least one embodiment, passive mode may be controller by further fan cooling of a condenser section of a heat pipe till no further benefit may be offered before an active mode is enabled.

[0088] In at least one embodiment, at least one processor may be adapted to receive sensor inputs from temperature sensors and from power sensors associated with a computing device. In at least one embodiment, a temperature sensor senses a temperature of a computing device and provides input to at least one processor. In at least one embodiment, at least one processor can determine a first cooling requirement and a first power requirement associated with a passive mode.

[0089] In at least one embodiment, a first cooling requirement may be a threshold within which passive mode cooling is sufficient to keep at least one computing device within operating temperatures. In at least one embodiment, a second cooling requirement and a second power requirement associated with an active mode. In at least one embodiment, a second cooling requirement and a second power requirement may be associated with both an active mode and a passive mode for a cold plate. In at least one embodiment, a second cooling requirement may be a threshold within which both modes of cooling are sufficient to keep at least one computing device within operating temperatures. In at least one embodiment such determination of cooling requirements may be made based in part on sensor inputs.

[0090] In at least one embodiment, a second fluid includes a phase-changeable fluid may be operative in a vapor phase within an evaporator section of a heat pipe, and a liquid or partly liquid phase may be formed for a second fluid within a condenser section of a heat pipe. In at least one embodiment, a heat pipe includes a wicking material to enable passive movement of a second fluid from a condenser section to an evaporator section as described in reference to FIG. 2B.

[0091] In at least one embodiment, a condenser section **294B** can function as a liquid-to-air heat exchanger, along with a provided fan **298**. In at least one embodiment, a liquid-to-air heat exchanger may be associated with a heat

pipe to enable dissipation of heat or cooling of a second fluid at a condenser section. In at least one embodiment, a condenser section may be located remotely from a cold plate.

[0092] In at least one embodiment, rack-level features 300 as illustrated in FIG. 3 can be associated with an intelligent cold plate system having active and passive features for a datacenter cooling system. In at least one embodiment, rack-level features 300 include a rack 302 having brackets 304, 306 to hang cooling manifolds 314A, B. In at least one embodiment, while a rack 330 is separately illustrated from a rack 302, this rack 330 may be illustrative of a rear perspective view of a rack 302. In at least one embodiment, as such, brackets 334, 336 provided on rack 330 are perspective views of brackets 304, 306 provided on rack 302. In at least one embodiment, brackets 304, 306 provided for a rack are flat structures against an inner wall of a rack. In at least one embodiment, brackets 304, 306 provided for a rack extend from an inner wall of a rack. In at least one embodiment, brackets 304, 306 provided for a rack are affixed to an inner wall of a rack and have multiple mounting points facing one or more directions, including inside or towards a rear of a rack.

[0093] In at least one embodiment, cooling manifolds 314A, B may be provided to pass secondary coolant or local coolant between server-level features 200 (and illustrated in FIG. 3 as server trays or boxes 308) and a CDU (such as CDU 406 of FIG. 4) of a secondary cooling loop or a local cooling loop of a datacenter cooling system. In at least one embodiment, different CDUs may serve different racks. In at least one embodiment, different rack cooling manifolds may be distinctly part of a secondary cooling loop and a local cooling loop. In at least one embodiment, at least one server tray or box 308 (such as a bottom-most server tray or box 308 in rack 302) may be designated as a control system for an intelligent cold plate system having active and passive features so that, if a local coolant is used, such a system may be isolated from a secondary cooling loop for an active mode of a cold plate. In at least one embodiment, a control system in a server tray or box 308 may include safety feature (such as sensors to provide sensor data or proper function), communication features (to communicate with at least one flow controller for an active mode and to communicate to an external monitor), power features to power one or more flow controllers and at least one processor (and its related features), and control features offered by at least one processor that may be associated with at least one flow controller.

[0094] In at least one embodiment, row manifold 350 may be part of a secondary cooling loop to feed an inlet rack manifold 314A via provided lines 310A, 310. In at least one embodiment, secondary coolant proceeds via a provided line 316 to cold plate 326 to extract heat from associated computing device 324 within a server 308 in an active mode of operation of a cold plate 326; and proceeds via a provided line 318 to outlet rack manifold 314B and through provided lines 312, 312A, and back into a same or a different row manifold 350. In at least one embodiment, an intelligent cold plate system with active and passive features can work independent of a secondary cooling loop via provided lines 312B, 310B for a local cooling loop. In at least one embodiment, one or more diverter flow controllers 310C, 312C isolates each of a secondary cooling loop and a local cooling loop. In at least one embodiment, an intelligent cold plate system with active and passive features can work indepen-

dent of a secondary cooling loop and a local cooling loop by action of a heat pipe alone. In at least one embodiment, this may be to an extent that cooling requirement for at least one associated computing device is within a temperature range addressable by a heat pipe alone.

[0095] In at least one embodiment, datacenter-level features 400 as illustrated in FIG. 4 can be associated with an intelligent cold plate system having active and passive features for a datacenter cooling system. In at least one embodiment, datacenter-level features 400, within a datacenter 402, may include racks 404 for hosting one or more server trays or boxes; one or more CDUs 406 for exchanging heat between a secondary cooling loop 412 and a primary cooling loop 422; one or more row manifolds 410 for distributing coolant from a CDU 406; and associated various flow controllers 424, and inlet and outlet lines 412, 414, 416, 418.

[0096] In at least one embodiment, an intelligent cold plate system with active and passive features are provided on each of rear doors of each of provided racks 404 in a datacenter 402. In at least one embodiment, an aisle behind racks 404 is a hot aisle for discharging heat from at least one computing device in at least one rack during a first and a second mode of operation of a datacenter cooling system. In at least one embodiment, different row manifolds may be associated with different racks. In at least one embodiment, different coolant may be a chemical match or mismatch with respect to a local coolant. In at least one embodiment, different fluid sources are provided as redundant features to different CDUs depending on chemistries of different secondary coolant used with each of different provided CDUs. In at least one embodiment, there need not be a secondary cooling loop and CDU for one or more racks 404. In at least one embodiment, these racks not associated with a secondary cooling loop may be sufficiently addressed by an intelligent cold plate system with active and passive features.

[0097] In at least one embodiment, at least one processor enables an intelligent cold plate system having active and passive features for a datacenter cooling system. In at least one embodiment, at least one processor may include one or more circuits and may be associated with a cold plate. In at least one embodiment, a cold plate includes microchannels and a heat pipe. In at least one embodiment, a processor can cause an active mode of operation for a cold plate so that microchannels support a first fluid forced through microchannels therein, by at least one flow controller. In at least one embodiment, at least one processor may cause an active mode to supplement a passive mode of operation of a cold plate. In at least one embodiment, at least one processor can also vary an effect within a passive mode of operation for a cold plate.

[0098] In at least one embodiment, a passive mode of operation may be always enabled by a second fluid (different from a first fluid allowed in microchannels of a cold plate) that is caused to circulate through a heat pipe due to differential temperature and other actions therein. In at least one embodiment, a heat pipe supports a second fluid forced to circulate therein by differential temperature and capillary actions between an evaporator section and a condenser section of a heat pipe. In at least one embodiment, enabling a higher differential temperature may cause faster circulation of second fluid within a heat pipe. In at least one embodiment, such higher differential temperature may be enabled by a fan providing forced air cooling and that may be

controlled by a processor providing inputs to a power controller associated with a fan and a heat pipe. In at least one embodiment, however, a heat pipe passively circulates second fluid by default without any asserted change of a differential temperature associated therewith.

[0099] In at least one embodiment, an output of one or more circuits can provide signals for at least one flow controller associated with a cold plate to enable flow of a first fluid through provided microchannels. In at least one embodiment, an input of one or more circuits may receive sensor inputs from sensors associated with at least one computing device and a second fluid. In at least one embodiment, a processor can determine a cooling requirement based in part on sensor inputs.

[0100] In at least one embodiment, one or more neural networks may be provided to receive sensor inputs and to infer cooling requirements as described in reference to at least FIGS. 2A and 2B. In at least one embodiment, one or more neural networks can infer a power requirement threshold, such as a power requirement that known to cause higher heat generation, for at least one computing device. In at least one embodiment, one or more circuits of a processor can cause an active mode for a cold plate based at least in part on an inference of a power requirement.

[0101] In at least one embodiment, a processor includes one or more circuits and may be associated with a cold plate that has microchannels and a heat pipe. In at least one embodiment, one or more circuits can train one or more neural networks to infer a cooling requirement from sensor inputs of sensors associated with a computing device or a cold plate. In at least one embodiment, a processor can cause an active mode of operation for a cold plate so that microchannels therein support a first fluid forced through microchannels by at least one flow controller. In at least one embodiment, a processor can supplement a passive mode of operation of a cold plate by an active mode.

[0102] In at least one embodiment, an output of one or more circuits can provide signals for at least one flow controller to enable flow of first fluid through microchannels of a cold plate. In at least one embodiment, one or more neural networks may be part of one or more circuits of at least one processor. In at least one embodiment, one or more neural networks can receive sensor inputs from sensors associated with at least one cold plate or a computing device. In at least one embodiment, one or more neural networks may be trained to infer cooling requirements as part of an analysis of prior sensor inputs and prior cooling requirements. In at least one embodiment, for a temperature sensed from a computing device, prior cooling requirements may be associated with maintaining a passive mode of cooling or with supplementing a passive mode with an active mode of cooling.

[0103] In at least one embodiment, an output of one or more circuits of a process can provide signals to cause at least one flow controller to activate or deactivate an active mode of cooling for a cold plate. In at least one embodiment, an input of one or more circuits of a process can receive sensor inputs associated with a temperature from at least one computing device or second fluid associated with a heat pipe. In at least one embodiment, one or more neural networks may be trained to infer a change in coolant state has occurred based in part on a temperature and on prior temperatures. In at least one embodiment, a temperature sensed by a sensor may be indicative of passive mode of

operation, by a heat pipe, not providing sufficient cooling for a cold plate (and its associated computing device). In at least one embodiment, one or more circuits can cause an active mode of operation along with an existing passive mode of operation for a cold plate.

[0104] In at least one embodiment, a control logic may be to determine a change in a coolant state, such as a failure or a change in a secondary cooling loop (such as a CDU and cooling manifolds), a local cooling loop, or a primary cooling loop (such as a chilling facility, cooling manifolds, and also an associated CDU). In at least one embodiment, a failure may also occur with a cooling manifold requiring replacement. In at least one embodiment, a change in a cooling loop may occur if coolant flow is stopped as only passive cooling is required under certain heat generation thresholds that may be addressed by passive cooling alone. In at least one embodiment, as a passive mode may be occurring by default, a processor determines if to enable an active mode if an active mode is not already ongoing for a cold plate.

[0105] In at least one embodiment, a processor may engage a further local coolant in addition to a secondary coolant as part of a response to a cooling requirement if still further cooling is required in a second mode, in supplemental action to an existing passive mode of operation of a cold plate. In at least one embodiment, a control logic may cause at least one flow controller to provide a coolant response, such as by engaging a first mode for a cold plate, in addition to a passive mode already existing with a cold plate. In at least one embodiment, a secondary or local cooling loop having a fluid source provides local coolant or secondary coolant for at least one computing device via a cold plate and its microchannels in a second mode to supplement a first mode (passive mode) of its operation.

[0106] In at least one embodiment, a processor may include one or more circuits that further includes one or more neural networks. In at least one embodiment, one or more neural networks of a processor can infer a cooling requirement from sensor inputs of sensors associated with a computing device or a cold plate. In at least one embodiment, an input of one or more circuits of a processor can receive sensor inputs associated with a temperature from at least one computing device or a second fluid (of a heat pipe). In at least one embodiment, one or more neural networks can infer a change in coolant state has occurred based in part on a temperature received and on prior temperatures.

[0107] In at least one embodiment, if a change is inferred indicating that no coolant is flowing under a second mode, but that temperature has been rising for multiple cycles of sensing from sensors, then a processor can cause an active mode of operation for a cold plate. In at least one embodiment, this may be caused by a malfunction in a passive mode that is addressed by an active mode available in a cold plate. In at least one embodiment, an active mode of operation is so that microchannels within a cold plate support flow of a first fluid by an action of at least one flow controller. In at least one embodiment, an active mode supplements a passive mode of operation of a cold plate. In at least one embodiment, a passive mode for a cold plate is enabled by a heat pipe in which circulation of a second fluid therein occurs by differential temperature and capillary actions between an evaporator section and a condenser section of a heat pipe.

[0108] In at least one embodiment, each of at least one processor described throughout FIGS. 1-4 has inference and/or training logic **1815** that may include, without limitation, code and/or data storage **1801** to store forward and/or output weight and/or input/output data, and/or other parameters to configure neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, training logic **1815** may include, or be coupled to code and/or data storage **1801** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information may be to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which such code corresponds. In at least one embodiment, code and/or data storage **1801** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during forward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, any portion of code and/or data storage **1801** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

[0109] In at least one embodiment, an inference and/or training logic **1815** of at least one processor may be part of a building management system (BMS) for controlling flow controllers at one or more of a server-level, a rack-level, cold plate-level, and a row-level. In at least one embodiment, a determination to engage a flow controller associated with a local cooling loop, an intelligent cold plate system with active and passive features, a CDU, or other cooling manifolds may be provided to one or more neural networks of an inference and/or training logic **1815** to cause one or more neural networks to infer which flow controllers to gracefully engage or disengage for coolant requirements for one or more cold plates, servers, or racks from either a passive mode of a cold plate or an active mode combined with a passive mode for a cold plate. In at least one embodiment, increase or decrease of fluid flow through a cold plate may be enabled by flow controllers that are controlled by an inference and/or training logic **1815** of at least one processor associated with control logic that is associated with a local cooling loop, a secondary cooling loop, or even a heat pipe.

[0110] In at least one embodiment, at least one processor may be associated with a local cooling loop and with a secondary cooling loop. In at least one embodiment, at least one processor may be associated with an intelligent cold plate system having active and passive features. In at least one embodiment, at least one processor includes control logic, such as inference and/or training logic **1815** and is associated with at least one flow controller. In at least one embodiment, at least one flow controller may have their own respective processor or micro controller. In at least one embodiment, a processor or a micro controller performs instructions sent to it from a control logic. In at least one embodiment, a control logic may be to determine a change in a coolant state, such as a failure or a change in a secondary cooling loop (such as a CDU and cooling manifolds), a local cooling loop, or a primary cooling loop (such as a chilling facility, cooling manifolds, and also an associated CDU). In

at least one embodiment, a failure or change may also occur with a cooling manifold requiring replacement or when cooling by a passive mode is sufficient. In at least one embodiment, a control logic may cause at least one flow controller to provide a coolant response, such as by engaging an active mode or disengaging an active mode for a cold plate.

[0111] In at least one embodiment, a control logic may cause a first signal to at least one flow controller to enable a stopping of a secondary coolant from a secondary cooling loop as part of a coolant response. In at least one embodiment, a control logic may cause a second signal to at least one flow controller to enable a starting of a local coolant from a local cooling loop as part of a coolant response. In at least one embodiment, an intelligent cold plate system having active and passive features may be used to provide redundant cooling to a secondary cooling loop by active or passive features therein using a local coolant. In at least one embodiment, a local coolant provides sufficient cooling to last a period determined by a service level agreement (SLA) providing downtime for a primary or a secondary cooling loop. In at least one embodiment, these signals may be part of an active mode initiation for a dual-cooling cold plate. In at least one embodiment, a control logic may receive sensor inputs from sensors associated with secondary coolant of a CDU, local coolant, a heat pipe fluid, and/or at least one computing device. In at least one embodiment, at least one processor can determine a change in a coolant state based in part on sensor inputs. In at least one embodiment, one or more neural networks of an inference and/or training logic **1815** may be adapted to receive sensor inputs and to infer a change in a coolant state.

[0112] In at least one embodiment, at least one processor may include one or more circuits for one or more neural networks, such as an inference and/or training logic **1815**. In at least one embodiment, an inference and/or training logic **1815** may be adapted to infer, from sensor inputs associated with at least one server or at least one rack, a change in a coolant state, such as coolant from a CDU being effective, ineffective, or retaining too much heat upon entry into a rack. In at least one embodiment, one or more circuits may be adapted to cause at least one flow controller to provide a coolant response from a local cooling loop to supplement a heat pipe.

[0113] In at least one embodiment, control logic associated with one or more circuits may cause a first signal (along with any associated signals) to at least one flow controller to enable a coolant response—either from a secondary cooling loop or a local cooling loop having an intelligent cold plate system with active and passive features. In at least one embodiment, a second signal may be provided to at least one flow controller that may be associated with a heat pipe to disengage it from operation so that there is no passive mode for a cold plate. In at least one embodiment, a distributed or an integrated architecture is enabled by one or more circuits of at least one processor. In at least one embodiment, a distributed architecture may be supported by distinctly located circuits of one or more circuits.

[0114] In at least one embodiment, one or more neural networks of an inference and/or training logic **1815** may be adapted to infer that an increase or a decrease in cooling requirements of at least one computing component of at least one server. In at least one embodiment, one or more circuits may be adapted to cause a cooling loop to economically

address decreased cooling requirements or to supplement increased cooling requirements for at least one computing component. In at least one embodiment, enabling a cooling loop represents a coolant response from a local cooling loop to preempt a respective increase or a respective decrease in cooling requirements of at least one computing component of at least one server based in part on workload sent to at least one computing component.

[0115] In at least one embodiment, at least one processor includes one or more circuits, such as an inference and/or training logic 1815, to train one or more neural networks to make inferences from provided data. In at least one embodiment, inference and/or training logic 1815 may infer, from sensor inputs associated with at least one server or at least one rack, a change in a coolant state. In at least one embodiment, an inference may be used to enable one or more circuits to cause at least one flow controller of a local cooling loop to provide a coolant response. In at least one embodiment, a coolant response may be to cause a coolant response from a local cooling loop to absorb heat into a local coolant of a cooling manifold, and to exchange absorbed heat to an environment, instead of a secondary cooling loop having a CDU.

[0116] In at least one embodiment, one or more circuits may be adapted to train one or more neural networks to infer that an increase or a decrease in cooling requirements of at least one computing component of at least one server. In at least one embodiment, one or more circuits may be adapted to train one or more neural networks to infer that an increase or a decrease in flow output from a secondary cooling loop is associated with an improper flow of secondary coolant because of a failed CDU or a respective increase or a respective decrease in power requirements of at least one computing component of at least one server.

[0117] In at least one embodiment, one or more neural networks may be trained to make inferences by prior associated heat features or cooling requirements from computing devices, servers, or racks, and cooling capacity or capabilities indicated by a fluid source of a local cooling loop. In at least one embodiment, prior cooling requirements satisfied by a local cooling loop may be used to cause one or more neural networks to make similar inferences for future similar cooling requirements (in consideration of small variations there from) to be satisfied by adjusting one or more flow controllers to engage a local cooling loop.

[0118] FIG. 5 illustrates a method 500 associated with a datacenter cooling system of FIG. 2A—4, according to at least one embodiment. In at least one embodiment, a method 500 includes a step 502 for providing a cold plate having microchannels and a heat pipe. In at least one embodiment, step 504 is for determining at least one cooling requirement for at least one computing device. In at least one embodiment, a step 506 may be performed to verify that a determination for at least one cooling requirement for at least one computing device exists. In at least one embodiment, step 508 may be then performed for enabling microchannels of a cold plate to support a first fluid in an active mode of operation of a cold plate. In at least one embodiment, a first fluid in an active mode may be forced through microchannels of a cold plate by at least one flow controller. In at least one embodiment, step 504 may be performed if a verification in step 506 provides that further determination is required.

[0119] In at least one embodiment, method 500 further includes step 510 for enabling a heat pipe to support a second fluid in a passive mode of operation of a cold plate. In at least one embodiment, second fluid may be forced to circulate in a heat pipe by differential temperature and capillary actions between an evaporator section and a condenser section of a heat pipe. In at least one embodiment, an active mode supplements a passive mode of operation of a cold plate. In at least one embodiment, a passive mode may be disengaged from a cold plate so that only an active mode may be provided. In at least one embodiment, step 510 may be performed by providing a wicking mesh or providing capillary tubes within shells of a pipe, while leaving a center portion of a pipe hollow. In at least one embodiment, metal powder may be provided within shells (such as inner surfaces of concentric tubes) and may be sintered to enable a wicking mesh to provide a sintered heat pipe. In at least one embodiment, sintering for a sintered wick is determined according to a second fluid to be used, according to a length of a heat pipe, and according to an amount of heat transfer to be enabled for a heat pipe. In at least one embodiment, sintering may be performed according to a temperature, a time, an atmosphere, and a position determined for a wicking mesh to be formed within concentric tubes forming a heat pipe. In at least one embodiment, meshing or springing may be other processes to provide a wicking mesh or capillary tubes within concentric tubes to form a heat pipe.

[0120] In at least one embodiment, method 500 may include a further step or a sub-step for determining, using at least one processor, a temperature associated with a computing device or a cold plate. In at least one embodiment, this information may be used in a further step or sub-step for causing at least one flow controller associated with a cold plate to pass first fluid through microchannels therein and to enable an active mode for a cold plate, as a result. In at least one embodiment, method 500 may include a further step or a sub-step for determining, using at least one processor, a temperature or power associated with a computing device. In at least one embodiment, this may result in a step or sub-step of causing one or more of an active mode or a combination of an active mode and a passive mode for a cold plate.

[0121] In at least one embodiment, method 500 may include a further step or a sub-step for receiving, in at least one processor, sensor inputs from sensors associated with a computing device or a cold plate. In at least one embodiment, method 500 may include a further step or a sub-step for determining, by at least one processor, a first cooling requirement associated with a passive mode, a second cooling requirement associated both modes, forming a combination of an active mode and a passive mode, based in part on sensor inputs received.

[0122] In at least one embodiment, method 500 may include a further step or a sub-step for receiving, in at least one processor, sensor inputs from temperature sensors and from power sensors associated with a computing device. In at least one embodiment, sensor inputs may be used for determining various cooling requirements. In at least one embodiment, at least one processor may use such information to determine a first cooling requirement and a first power requirement associated with a passive mode for a cold plate. In at least one embodiment, at least one processor may use such information to determine a second cooling require-

ment and a second power requirement associated with both modes, an active mode and a passive mode, based in part on sensor inputs received.

Servers and Data Centers

[0123] The following figures set forth, without limitation, exemplary network server and datacenter based systems that can be used to implement at least one embodiment.

[0124] FIG. 6 illustrates a distributed system 600, in accordance with at least one embodiment. In at least one embodiment, distributed system 600 includes one or more client computing devices 602, 604, 606, and 608, which are configured to execute and operate a client application such as a web browser, proprietary client, and/or variations thereof over one or more network(s) 610. In at least one embodiment, server 612 may be communicatively coupled with remote client computing devices 602, 604, 606, and 608 via network 610.

[0125] In at least one embodiment, server 612 may be adapted to run one or more services or software applications such as services and applications that may manage session activity of single sign-on (SSO) access across multiple datacenters. In at least one embodiment, server 612 may also provide other services or software applications can include non-virtual and virtual environments. In at least one embodiment, these services may be offered as web-based or cloud services or under a Software as a Service (SaaS) model to users of client computing devices 602, 604, 606, and/or 608. In at least one embodiment, users operating client computing devices 602, 604, 606, and/or 608 may in turn utilize one or more client applications to interact with server 612 to utilize services provided by these components.

[0126] In at least one embodiment, software components 618, 620 and 622 of system 600 are implemented on server 612. In at least one embodiment, one or more components of system 600 and/or services provided by these components may also be implemented by one or more of client computing devices 602, 604, 606, and/or 608. In at least one embodiment, users operating client computing devices may then utilize one or more client applications to use services provided by these components. In at least one embodiment, these components may be implemented in hardware, firmware, software, or combinations thereof. It should be appreciated that various different system configurations are possible, which may be different from distributed system 600. The embodiment shown in FIG. 6 is thus at least one embodiment of a distributed system for implementing an embodiment system and is not intended to be limiting.

[0127] In at least one embodiment, client computing devices 602, 604, 606, and/or 608 may include various types of computing systems. In at least one embodiment, a client computing device may include portable handheld devices (e.g., an iPhone®, cellular telephone, an iPad®, computing tablet, a personal digital assistant (PDA)) or wearable devices (e.g., a Google Glass® head mounted display), running software such as Microsoft Windows Mobile®, and/or a variety of mobile operating systems such as iOS, Windows Phone, Android, BlackBerry 10, Palm OS, and/or variations thereof. In at least one embodiment, devices may support various applications such as various Internet-related apps, e-mail, short message service (SMS) applications, and may use various other communication protocols. In at least one embodiment, client computing devices may also include general purpose personal computers including, by way of at

least one embodiment, personal computers and/or laptop computers running various versions of Microsoft Windows®, Apple Macintosh®, and/or Linux operating systems.

[0128] In at least one embodiment, client computing devices can be workstation computers running any of a variety of commercially-available UNIX® or UNIX-like operating systems, including without limitation a variety of GNU/Linux operating systems, such as Google Chrome OS. In at least one embodiment, client computing devices may also include electronic devices such as a thin-client computer, an Internet-enabled gaming system (e.g., a Microsoft Xbox gaming console with or without a Kinect® gesture input device), and/or a personal messaging device, capable of communicating over network(s) 610. Although distributed system 600 in FIG. 6 is shown with four client computing devices, any number of client computing devices may be supported. Other devices, such as devices with sensors, etc., may interact with server 612.

[0129] In at least one embodiment, network(s) 610 in distributed system 600 may be any type of network that can support data communications using any of a variety of available protocols, including without limitation TCP/IP (transmission control protocol/Internet protocol), SNA (systems network architecture), IPX (Internet packet exchange), AppleTalk, and/or variations thereof. In at least one embodiment, network(s) 610 can be a local area network (LAN), networks based on Ethernet, Token-Ring, a wide-area network, Internet, a virtual network, a virtual private network (VPN), an intranet, an extranet, a public switched telephone network (PSTN), an infra-red network, a wireless network (e.g., a network operating under any of the Institute of Electrical and Electronics (IEEE) 802.11 suite of protocols, Bluetooth®, and/or any other wireless protocol), and/or any combination of these and/or other networks.

[0130] In at least one embodiment, server 612 may be composed of one or more general purpose computers, specialized server computers (including, by way of at least one embodiment, PC (personal computer) servers, UNIX® servers, mid-range servers, mainframe computers, rack-mounted servers, etc.), server farms, server clusters, or any other appropriate arrangement and/or combination. In at least one embodiment, server 612 can include one or more virtual machines running virtual operating systems, or other computing architectures involving virtualization. In at least one embodiment, one or more flexible pools of logical storage devices can be virtualized to maintain virtual storage devices for a server. In at least one embodiment, virtual networks can be controlled by server 612 using software defined networking. In at least one embodiment, server 612 may be adapted to run one or more services or software applications.

[0131] In at least one embodiment, server 612 may run any operating system, as well as any commercially available server operating system. In at least one embodiment, server 612 may also run any of a variety of additional server applications and/or mid-tier applications, including HTTP (hypertext transport protocol) servers, FTP (file transfer protocol) servers, CGI (common gateway interface) servers, JAVA® servers, database servers, and/or variations thereof. In at least one embodiment, exemplary database servers include without limitation those commercially available from Oracle, Microsoft, Sybase, IBM (International Business Machines), and/or variations thereof.

[0132] In at least one embodiment, server 612 may include one or more applications to analyze and consolidate data feeds and/or event updates received from users of client computing devices 602, 604, 606, and 608. In at least one embodiment, data feeds and/or event updates may include, but are not limited to, Twitter® feeds, Facebook® updates or real-time updates received from one or more third party information sources and continuous data streams, which may include real-time events related to sensor data applications, financial tickers, network performance measuring tools (e.g., network monitoring and traffic management applications), clickstream analysis tools, automobile traffic monitoring, and/or variations thereof. In at least one embodiment, server 612 may also include one or more applications to display data feeds and/or real-time events via one or more display devices of client computing devices 602, 604, 606, and 608.

[0133] In at least one embodiment, distributed system 600 may also include one or more databases 614 and 616. In at least one embodiment, databases may provide a mechanism for storing information such as user interactions information, usage patterns information, adaptation rules information, and other information. In at least one embodiment, databases 614 and 616 may reside in a variety of locations. In at least one embodiment, one or more of databases 614 and 616 may reside on a non-transitory storage medium local to (and/or resident in) server 612. In at least one embodiment, databases 614 and 616 may be remote from server 612 and in communication with server 612 via a network-based or dedicated connection. In at least one embodiment, databases 614 and 616 may reside in a storage-area network (SAN). In at least one embodiment, any necessary files for performing functions attributed to server 612 may be stored locally on server 612 and/or remotely, as appropriate. In at least one embodiment, databases 614 and 616 may include relational databases, such as databases that are adapted to store, update, and retrieve data in response to SQL-formatted commands.

[0134] FIG. 7 illustrates an exemplary datacenter 700, in accordance with at least one embodiment. In at least one embodiment, datacenter 700 includes, without limitation, a datacenter infrastructure layer 710, a framework layer 720, a software layer 730 and an application layer 740.

[0135] In at least one embodiment, as shown in FIG. 7, datacenter infrastructure layer 710 may include a resource orchestrator 712, grouped computing resources 714, and node computing resources (“node C.R.s”) 716(1)-716(N), where “N” represents any whole, positive integer. In at least one embodiment, node C.R.s 716(1)-716(N) may include, but are not limited to, any number of central processing units (“CPUs”) or other processors (including accelerators, field programmable gate arrays (“FPGAs”), graphics processors, etc.), memory devices (e.g., dynamic read-only memory), storage devices (e.g., solid state or disk drives), network input/output (“NW I/O”) devices, network switches, virtual machines (“VMs”), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s 716(1)-716(N) may be a server having one or more of above-mentioned computing resources.

[0136] In at least one embodiment, grouped computing resources 714 may include separate groupings of node C.R.s housed within one or more racks (not shown), or many racks housed in datacenters at various geographical locations (also

not shown). Separate groupings of node C.R.s within grouped computing resources 714 may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s including CPUs or processors may grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks may also include any number of power modules, cooling modules, and network switches, in any combination.

[0137] In at least one embodiment, resource orchestrator 712 may configure or otherwise control one or more node C.R.s 716(1)-716(N) and/or grouped computing resources 714. In at least one embodiment, resource orchestrator 712 may include a software design infrastructure (“SDI”) management entity for datacenter 700. In at least one embodiment, resource orchestrator 712 may include hardware, software or some combination thereof.

[0138] In at least one embodiment, as shown in FIG. 7, framework layer 720 includes, without limitation, a job scheduler 732, a configuration manager 734, a resource manager 736 and a distributed file system 738. In at least one embodiment, framework layer 720 may include a framework to support software 752 of software layer 730 and/or one or more application(s) 742 of application layer 740. In at least one embodiment, software 752 or application(s) 742 may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. In at least one embodiment, framework layer 720 may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter “Spark”) that may utilize distributed file system 738 for large-scale data processing (e.g., “big data”). In at least one embodiment, job scheduler 732 may include a Spark driver to facilitate scheduling of workloads supported by various layers of datacenter 700. In at least one embodiment, configuration manager 734 may be capable of configuring different layers such as software layer 730 and framework layer 720, including Spark and distributed file system 738 for supporting large-scale data processing. In at least one embodiment, resource manager 736 may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system 738 and job scheduler 732. In at least one embodiment, clustered or grouped computing resources may include grouped computing resource 714 at datacenter infrastructure layer 710. In at least one embodiment, resource manager 736 may coordinate with resource orchestrator 712 to manage these mapped or allocated computing resources.

[0139] In at least one embodiment, software 752 included in software layer 730 may include software used by at least portions of node C.R.s 716(1)-716(N), grouped computing resources 714, and/or distributed file system 738 of framework layer 720. One or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

[0140] In at least one embodiment, application(s) 742 included in application layer 740 may include one or more types of applications used by at least portions of node C.R.s 716(1)-716(N), grouped computing resources 714, and/or distributed file system 738 of framework layer 720. In at least one or more types of applications may include, without

limitation, CUDA applications, 5G network applications, artificial intelligence application, datacenter applications, and/or variations thereof.

[0141] In at least one embodiment, any of configuration manager 734, resource manager 736, and resource orchestrator 712 may implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions may relieve a datacenter operator of datacenter 700 from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor performing portions of a datacenter.

[0142] FIG. 8 illustrates a client-server network 804 formed by a plurality of network server computers 802 which are interlinked, in accordance with at least one embodiment. In at least one embodiment, each network server computer 802 stores data accessible to other network server computers 802 and to client computers 806 and networks 808 which link into a wide area network 804. In at least one embodiment, configuration of a client-server network 804 may change over time as client computers 806 and one or more networks 808 connect and disconnect from a network 804, and as one or more trunk line server computers 802 are added or removed from a network 804. In at least one embodiment, when a client computer 806 and a network 808 are connected with network server computers 802, client-server network includes such client computer 806 and network 808. In at least one embodiment, the term computer includes any device or machine capable of accepting data, applying prescribed processes to data, and supplying results of processes.

[0143] In at least one embodiment, client-server network 804 stores information which is accessible to network server computers 802, remote networks 808 and client computers 806. In at least one embodiment, network server computers 802 are formed by main frame computers minicomputers, and/or microcomputers having one or more processors each. In at least one embodiment, server computers 802 are linked together by wired and/or wireless transfer media, such as conductive wire, fiber optic cable, and/or microwave transmission media, satellite transmission media or other conductive, optic or electromagnetic wave transmission media. In at least one embodiment, client computers 806 access a network server computer 802 by a similar wired or a wireless transfer medium. In at least one embodiment, a client computer 806 may link into a client-server network 804 using a modem and a standard telephone communication network. In at least one embodiment, alternative carrier systems such as cable and satellite communication systems also may be used to link into client-server network 804. In at least one embodiment, other private or time-shared carrier systems may be used. In at least one embodiment, network 804 is a global information network, such as the Internet. In at least one embodiment, network is a private intranet using similar protocols as the Internet, but with added security measures and restricted access controls. In at least one embodiment, network 804 is a private, or semi-private network using proprietary communication protocols.

[0144] In at least one embodiment, client computer 806 is any end user computer, and may also be a mainframe computer, mini-computer or microcomputer having one or more microprocessors. In at least one embodiment, server computer 802 may at times function as a client computer accessing another server computer 802. In at least one

embodiment, remote network 808 may be a local area network, a network added into a wide area network through an independent service provider (ISP) for the Internet, or another group of computers interconnected by wired or wireless transfer media having a configuration which is either fixed or changing over time. In at least one embodiment, client computers 806 may link into and access a network 804 independently or through a remote network 808.

[0145] FIG. 9 illustrates a computer network 908 connecting one or more computing machines, in accordance with at least one embodiment. In at least one embodiment, network 908 may be any type of electronically connected group of computers including, for instance, the following networks: Internet, Intranet, Local Area Networks (LAN), Wide Area Networks (WAN) or an interconnected combination of these network types. In at least one embodiment, connectivity within a network 908 may be a remote modem, Ethernet (IEEE 802.3), Token Ring (IEEE 802.5), Fiber Distributed Datalink Interface (FDDI), Asynchronous Transfer Mode (ATM), or any other communication protocol. In at least one embodiment, computing devices linked to a network may be desktop, server, portable, handheld, set-top box, personal digital assistant (PDA), a terminal, or any other desired type or configuration. In at least one embodiment, depending on their functionality, network connected devices may vary widely in processing power, internal memory, and other performance aspects.

[0146] In at least one embodiment, communications within a network and to or from computing devices connected to a network may be either wired or wireless. In at least one embodiment, network 908 may include, at least in part, the world-wide public Internet which generally connects a plurality of users in accordance with a client-server model in accordance with a transmission control protocol/internet protocol (TCP/IP) specification. In at least one embodiment, client-server network is a dominant model for communicating between two computers. In at least one embodiment, a client computer (“client”) issues one or more commands to a server computer (“server”). In at least one embodiment, server fulfills client commands by accessing available network resources and returning information to a client pursuant to client commands. In at least one embodiment, client computer systems and network resources resident on network servers are assigned a network address for identification during communications between elements of a network. In at least one embodiment, communications from other network connected systems to servers will include a network address of a relevant server/network resource as part of communication so that an appropriate destination of a data/request is identified as a recipient. In at least one embodiment, when a network 908 comprises the global Internet, a network address is an IP address in a TCP/IP format which may, at least in part, route data to an e-mail account, a website, or other Internet tool resident on a server. In at least one embodiment, information and services which are resident on network servers may be available to a web browser of a client computer through a domain name (e.g. www.site.com) which maps to an IP address of a network server.

[0147] In at least one embodiment, a plurality of clients 902, 904, and 906 are connected to a network 908 via respective communication links. In at least one embodiment, each of these clients may access a network 908 via any

desired form of communication, such as via a dial-up modem connection, cable link, a digital subscriber line (DSL), wireless or satellite link, or any other form of communication. In at least one embodiment, each client may communicate using any machine that is compatible with a network **908**, such as a personal computer (PC), work station, dedicated terminal, personal data assistant (PDA), or other similar equipment. In at least one embodiment, clients **902**, **904**, and **906** may or may not be located in a same geographical area.

[0148] In at least one embodiment, a plurality of servers **910**, **912**, and **914** are connected to a network **918** to serve clients that are in communication with a network **918**. In at least one embodiment, each server is typically a powerful computer or device that manages network resources and responds to client commands. In at least one embodiment, servers include computer readable data storage media such as hard disk drives and RAM memory that store program instructions and data. In at least one embodiment, servers **910**, **912**, **914** run application programs that respond to client commands. In at least one embodiment, server **910** may run a web server application for responding to client requests for HTML pages and may also run a mail server application for receiving and routing electronic mail. In at least one embodiment, other application programs, such as an FTP server or a media server for streaming audio/video data to clients may also be running on a server **910**. In at least one embodiment, different servers may be dedicated to performing different tasks. In at least one embodiment, server **910** may be a dedicated web server that manages resources relating to web sites for various users, whereas a server **912** may be dedicated to provide electronic mail (email) management. In at least one embodiment, other servers may be dedicated for media (audio, video, etc.), file transfer protocol (FTP), or a combination of any two or more services that are typically available or provided over a network. In at least one embodiment, each server may be in a location that is the same as or different from that of other servers. In at least one embodiment, there may be multiple servers that perform mirrored tasks for users, thereby relieving congestion or minimizing traffic directed to and from a single server. In at least one embodiment, servers **910**, **912**, **914** are under control of a web hosting provider in a business of maintaining and delivering third party content over a network **918**.

[0149] In at least one embodiment, web hosting providers deliver services to two different types of clients. In at least one embodiment, one type, which may be referred to as a browser, requests content from servers **910**, **912**, **914** such as web pages, email messages, video clips, etc. In at least one embodiment, a second type, which may be referred to as a user, hires a web hosting provider to maintain a network resource such as a web site, and to make it available to browsers. In at least one embodiment, users contract with a web hosting provider to make memory space, processor capacity, and communication bandwidth available for their desired network resource in accordance with an amount of server resources a user desires to utilize.

[0150] In at least one embodiment, in order for a web hosting provider to provide services for both of these clients, application programs which manage a network resources hosted by servers must be properly configured. In at least one embodiment, program configuration process involves defining a set of parameters which control, at least in part,

an application program's response to browser requests and which also define, at least in part, a server resources available to a particular user.

[0151] In one embodiment, an intranet server **916** is in communication with a network **908** via a communication link. In at least one embodiment, intranet server **916** is in communication with a server manager **918**. In at least one embodiment, server manager **918** comprises a database of an application program configuration parameters which are being utilized in servers **910**, **912**, **914**. In at least one embodiment, users modify a database **920** via an intranet **916**, and a server manager **918** interacts with servers **910**, **912**, **914** to modify application program parameters so that they match a content of a database. In at least one embodiment, a user logs onto an intranet server **916** by connecting to an intranet **916** via computer **902** and entering authentication information, such as a username and password.

[0152] In at least one embodiment, when a user wishes to sign up for new service or modify an existing service, an intranet server **916** authenticates a user and provides a user with an interactive screen display/control panel that allows a user to access configuration parameters for a particular application program. In at least one embodiment, a user is presented with a number of modifiable text boxes that describe aspects of a configuration of a user's web site or other network resource. In at least one embodiment, if a user desires to increase memory space reserved on a server for its web site, a user is provided with a field in which a user specifies a desired memory space. In at least one embodiment, in response to receiving this information, an intranet server **916** updates a database **920**. In at least one embodiment, server manager **918** forwards this information to an appropriate server, and a new parameter is used during application program operation. In at least one embodiment, an intranet server **916** is configured to provide users with access to configuration parameters of hosted network resources (e.g., web pages, email, FTP sites, media sites, etc.), for which a user has contracted with a web hosting service provider.

[0153] FIG. 10A illustrates a networked computer system **1000A**, in accordance with at least one embodiment. In at least one embodiment, networked computer system **1000A** comprises a plurality of nodes or personal computers ("PCs") **1002**, **1018**, **1020**. In at least one embodiment, personal computer or node **1002** comprises a processor **1014**, memory **1016**, video camera **1004**, microphone **1006**, mouse **1008**, speakers **1010**, and monitor **1012**. In at least one embodiment, PCs **1002**, **1018**, **1020** may each run one or more desktop servers of an internal network within a given company, for instance, or may be servers of a general network not limited to a specific environment. In at least one embodiment, there is one server per PC node of a network, so that each PC node of a network represents a particular network server, having a particular network URL address. In at least one embodiment, each server defaults to a default web page for that server's user, which may itself contain embedded URLs pointing to further subpages of that user on that server, or to other servers or pages on other servers on a network.

[0154] In at least one embodiment, nodes **1002**, **1018**, **1020** and other nodes of a network are interconnected via medium **1022**. In at least one embodiment, medium **1022** may be, a communication channel such as an Integrated Services Digital Network ("ISDN"). In at least one embodi-

ment, various nodes of a networked computer system may be connected through a variety of communication media, including local area networks (“LANs”), plain-old telephone lines (“POTS”), sometimes referred to as public switched telephone networks (“PSTN”), and/or variations thereof. In at least one embodiment, various nodes of a network may also constitute computer system users inter-connected via a network such as the Internet. In at least one embodiment, each server on a network (running from a particular node of a network at a given instance) has a unique address or identification within a network, which may be specifiable in terms of an URL.

[0155] In at least one embodiment, a plurality of multi-point conferencing units (“MCUs”) may thus be utilized to transmit data to and from various nodes or “endpoints” of a conferencing system. In at least one embodiment, nodes and/or MCUs may be interconnected via an ISDN link or through a local area network (“LAN”), in addition to various other communications media such as nodes connected through the Internet. In at least one embodiment, nodes of a conferencing system may, in general, be connected directly to a communications medium such as a LAN or through an MCU, and that a conferencing system may comprise other nodes or elements such as routers, servers, and/or variations thereof.

[0156] In at least one embodiment, processor 1014 is a general-purpose programmable processor. In at least one embodiment, processors of nodes of networked computer system 1000A may also be special-purpose video processors. In at least one embodiment, various peripherals and components of a node such as those of node 1002 may vary from those of other nodes. In at least one embodiment, node 1018 and node 1020 may be configured identically to or differently than node 1002. In at least one embodiment, a node may be implemented on any suitable computer system in addition to PC systems.

[0157] FIG. 10B illustrates a networked computer system 1000B, in accordance with at least one embodiment. In at least one embodiment, system 1000B illustrates a network such as LAN 1024, which may be used to interconnect a variety of nodes that may communicate with each other. In at least one embodiment, attached to LAN 1024 are a plurality of nodes such as PC nodes 1026, 1028, 1030. In at least one embodiment, a node may also be connected to the LAN via a network server or other means. In at least one embodiment, system 1000B comprises other types of nodes or elements, for at least one embodiment including routers, servers, and nodes.

[0158] FIG. 10C illustrates a networked computer system 1000C, in accordance with at least one embodiment. In at least one embodiment, system 1000C illustrates a WWW system having communications across a backbone communications network such as Internet 1032, which may be used to interconnect a variety of nodes of a network. In at least one embodiment, WWW is a set of protocols operating on top of the Internet, and allows a graphical interface system to operate thereon for accessing information through the Internet. In at least one embodiment, attached to Internet 1032 in WWW are a plurality of nodes such as PCs 1040, 1042, 1044. In at least one embodiment, a node is interfaced to other nodes of WWW through a WWW HTTP server such as servers 1034, 1036. In at least one embodiment, PC 1044 may be a PC forming a node of network 1032 and itself

running its server 1036, although PC 1044 and server 1036 are illustrated separately in FIG. 10C for illustrative purposes.

[0159] In at least one embodiment, WWW is a distributed type of application, characterized by WWW HTTP, WWW’s protocol, which runs on top of the Internet’s transmission control protocol/Internet protocol (“TCP/IP”). In at least one embodiment, WWW may thus be characterized by a set of protocols (i.e., HTTP) running on the Internet as its “backbone.”

[0160] In at least one embodiment, a web browser is an application running on a node of a network that, in WWW-compatible type network systems, allows users of a particular server or node to view such information and thus allows a user to search graphical and text-based files that are linked together using hypertext links that are embedded in documents or files available from servers on a network that understand HTTP. In at least one embodiment, when a given web page of a first server associated with a first node is retrieved by a user using another server on a network such as the Internet, a document retrieved may have various hypertext links embedded therein and a local copy of a page is created local to a retrieving user. In at least one embodiment, when a user clicks on a hypertext link, locally-stored information related to a selected hypertext link is typically sufficient to allow a user’s machine to open a connection across the Internet to a server indicated by a hypertext link.

[0161] In at least one embodiment, more than one user may be coupled to each HTTP server, through a LAN such as LAN 1038 as illustrated with respect to WWW HTTP server 1034. In at least one embodiment, system 1000C may also comprise other types of nodes or elements. In at least one embodiment, a WWW HTTP server is an application running on a machine, such as a PC. In at least one embodiment, each user may be considered to have a unique “server,” as illustrated with respect to PC 1044. In at least one embodiment, a server may be considered to be a server such as WWW HTTP server 1034, which provides access to a network for a LAN or plurality of nodes or plurality of LANs. In at least one embodiment, there are a plurality of users, each having a desktop PC or node of a network, each desktop PC potentially establishing a server for a user thereof. In at least one embodiment, each server is associated with a particular network address or URL, which, when accessed, provides a default web page for that user. In at least one embodiment, a web page may contain further links (embedded URLs) pointing to further subpages of that user on that server, or to other servers on a network or to pages on other servers on a network.

Cloud Computing and Services

[0162] The following figures set forth, without limitation, exemplary cloud-based systems that can be used to implement at least one embodiment.

[0163] In at least one embodiment, cloud computing is a style of computing in which dynamically scalable and often virtualized resources are provided as a service over the Internet. In at least one embodiment, users need not have knowledge of, expertise in, or control over technology infrastructure, which can be referred to as “in the cloud,” that supports them. In at least one embodiment, cloud computing incorporates infrastructure as a service, platform as a service, software as a service, and other variations that have a common theme of reliance on the Internet for

satisfying computing needs of users. In at least one embodiment, a typical cloud deployment, such as in a private cloud (e.g., enterprise network), or a datacenter (DC) in a public cloud (e.g., Internet) can consist of thousands of servers (or alternatively, VMs), hundreds of Ethernet, Fiber Channel or Fiber Channel over Ethernet (FCoE) ports, switching and storage infrastructure, etc. In at least one embodiment, cloud can also consist of network services infrastructure like IPsec VPN hubs, firewalls, load balancers, wide area network (WAN) optimizers etc. In at least one embodiment, remote subscribers can access cloud applications and services securely by connecting via a VPN tunnel, such as an IPsec VPN tunnel.

[0164] In at least one embodiment, cloud computing is a model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.

[0165] In at least one embodiment, cloud computing is characterized by on-demand self-service, in which a consumer can unilaterally provision computing capabilities, such as server time and network storage, as needed automatically without requiring human inter-action with each service's provider. In at least one embodiment, cloud computing is characterized by broad network access, in which capabilities are available over a network and accessed through standard mechanisms that promote use by heterogeneous thin or thick client platforms (e.g., mobile phones, laptops, and PDAs). In at least one embodiment, cloud computing is characterized by resource pooling, in which a provider's computing resources are pooled to serve multiple consumers using a multi-tenant model, with different physical and virtual resources dynamically as-signed and reasigned according to consumer demand. In at least one embodiment, there is a sense of location independence in that a customer generally has no control or knowledge over an exact location of provided resources, but may be able to specify location at a higher level of abstraction (e.g., country, state, or datacenter).

[0166] In at least one embodiment, resources include storage, processing, memory, network bandwidth, and virtual machines. In at least one embodiment, cloud computing is characterized by rapid elasticity, in which capabilities can be rapidly and elastically provisioned, in some cases automatically, to quickly scale out and rapidly released to quickly scale in. In at least one embodiment, to a consumer, capabilities available for provisioning often appear to be unlimited and can be purchased in any quantity at any time. In at least one embodiment, cloud computing is characterized by measured service, in which cloud systems automatically control and optimize resource use by leveraging a metering capability at some level of abstraction appropriate to a type of service (e.g., storage, processing, bandwidth, and active user accounts). In at least one embodiment, resource usage can be monitored, controlled, and reported providing transparency for both a provider and consumer of a utilized service.

[0167] In at least one embodiment, cloud computing may be associated with various services. In at least one embodiment, cloud Software as a Service (SaaS) may refer to as service in which a capability provided to a consumer is to use a provider's applications running on a cloud infrastructure. In at least one embodiment, applications are accessible

from various client devices through a thin client interface such as a web browser (e.g., web-based email). In at least one embodiment, consumer does not manage or control underlying cloud infrastructure including network, servers, operating systems, storage, or even individual application capabilities, with a possible exception of limited user-specific application configuration settings.

[0168] In at least one embodiment, cloud Platform as a Service (PaaS) may refer to a service in which a capability provided to a consumer is to deploy onto cloud infrastructure consumer-created or acquired applications created using programming languages and tools supported by a provider. In at least one embodiment, consumer does not manage or control underlying cloud infrastructure including networks, servers, operating systems, or storage, but has control over deployed applications and possibly application hosting environment configurations.

[0169] In at least one embodiment, cloud Infrastructure as a Service (IaaS) may refer to a service in which a capability provided to a consumer is to provision processing, storage, networks, and other fundamental computing resources where a consumer is able to deploy and run arbitrary software, which can include operating systems and applications. In at least one embodiment, consumer does not manage or control underlying cloud infrastructure, but has control over operating systems, storage, deployed applications, and possibly limited control of select networking components (e.g., host firewalls).

[0170] In at least one embodiment, cloud computing may be deployed in various ways. In at least one embodiment, a private cloud may refer to a cloud infrastructure that is operated solely for an organization. In at least one embodiment, a private cloud may be managed by an organization or a third party and may exist on-premises or off-premises. In at least one embodiment, a community cloud may refer to a cloud infrastructure that is shared by several organizations and supports a specific community that has shared concerns (e.g., mission, security requirements, policy, and compliance considerations). In at least one embodiment, a community cloud may be managed by organizations or a third party and may exist on-premises or off-premises. In at least one embodiment, a public cloud may refer to a cloud infrastructure that is made available to a general public or a large industry group and is owned by an organization providing cloud services. In at least one embodiment, a hybrid cloud may refer to a cloud infrastructure is a composition of two or more clouds (private, community, or public) that remain unique entities, but are bound together by standardized or proprietary technology that enables data and application portability (e.g., cloud bursting for load-balancing between clouds). In at least one embodiment, a cloud computing environment is service oriented with a focus on statelessness, low coupling, modularity, and semantic interoperability.

[0171] FIG. 11 illustrates one or more components of a system environment 1100 in which services may be offered as third party network services, in accordance with at least one embodiment. In at least one embodiment, a third party network may be referred to as a cloud, cloud network, cloud computing network, and/or variations thereof. In at least one embodiment, system environment 1100 includes one or more client computing devices 1104, 1106, and 1108 that may be used by users to interact with a third party network infrastructure system 1102 that provides third party network

services, which may be referred to as cloud computing services. In at least one embodiment, third party network infrastructure system 1102 may comprise one or more computers and/or servers.

[0172] It should be appreciated that third party network infrastructure system 1102 depicted in FIG. 11 may have other components than those depicted. Further, FIG. 11 depicts an embodiment of a third party network infrastructure system. In at least one embodiment, third party network infrastructure system 1102 may have more or fewer components than depicted in FIG. 11, may combine two or more components, or may have a different configuration or arrangement of components.

[0173] In at least one embodiment, client computing devices 1104, 1106, and 1108 may be configured to operate a client application such as a web browser, a proprietary client application, or some other application, which may be used by a user of a client computing device to interact with third party network infrastructure system 1102 to use services provided by third party network infrastructure system 1102. Although exemplary system environment 1100 is shown with three client computing devices, any number of client computing devices may be supported. In at least one embodiment, other devices such as devices with sensors, etc. may interact with third party network infrastructure system 1102. In at least one embodiment, network(s) 1110 may facilitate communications and exchange of data between client computing devices 1104, 1106, and 1108 and third party network infrastructure system 1102.

[0174] In at least one embodiment, services provided by third party network infrastructure system 1102 may include a host of services that are made available to users of a third party network infrastructure system on demand. In at least one embodiment, various services may also be offered including without limitation online data storage and backup solutions, Web-based e-mail services, hosted office suites and document collaboration services, database management and processing, managed technical support services, and/or variations thereof. In at least one embodiment, services provided by a third party network infrastructure system can dynamically scale to meet needs of its users.

[0175] In at least one embodiment, a specific instantiation of a service provided by third party network infrastructure system 1102 may be referred to as a “service instance.” In at least one embodiment, in general, any service made available to a user via a communication network, such as the Internet, from a third party network service provider’s system is referred to as a “third party network service.” In at least one embodiment, in a public third party network environment, servers and systems that make up a third party network service provider’s system are different from a customer’s own on-premises servers and systems. In at least one embodiment, a third party network service provider’s system may host an application, and a user may, via a communication network such as the Internet, on demand, order and use an application.

[0176] In at least one embodiment, a service in a computer network third party network infrastructure may include protected computer network access to storage, a hosted database, a hosted web server, a software application, or other service provided by a third party network vendor to a user. In at least one embodiment, a service can include password-protected access to remote storage on a third party network through the Internet. In at least one embodiment, a

service can include a web service-based hosted relational database and a script-language middleware engine for private use by a networked developer. In at least one embodiment, a service can include access to an email software application hosted on a third party network vendor’s web site.

[0177] In at least one embodiment, third party network infrastructure system 1102 may include a suite of applications, middleware, and database service offerings that are delivered to a customer in a self-service, subscription-based, elastically scalable, reliable, highly available, and secure manner. In at least one embodiment, third party network infrastructure system 1102 may also provide “big data” related computation and analysis services. In at least one embodiment, term “big data” is generally used to refer to extremely large data sets that can be stored and manipulated by analysts and researchers to visualize large amounts of data, detect trends, and/or otherwise interact with data. In at least one embodiment, big data and related applications can be hosted and/or manipulated by an infrastructure system on many levels and at different scales. In at least one embodiment, tens, hundreds, or thousands of processors linked in parallel can act upon such data in order to present it or simulate external forces on data or what it represents. In at least one embodiment, these data sets can involve structured data, such as that organized in a database or otherwise according to a structured model, and/or unstructured data (e.g., emails, images, data blobs (binary large objects), web pages, complex event processing). In at least one embodiment, by leveraging an ability of an embodiment to relatively quickly focus more (or fewer) computing resources upon an objective, a third party network infrastructure system may be better available to carry out tasks on large data sets based on demand from a business, government agency, research organization, private individual, group of like-minded individuals or organizations, or other entity.

[0178] In at least one embodiment, third party network infrastructure system 1102 may be adapted to automatically provision, manage and track a customer’s subscription to services offered by third party network infrastructure system 1102. In at least one embodiment, third party network infrastructure system 1102 may provide third party network services via different deployment models. In at least one embodiment, services may be provided under a public third party network model in which third party network infrastructure system 1102 is owned by an organization selling third party network services and services are made available to a general public or different industry enterprises. In at least one embodiment, services may be provided under a private third party network model in which third party network infrastructure system 1102 is operated solely for a single organization and may provide services for one or more entities within an organization. In at least one embodiment, third party network services may also be provided under a community third party network model in which third party network infrastructure system 1102 and services provided by third party network infrastructure system 1102 are shared by several organizations in a related community. In at least one embodiment, third party network services may also be provided under a hybrid third party network model, which is a combination of two or more different models.

[0179] In at least one embodiment, services provided by third party network infrastructure system 1102 may include one or more services provided under Software as a Service

(SaaS) category, Platform as a Service (PaaS) category, Infrastructure as a Service (IaaS) category, or other categories of services including hybrid services. In at least one embodiment, a customer, via a subscription order, may order one or more services provided by third party network infrastructure system **1102**. In at least one embodiment, third party network infrastructure system **1102** then performs processing to provide services in a customer's subscription order.

[0180] In at least one embodiment, services provided by third party network infrastructure system **1102** may include, without limitation, application services, platform services and infrastructure services. In at least one embodiment, application services may be provided by a third party network infrastructure system via a SaaS platform. In at least one embodiment, SaaS platform may be configured to provide third party network services that fall under a SaaS category. In at least one embodiment, SaaS platform may provide capabilities to build and deliver a suite of on-demand applications on an integrated development and deployment platform. In at least one embodiment, SaaS platform may manage and control underlying software and infrastructure for providing SaaS services. In at least one embodiment, by utilizing services provided by a SaaS platform, customers can utilize applications executing on a third party network infrastructure system. In at least one embodiment, customers can acquire an application services without a need for customers to purchase separate licenses and support. In at least one embodiment, various different SaaS services may be provided. In at least one embodiment, this may include, without limitation, services that provide solutions for sales performance management, enterprise integration, and business flexibility for large organizations.

[0181] In at least one embodiment, platform services may be provided by third party network infrastructure system **1102** via a PaaS platform. In at least one embodiment, PaaS platform may be configured to provide third party network services that fall under a PaaS category. In at least one embodiment, platform services may include without limitation services that enable organizations to consolidate existing applications on a shared, common architecture, as well as an ability to build new applications that leverage shared services provided by a platform. In at least one embodiment, PaaS platform may manage and control underlying software and infrastructure for providing PaaS services. In at least one embodiment, customers can acquire PaaS services provided by third party network infrastructure system **1102** without a need for customers to purchase separate licenses and support.

[0182] In at least one embodiment, by utilizing services provided by a PaaS platform, customers can employ programming languages and tools supported by a third party network infrastructure system and also control deployed services. In at least one embodiment, platform services provided by a third party network infrastructure system may include database third party network services, middleware third party network services and third party network services. In at least one embodiment, database third party network services may support shared service deployment models that enable organizations to pool database resources and offer customers a Database as a Service in a form of a database third party network. In at least one embodiment, middleware third party network services may provide a platform for customers to develop and deploy various busi-

ness applications, and third party network services may provide a platform for customers to deploy applications, in a third party network infrastructure system.

[0183] In at least one embodiment, various different infrastructure services may be provided by an IaaS platform in a third party network infrastructure system. In at least one embodiment, infrastructure services facilitate management and control of underlying computing resources, such as storage, networks, and other fundamental computing resources for customers utilizing services provided by a SaaS platform and a PaaS platform.

[0184] In at least one embodiment, third party network infrastructure system **1102** may also include infrastructure resources **1130** for providing resources used to provide various services to customers of a third party network infrastructure system. In at least one embodiment, infrastructure resources **1130** may include pre-integrated and optimized combinations of hardware, such as servers, storage, and networking resources to execute services provided by a PaaS platform and a SaaS platform, and other resources.

[0185] In at least one embodiment, resources in third party network infrastructure system **1102** may be shared by multiple users and dynamically re-allocated per demand. In at least one embodiment, resources may be allocated to users in different time zones. In at least one embodiment, third party network infrastructure system **1102** may enable a first set of users in a first time zone to utilize resources of a third party network infrastructure system for a specified number of hours and then enable a re-allocation of same resources to another set of users located in a different time zone, thereby maximizing utilization of resources.

[0186] In at least one embodiment, a number of internal shared services **1132** may be provided that are shared by different components or modules of third party network infrastructure system **1102** to enable provision of services by third party network infrastructure system **1102**. In at least one embodiment, these internal shared services may include, without limitation, a security and identity service, an integration service, an enterprise repository service, an enterprise manager service, a virus scanning and white list service, a high availability, backup and recovery service, service for enabling third party network support, an email service, a notification service, a file transfer service, and/or variations thereof.

[0187] In at least one embodiment, third party network infrastructure system **1102** may provide comprehensive management of third party network services (e.g., SaaS, PaaS, and IaaS services) in a third party network infrastructure system. In at least one embodiment, third party network management functionality may include capabilities for provisioning, managing and tracking a customer's subscription received by third party network infrastructure system **1102**, and/or variations thereof.

[0188] In at least one embodiment, as depicted in FIG. **11**, third party network management functionality may be provided by one or more modules, such as an order management module **1120**, an order orchestration module **1122**, an order provisioning module **1124**, an order management and monitoring module **1126**, and an identity management module **1128**. In at least one embodiment, these modules may include or be provided using one or more computers and/or servers, which may be general purpose computers, specialized server computers, server farms, server clusters, or any other appropriate arrangement and/or combination.

[0189] In at least one embodiment, at step 1134, a customer using a client device, such as client computing devices 1104, 1106 or 1108, may interact with third party network infrastructure system 1102 by requesting one or more services provided by third party network infrastructure system 1102 and placing an order for a subscription for one or more services offered by third party network infrastructure system 1102. In at least one embodiment, a customer may access a third party network User Interface (UI) such as third party network UI 1112, third party network UI 1114 and/or third party network UI 1116 and place a subscription order via these UIs. In at least one embodiment, order information received by third party network infrastructure system 1102 in response to a customer placing an order may include information identifying a customer and one or more services offered by a third party network infrastructure system 1102 that a customer intends to subscribe to.

[0190] In at least one embodiment, at step 1136, an order information received from a customer may be stored in an order database 1118. In at least one embodiment, if this is a new order, a new record may be created for an order. In at least one embodiment, order database 1118 can be one of several databases operated by third party network infrastructure system 1118 and operated in conjunction with other system elements.

[0191] In at least one embodiment, at step 1138, an order information may be forwarded to an order management module 1120 that may be configured to perform billing and accounting functions related to an order, such as verifying an order, and upon verification, booking an order.

[0192] In at least one embodiment, at step 1140, information regarding an order may be communicated to an order orchestration module 1122 that is configured to orchestrate provisioning of services and resources for an order placed by a customer. In at least one embodiment, order orchestration module 1122 may use services of order provisioning module 1124 for provisioning. In at least one embodiment, order orchestration module 1122 enables management of business processes associated with each order and applies business logic to determine whether an order should proceed to provisioning.

[0193] In at least one embodiment, at step 1142, upon receiving an order for a new subscription, order orchestration module 1122 sends a request to order provisioning module 1124 to allocate resources and configure resources needed to fulfill a subscription order. In at least one embodiment, order provisioning module 1124 enables an allocation of resources for services ordered by a customer. In at least one embodiment, order provisioning module 1124 provides a level of abstraction between third party network services provided by third party network infrastructure system 1100 and a physical implementation layer that is used to provision resources for providing requested services. In at least one embodiment, this enables order orchestration module 1122 to be isolated from implementation details, such as whether or not services and resources are actually provisioned in real-time or pre-provisioned and only allocated/assigned upon request.

[0194] In at least one embodiment, at step 1144, once services and resources are provisioned, a notification may be sent to subscribing customers indicating that a requested service is now ready for use. In at least one embodiment, information (e.g. a link) may be sent to a customer that enables a customer to start using requested services.

[0195] In at least one embodiment, at step 1146, a customer's subscription order may be managed and tracked by an order management and monitoring module 1126. In at least one embodiment, order management and monitoring module 1126 may be configured to collect usage statistics regarding a customer use of subscribed services. In at least one embodiment, statistics may be collected for an amount of storage used, an amount data transferred, a number of users, and an amount of system up time and system down time, and/or variations thereof.

[0196] In at least one embodiment, third party network infrastructure system 1100 may include an identity management module 1128 that is configured to provide identity services, such as access management and authorization services in third party network infrastructure system 1100. In at least one embodiment, identity management module 1128 may control information about customers who wish to utilize services provided by third party network infrastructure system 1102. In at least one embodiment, such information can include information that authenticates identities of such customers and information that describes which actions those customers are authorized to perform relative to various system resources (e.g., files, directories, applications, communication ports, memory segments, etc.). In at least one embodiment, identity management module 1128 may also include management of descriptive information about each customer and about how and by whom that descriptive information can be accessed and modified.

[0197] FIG. 12 illustrates a cloud computing environment 1202, in accordance with at least one embodiment. In at least one embodiment, cloud computing environment 1202 comprises one or more computer system/servers 1204 with which computing devices such as, personal digital assistant (PDA) or cellular telephone 1206A, desktop computer 1206B, laptop computer 1206C, and/or automobile computer system 1206N communicate. In at least one embodiment, this allows for infrastructure, platforms and/or software to be offered as services from cloud computing environment 1202, so as to not require each client to separately maintain such resources. It is understood that types of computing devices 1206A-N shown in FIG. 12 are intended to be illustrative only and that cloud computing environment 1202 can communicate with any type of computerized device over any type of network and/or network/addressable connection (e.g., using a web browser).

[0198] In at least one embodiment, a computer system/server 1204, which can be denoted as a cloud computing node, is operational with numerous other general purpose or special purpose computing system environments or configurations. In at least one embodiment, computing systems, environments, and/or configurations that may be suitable for use with computer system/server 1204 include, but are not limited to, personal computer systems, server computer systems, thin clients, thick clients, hand-held or laptop devices, multiprocessor systems, microprocessor-based systems, set top boxes, programmable consumer electronics, network PCs, minicomputer systems, mainframe computer systems, and distributed cloud computing environments that include any of the above systems or devices, and/or variations thereof.

[0199] In at least one embodiment, computer system/server 1204 may be described in a general context of computer system-executable instructions, such as program modules, being executed by a computer system. In at least

one embodiment, program modules include routines, programs, objects, components, logic, data structures, and so on, that perform particular tasks or implement particular abstract data types. In at least one embodiment, exemplary computer system/server **1204** may be practiced in distributed loud computing environments where tasks are performed by remote processing devices that are linked through a communications network. In at least one embodiment, in a distributed cloud computing environment, program modules may be located in both local and remote computer system storage media including memory storage devices.

[0200] FIG. **13** illustrates a set of functional abstraction layers provided by cloud computing environment **1202** (FIG. **12**), in accordance with at least one embodiment. It should be understood in advance that components, layers, and functions shown in FIG. **13** are intended to be illustrative only, and components, layers, and functions may vary.

[0201] In at least one embodiment, hardware and software layer **1302** includes hardware and software components. In at least one embodiment, hardware components include mainframes, various RISC (Reduced Instruction Set Computer) architecture based servers, various computing systems, supercomputing systems, storage devices, networks, networking components, and/or variations thereof. In at least one embodiment, software components include network application server software, various application server software, various database software, and/or variations thereof.

[0202] In at least one embodiment, virtualization layer **1304** provides an abstraction layer from which following exemplary virtual entities may be provided: virtual servers, virtual storage, virtual networks, including virtual private networks, virtual applications, virtual clients, and/or variations thereof.

[0203] In at least one embodiment, management layer **1306** provides various functions. In at least one embodiment, resource provisioning provides dynamic procurement of computing resources and other resources that are utilized to perform tasks within a cloud computing environment. In at least one embodiment, metering provides usage tracking as resources are utilized within a cloud computing environment, and billing or invoicing for consumption of these resources. In at least one embodiment, resources may comprise application software licenses. In at least one embodiment, security provides identity verification for users and tasks, as well as protection for data and other resources. In at least one embodiment, user interface provides access to a cloud computing environment for both users and system administrators. In at least one embodiment, service level management provides cloud computing resource allocation and management such that required service levels are met. In at least one embodiment, Service Level Agreement (SLA) management provides pre-arrangement for, and procurement of, cloud computing resources for which a future requirement is anticipated in accordance with an SLA.

[0204] In at least one embodiment, workloads layer **1308** provides functionality for which a cloud computing environment is utilized. In at least one embodiment, workloads and functions which may be provided from this layer include: mapping and navigation, software development and management, educational services, data analytics and processing, transaction processing, and service delivery.

Supercomputing

[0205] The following figures set forth, without limitation, exemplary supercomputer-based systems that can be used to implement at least one embodiment.

[0206] In at least one embodiment, a supercomputer may refer to a hardware system exhibiting substantial parallelism and comprising at least one chip, where chips in a system are interconnected by a network and are placed in hierarchically organized enclosures. In at least one embodiment, a large hardware system filling a machine room, with several racks, each containing several boards/rack modules, each containing several chips, all interconnected by a scalable network, is at least one embodiment of a supercomputer. In at least one embodiment, a single rack of such a large hardware system is at least one other embodiment of a supercomputer. In at least one embodiment, a single chip exhibiting substantial parallelism and containing several hardware components can equally be considered to be a supercomputer, since as feature sizes may decrease, an amount of hardware that can be incorporated in a single chip may also increase.

[0207] FIG. **14** illustrates a supercomputer at a chip level, in accordance with at least one embodiment. In at least one embodiment, inside an FPGA or ASIC chip, main computation is performed within finite state machines (**1404**) called thread units. In at least one embodiment, task and synchronization networks (**1402**) connect finite state machines and are used to dispatch threads and execute operations in correct order. In at least one embodiment, a multi-level partitioned on-chip cache hierarchy (**1408**, **1412**) is accessed using memory networks (**1406**, **1410**). In at least one embodiment, off-chip memory is accessed using memory controllers (**1416**) and an off-chip memory network (**1414**). In at least one embodiment, I/O controller (**1418**) is used for cross-chip communication when a design does not fit in a single logic chip.

[0208] FIG. **15** illustrates a supercomputer at a rack module level, in accordance with at least one embodiment. In at least one embodiment, within a rack module, there are multiple FPGA or ASIC chips (**1502**) that are connected to one or more DRAM units (**1504**) which constitute main accelerator memory. In at least one embodiment, each FPGA/ASIC chip is connected to its neighbor FPGA/ASIC chip using wide busses on a board, with differential high speed signaling (**1506**). In at least one embodiment, each FPGA/ASIC chip is also connected to at least one high-speed serial communication cable.

[0209] FIG. **16** illustrates a supercomputer at a rack level, in accordance with at least one embodiment. FIG. **17** illustrates a supercomputer at a whole system level, in accordance with at least one embodiment. In at least one embodiment, referring to FIG. **16** and FIG. **17**, between rack modules in a rack and across racks throughout an entire system, high-speed serial optical or copper cables (**1602**, **1702**) are used to realize a scalable, possibly incomplete hypercube network. In at least one embodiment, one of FPGA/ASIC chips of an accelerator is connected to a host system through a PCI-Express connection (**1704**). In at least one embodiment, host system comprises a host microprocessor (**1708**) that a software part of an application runs on and a memory consisting of one or more host memory DRAM units (**1706**) that is kept coherent with memory on an accelerator. In at least one embodiment, host system can be a separate module on one of racks, or can be integrated with one of a supercomputer's modules. In at least one

embodiment, cube-connected cycles topology provide communication links to create a hypercube network for a large supercomputer. In at least one embodiment, a small group of FPGA/ASIC chips on a rack module can act as a single hypercube node, such that a total number of external links of each group is increased, compared to a single chip. In at least one embodiment, a group contains chips A, B, C and D on a rack module with internal wide differential busses connecting A, B, C and D in a torus organization. In at least one embodiment, there are 12 serial communication cables connecting a rack module to an outside world. In at least one embodiment, chip A on a rack module connects to serial communication cables 0, 1, 2. In at least one embodiment, chip B connects to cables 3, 4, 5. In at least one embodiment, chip C connects to 6, 7, 8. In at least one embodiment, chip D connects to 9, 10, 11. In at least one embodiment, an entire group {A, B, C, D} constituting a rack module can form a hypercube node within a supercomputer system, with up to $2^{12}=4096$ rack modules (16384 FPGA/ASIC chips). In at least one embodiment, for chip A to send a message out on link 4 of group {A, B, C, D}, a message has to be routed first to chip B with an on-board differential wide bus connection. In at least one embodiment, a message arriving into a group {A, B, C, D} on link 4 (i.e., arriving at B) destined to chip A, also has to be routed first to a correct destination chip (A) internally within a group {A, B, C, D}. In at least one embodiment, parallel supercomputer systems of other sizes may also be implemented.

Artificial Intelligence

[0210] The following figures set forth, without limitation, exemplary artificial intelligence-based systems that can be used to implement at least one embodiment.

[0211] FIG. 18A illustrates inference and/or training logic **1815** used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1815** are provided below in conjunction with FIGS. 18A and/or 18B.

[0212] In at least one embodiment, inference and/or training logic **1815** may include, without limitation, code and/or data storage **1801** to store forward and/or output weight and/or input/output data, and/or other parameters to configure neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, training logic **1815** may include, or be coupled to code and/or data storage **1801** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which such code corresponds. In at least one embodiment code and/or data storage **1801** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during forward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, any portion of code and/or data storage **1801** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

[0213] In at least one embodiment, any portion of code and/or data storage **1801** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or code and/or data storage **1801** may be cache memory, dynamic randomly addressable memory ("DRAM"), static randomly addressable memory ("SRAM"), non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, a choice of whether code and/or code and/or data storage **1801** is internal or external to a processor, in at least one embodiment, or comprising DRAM, SRAM, flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

[0214] In at least one embodiment, inference and/or training logic **1815** may include, without limitation, a code and/or data storage **1805** to store backward and/or output weight and/or input/output data corresponding to neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, code and/or data storage **1805** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during backward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, training logic **1815** may include, or be coupled to code and/or data storage **1805** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)).

[0215] In at least one embodiment, code, such as graph code, causes loading of weight or other parameter information into processor ALUs based on an architecture of a neural network to which such code corresponds. In at least one embodiment, any portion of code and/or data storage **1805** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. In at least one embodiment, any portion of code and/or data storage **1805** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or data storage **1805** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, a choice of whether code and/or data storage **1805** is internal or external to a processor, in at least one embodiment, or comprising DRAM, SRAM, flash memory or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

[0216] In at least one embodiment, code and/or data storage **1801** and code and/or data storage **1805** may be separate storage structures. In at least one embodiment, code and/or data storage **1801** and code and/or data storage **1805** may be a combined storage structure. In at least one embodiment, code and/or data storage **1801** and code and/or data storage **1805** may be partially combined and partially separate. In at least one embodiment, any portion of code and/or data storage **1801** and code and/or data storage **1805** may be

included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

[0217] In at least one embodiment, inference and/or training logic **1815** may include, without limitation, one or more arithmetic logic unit(s) ("ALU(s)") **1810**, including integer and/or floating point units, to perform logical and/or mathematical operations based, at least in part on, or indicated by, training and/or inference code (e.g., graph code), a result of which may produce activations (e.g., output values from layers or neurons within a neural network) stored in an activation storage **1820** that are functions of input/output and/or weight parameter data stored in code and/or data storage **1801** and/or code and/or data storage **1805**. In at least one embodiment, activations stored in activation storage **1820** are generated according to linear algebraic and/or matrix-based mathematics performed by ALU(s) **1810** in response to performing instructions or other code, wherein weight values stored in code and/or data storage **1805** and/or data storage **1801** are used as operands along with other values, such as bias values, gradient information, momentum values, or other parameters or hyperparameters, any or all of which may be stored in code and/or data storage **1805** or code and/or data storage **1801** or another storage on or off-chip.

[0218] In at least one embodiment, ALU(s) **1810** are included within one or more processors or other hardware logic devices or circuits, whereas in another embodiment, ALU(s) **1810** may be external to a processor or other hardware logic device or circuit that uses them (e.g., a co-processor). In at least one embodiment, ALUs **1810** may be included within a processor's execution units or otherwise within a bank of ALUs accessible by a processor's execution units either within same processor or distributed between different processors of different types (e.g., central processing units, graphics processing units, fixed function units, etc.). In at least one embodiment, code and/or data storage **1801**, code and/or data storage **1805**, and activation storage **1820** may share a processor or other hardware logic device or circuit, whereas in another embodiment, they may be in different processors or other hardware logic devices or circuits, or some combination of same and different processors or other hardware logic devices or circuits. In at least one embodiment, any portion of activation storage **1820** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. Furthermore, inferencing and/or training code may be stored with other code accessible to a processor or other hardware logic or circuit and fetched and/or processed using a processor's fetch, decode, scheduling, execution, retirement and/or other logical circuits.

[0219] In at least one embodiment, activation storage **1820** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, activation storage **1820** may be completely or partially within or external to one or more processors or other logical circuits. In at least one embodiment, a choice of whether activation storage **1820** is internal or external to a processor, in at least one embodiment, or comprising DRAM, SRAM, flash memory or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

[0220] In at least one embodiment, inference and/or training logic **1815** illustrated in FIG. **18A** may be used in conjunction with an application-specific integrated circuit ("ASIC"), such as a TensorFlow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1815** illustrated in FIG. **18A** may be used in conjunction with central processing unit ("CPU") hardware, graphics processing unit ("GPU") hardware or other hardware, such as field programmable gate arrays ("FPGAs").

[0221] FIG. **18B** illustrates inference and/or training logic **1815**, according to at least one embodiment. In at least one embodiment, inference and/or training logic **1815** may include, without limitation, hardware logic in which computational resources are dedicated or otherwise exclusively used in conjunction with weight values or other information corresponding to one or more layers of neurons within a neural network. In at least one embodiment, inference and/or training logic **1815** illustrated in FIG. **18B** may be used in conjunction with an application-specific integrated circuit (ASIC), such as TensorFlow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1815** illustrated in FIG. **18B** may be used in conjunction with central processing unit (CPU) hardware, graphics processing unit (GPU) hardware or other hardware, such as field programmable gate arrays (FPGAs). In at least one embodiment, inference and/or training logic **1815** includes, without limitation, code and/or data storage **1801** and code and/or data storage **1805**, which may be used to store code (e.g., graph code), weight values and/or other information, including bias values, gradient information, momentum values, and/or other parameter or hyperparameter information. In at least one embodiment illustrated in FIG. **18B**, each of code and/or data storage **1801** and code and/or data storage **1805** is associated with a dedicated computational resource, such as computational hardware **1802** and computational hardware **1806**, respectively. In at least one embodiment, each of computational hardware **1802** and computational hardware **1806** comprises one or more ALUs that perform mathematical functions, such as linear algebraic functions, only on information stored in code and/or data storage **1801** and code and/or data storage **1805**, respectively, result of which is stored in activation storage **1820**.

[0222] In at least one embodiment, each of code and/or data storage **1801** and **1805** and corresponding computational hardware **1802** and **1806**, respectively, correspond to different layers of a neural network, such that resulting activation from one storage/computational pair **1801/1802** of code and/or data storage **1801** and computational hardware **1802** is provided as an input to a next storage/computational pair **1805/1806** of code and/or data storage **1805** and computational hardware **1806**, in order to mirror a conceptual organization of a neural network. In at least one embodiment, each of storage/computational pairs **1801/1802** and **1805/1806** may correspond to more than one neural network layer. In at least one embodiment, additional storage/computation pairs (not shown) subsequent to or in

parallel with storage/computation pairs **1801/1802** and **1805/1806** may be included in inference and/or training logic **1815**.

[0223] FIG. 19 illustrates training and deployment of a deep neural network, according to at least one embodiment. In at least one embodiment, untrained neural network **1906** is trained using a training dataset **1902**. In at least one embodiment, training framework **1904** is a PyTorch framework, whereas in other embodiments, training framework **1904** is a TensorFlow, Boost, Caffe, Microsoft Cognitive Toolkit/CNTK, MXNet, Chainer, Keras, Deeplearning4j, or other training framework. In at least one embodiment, training framework **1904** trains an untrained neural network **1906** and enables it to be trained using processing resources described herein to generate a trained neural network **1908**. In at least one embodiment, weights may be chosen randomly or by pre-training using a deep belief network. In at least one embodiment, training may be performed in either a supervised, partially supervised, or unsupervised manner.

[0224] In at least one embodiment, untrained neural network **1906** is trained using supervised learning, wherein training dataset **1902** includes an input paired with a desired output for an input, or where training dataset **1902** includes input having a known output and an output of neural network **1906** is manually graded. In at least one embodiment, untrained neural network **1906** is trained in a supervised manner and processes inputs from training dataset **1902** and compares resulting outputs against a set of expected or desired outputs. In at least one embodiment, errors are then propagated back through untrained neural network **1906**. In at least one embodiment, training framework **1904** adjusts weights that control untrained neural network **1906**. In at least one embodiment, training framework **1904** includes tools to monitor how well untrained neural network **1906** is converging towards a model, such as trained neural network **1908**, suitable to generating correct answers, such as in result **1914**, based on input data such as a new dataset **1912**. In at least one embodiment, training framework **1904** trains untrained neural network **1906** repeatedly while adjust weights to refine an output of untrained neural network **1906** using a loss function and adjustment algorithm, such as stochastic gradient descent. In at least one embodiment, training framework **1904** trains untrained neural network **1906** until untrained neural network **1906** achieves a desired accuracy. In at least one embodiment, trained neural network **1908** can then be deployed to implement any number of machine learning operations.

[0225] In at least one embodiment, untrained neural network **1906** is trained using unsupervised learning, wherein untrained neural network **1906** attempts to train itself using unlabeled data. In at least one embodiment, unsupervised learning training dataset **1902** will include input data without any associated output data or “ground truth” data. In at least one embodiment, untrained neural network **1906** can learn groupings within training dataset **1902** and can determine how individual inputs are related to untrained dataset **1902**. In at least one embodiment, unsupervised training can be used to generate a self-organizing map in trained neural network **1908** capable of performing operations useful in reducing dimensionality of new dataset **1912**. In at least one embodiment, unsupervised training can also be used to perform anomaly detection, which allows identification of

data points in new dataset **1912** that deviate from normal patterns of new dataset **1912**.

[0226] In at least one embodiment, semi-supervised learning may be used, which is a technique in which in training dataset **1902** includes a mix of labeled and unlabeled data. In at least one embodiment, training framework **1904** may be used to perform incremental learning, such as through transferred learning techniques. In at least one embodiment, incremental learning enables trained neural network **1908** to adapt to new dataset **1912** without forgetting knowledge instilled within trained neural network **1408** during initial training.

5G Networks

[0227] The following figures set forth, without limitation, exemplary 5G network-based systems that can be used to implement at least one embodiment.

[0228] FIG. 20 illustrates an architecture of a system **2000** of a network, in accordance with at least one embodiment. In at least one embodiment, system **2000** is shown to include a user equipment (UE) **2002** and a UE **2004**. In at least one embodiment, UEs **2002** and **2004** are illustrated as smartphones (e.g., handheld touchscreen mobile computing devices connectable to one or more cellular networks) but may also comprise any mobile or non-mobile computing device, such as Personal Data Assistants (PDAs), pagers, laptop computers, desktop computers, wireless handsets, or any computing device including a wireless communications interface.

[0229] In at least one embodiment, any of UEs **2002** and **2004** can comprise an Internet of Things (IoT) UE, which can comprise a network access layer designed for low-power IoT applications utilizing short-lived UE connections. In at least one embodiment, an IoT UE can utilize technologies such as machine-to-machine (M2M) or machine-type communications (MTC) for exchanging data with an MTC server or device via a public land mobile network (PLMN), Proximity-Based Service (ProSe) or device-to-device (D2D) communication, sensor networks, or IoT networks. In at least one embodiment, a M2M or MTC exchange of data may be a machine-initiated exchange of data. In at least one embodiment, an IoT network describes interconnecting IoT UEs, which may include uniquely identifiable embedded computing devices (within Internet infrastructure), with short-lived connections. In at least one embodiment, an IoT UEs may execute background applications (e.g., keep alive messages, status updates, etc.) to facilitate connections of an IoT network.

[0230] In at least one embodiment, UEs **2002** and **2004** may be configured to connect, e.g., communicatively couple, with a radio access network (RAN) **2016**. In at least one embodiment, RAN **2016** may be, in at least one embodiment, an Evolved Universal Mobile Telecommunications System (UMTS) Terrestrial Radio Access Network (E-UTRAN), a NextGen RAN (NG RAN), or some other type of RAN. In at least one embodiment, UEs **2002** and **2004** utilize connections **2012** and **2014**, respectively, each of which comprises a physical communications interface or layer. In at least one embodiment, connections **2012** and **2014** are illustrated as an air interface to enable communicative coupling, and can be consistent with cellular communications protocols, such as a Global System for Mobile Communications (GSM) protocol, a code-division multiple access (CDMA) network protocol, a Push-to-Talk (PTT)

protocol, a PTT over Cellular (POC) protocol, a Universal Mobile Telecommunications System (UMTS) protocol, a 3GPP Long Term Evolution (LTE) protocol, a fifth generation (5G) protocol, a New Radio (NR) protocol, and variations thereof.

[0231] In at least one embodiment, UEs **2002** and **2004** may further directly exchange communication data via a ProSe interface **2006**. In at least one embodiment, ProSe interface **2006** may alternatively be referred to as a sidelink interface comprising one or more logical channels, including but not limited to a Physical Sidelink Control Channel (PSCCH), a Physical Sidelink Shared Channel (PSSCH), a Physical Sidelink Discovery Channel (PSDCH), and a Physical Sidelink Broadcast Channel (PSBCH).

[0232] In at least one embodiment, UE **2004** is shown to be configured to access an access point (AP) **2010** via connection **2008**. In at least one embodiment, connection **2008** can comprise a local wireless connection, such as a connection consistent with any IEEE 802.11 protocol, wherein AP **2010** would comprise a wireless fidelity (WiFi®) router. In at least one embodiment, AP **2010** is shown to be connected to an Internet without connecting to a core network of a wireless system.

[0233] In at least one embodiment, RAN **2016** can include one or more access nodes that enable connections **2012** and **2014**. In at least one embodiment, these access nodes (ANs) can be referred to as base stations (BSs), NodeBs, evolved NodeBs (eNBs), next Generation NodeBs (gNB), RAN nodes, and so forth, and can comprise ground stations (e.g., terrestrial access points) or satellite stations providing coverage within a geographic area (e.g., a cell). In at least one embodiment, RAN **2016** may include one or more RAN nodes for providing macrocells, e.g., macro RAN node **2018**, and one or more RAN nodes for providing femtocells or picocells (e.g., cells having smaller coverage areas, smaller user capacity, or higher bandwidth compared to macrocells), e.g., low power (LP) RAN node **2020**.

[0234] In at least one embodiment, any of RAN nodes **2018** and **2020** can terminate an air interface protocol and can be a first point of contact for UEs **2002** and **2004**. In at least one embodiment, any of RAN nodes **2018** and **2020** can fulfill various logical functions for RAN **2016** including, but not limited to, radio network controller (RNC) functions such as radio bearer management, uplink and downlink dynamic radio resource management and data packet scheduling, and mobility management.

[0235] In at least one embodiment, UEs **2002** and **2004** can be configured to communicate using Orthogonal Frequency-Division Multiplexing (OFDM) communication signals with each other or with any of RAN nodes **2018** and **2020** over a multi-carrier communication channel in accordance various communication techniques, such as, but not limited to, an Orthogonal Frequency Division Multiple Access (OFDMA) communication technique (e.g., for downlink communications) or a Single Carrier Frequency Division Multiple Access (SC-FDMA) communication technique (e.g., for uplink and ProSe or sidelink communications), and/or variations thereof. In at least one embodiment, OFDM signals can comprise a plurality of orthogonal sub-carriers.

[0236] In at least one embodiment, a downlink resource grid can be used for downlink transmissions from any of RAN nodes **2018** and **2020** to UEs **2002** and **2004**, while uplink transmissions can utilize similar techniques. In at

least one embodiment, a grid can be a time frequency grid, called a resource grid or time-frequency resource grid, which is a physical resource in a downlink in each slot. In at least one embodiment, such a time frequency plane representation is a common practice for OFDM systems, which makes it intuitive for radio resource allocation. In at least one embodiment, each column and each row of a resource grid corresponds to one OFDM symbol and one OFDM subcarrier, respectively. In at least one embodiment, a duration of a resource grid in a time domain corresponds to one slot in a radio frame. In at least one embodiment, a smallest time-frequency unit in a resource grid is denoted as a resource element. In at least one embodiment, each resource grid comprises a number of resource blocks, which describe a mapping of certain physical channels to resource elements. In at least one embodiment, each resource block comprises a collection of resource elements. In at least one embodiment, in a frequency domain, this may represent a smallest quantity of resources that currently can be allocated. In at least one embodiment, there are several different physical downlink channels that are conveyed using such resource blocks.

[0237] In at least one embodiment, a physical downlink shared channel (PDSCH) may carry user data and higher-layer signaling to UEs **2002** and **2004**. In at least one embodiment, a physical downlink control channel (PDCCH) may carry information about a transport format and resource allocations related to PDSCH channel, among other things. In at least one embodiment, it may also inform UEs **2002** and **2004** about a transport format, resource allocation, and HARQ (Hybrid Automatic Repeat Request) information related to an uplink shared channel. In at least one embodiment, typically, downlink scheduling (assigning control and shared channel resource blocks to UE **2002** within a cell) may be performed at any of RAN nodes **2018** and **2020** based on channel quality information fed back from any of UEs **2002** and **2004**. In at least one embodiment, downlink resource assignment information may be sent on a PDCCH used for (e.g., assigned to) each of UEs **2002** and **2004**.

[0238] In at least one embodiment, a PDCCH may use control channel elements (CCEs) to convey control information. In at least one embodiment, before being mapped to resource elements, PDCCH complex valued symbols may first be organized into quadruplets, which may then be permuted using a sub-block interleaver for rate matching. In at least one embodiment, each PDCCH may be transmitted using one or more of these CCEs, where each CCE may correspond to nine sets of four physical resource elements known as resource element groups (REGs). In at least one embodiment, four Quadrature Phase Shift Keying (QPSK) symbols may be mapped to each REG. In at least one embodiment, PDCCH can be transmitted using one or more CCEs, depending on a size of a downlink control information (DCI) and a channel condition. In at least one embodiment, there can be four or more different PDCCH formats defined in LTE with different numbers of CCEs (e.g., aggregation level, $L=1, 2, 4$, or 8).

[0239] In at least one embodiment, an enhanced physical downlink control channel (EPDCCH) that uses PDSCH resources may be utilized for control information transmission. In at least one embodiment, EPDCCH may be transmitted using one or more enhanced control channel elements (ECCEs). In at least one embodiment, each ECCE may correspond to nine sets of four physical resource elements

known as an enhanced resource element groups (EREGs). In at least one embodiment, an ECCE may have other numbers of EREGs in some situations.

[0240] In at least one embodiment, RAN **2016** is shown to be communicatively coupled to a core network (CN) **2038** via an S1 interface **2022**. In at least one embodiment, CN **2038** may be an evolved packet core (EPC) network, a NextGen Packet Core (NPC) network, or some other type of CN. In at least one embodiment, S1 interface **2022** is split into two parts: S1-U interface **2026**, which carries traffic data between RAN nodes **2018** and **2020** and serving gateway (S-GW) **2030**, and a S1-mobility management entity (MME) interface **2024**, which is a signaling interface between RAN nodes **2018** and **2020** and MMEs **2028**.

[0241] In at least one embodiment, CN **2038** comprises MMES **2028**, S-GW **2030**, Packet Data Network (PDN) Gateway (P-GW) **2034**, and a home subscriber server (HSS) **2032**. In at least one embodiment, MMES **2028** may be similar in function to a control plane of legacy Serving General Packet Radio Service (GPRS) Support Nodes (SGSN). In at least one embodiment, MMES **2028** may manage mobility aspects in access such as gateway selection and tracking area list management. In at least one embodiment, HSS **2032** may comprise a database for network users, including subscription related information to support a network entities' handling of communication sessions. In at least one embodiment, CN **2038** may comprise one or several HSSs **2032**, depending on a number of mobile subscribers, on a capacity of an equipment, on an organization of a network, etc. In at least one embodiment, HSS **2032** can provide support for routing/roaming, authentication, authorization, naming/addressing resolution, location dependencies, etc.

[0242] In at least one embodiment, S-GW **2030** may terminate a S1 interface **2022** towards RAN **2016**, and routes data packets between RAN **2016** and CN **2038**. In at least one embodiment, S-GW **2030** may be a local mobility anchor point for inter-RAN node handovers and also may provide an anchor for inter-3GPP mobility. In at least one embodiment, other responsibilities may include lawful intercept, charging, and some policy enforcement.

[0243] In at least one embodiment, P-GW **2034** may terminate an SGI interface toward a PDN. In at least one embodiment, P-GW **2034** may route data packets between an EPC network **2038** and external networks such as a network including application server **2040** (alternatively referred to as application function (AF)) via an Internet Protocol (IP) interface **2042**. In at least one embodiment, application server **2040** may be an element offering applications that use IP bearer resources with a core network (e.g., UMTS Packet Services (PS) domain, LTE PS data services, etc.). In at least one embodiment, P-GW **2034** is shown to be communicatively coupled to an application server **2040** via an IP communications interface **2042**. In at least one embodiment, application server **2040** can also be configured to support one or more communication services (e.g., Voice-over-Internet Protocol (VoIP) sessions, PTT sessions, group communication sessions, social networking services, etc.) for UEs **2002** and **2004** via CN **2038**.

[0244] In at least one embodiment, P-GW **2034** may further be a node for policy enforcement and charging data collection. In at least one embodiment, policy and Charging Enforcement Function (PCRF) **2036** is a policy and charging control element of CN **2038**. In at least one embodiment, in

a non-roaming scenario, there may be a single PCRF in a Home Public Land Mobile Network (HPLMN) associated with a UE's Internet Protocol Connectivity Access Network (IP-CAN) session. In at least one embodiment, in a roaming scenario with local breakout of traffic, there may be two PCRFs associated with a UE's IP-CAN session: a Home PCRF (H-PCRF) within a HPLMN and a Visited PCRF (V-PCRF) within a Visited Public Land Mobile Network (VPLMN). In at least one embodiment, PCRF **2036** may be communicatively coupled to application server **2040** via P-GW **2034**. In at least one embodiment, application server **2040** may signal PCRF **2036** to indicate a new service flow and select an appropriate Quality of Service (QoS) and charging parameters. In at least one embodiment, PCRF **2036** may provision this rule into a Policy and Charging Enforcement Function (PCEF) (not shown) with an appropriate traffic flow template (TFT) and QoS class of identifier (QCI), which commences a QoS and charging as specified by application server **2040**.

[0245] FIG. **21** illustrates an architecture of a system **2100** of a network in accordance with some embodiments. In at least one embodiment, system **2100** is shown to include a UE **2102**, a 5G access node or RAN node (shown as (R)AN node **2108**), a User Plane Function (shown as UPF **2104**), a Data Network (DN **2106**), which may be, in at least one embodiment, operator services, Internet access or 3rd party services, and a 5G Core Network (5GC) (shown as CN **2110**).

[0246] In at least one embodiment, CN **2110** includes an Authentication Server Function (AUSF **2114**); a Core Access and Mobility Management Function (AMF **2112**); a Session Management Function (SMF **2118**); a Network Exposure Function (NEF **2116**); a Policy Control Function (PCF **2122**); a Network Function (NF) Repository Function (NRF **2120**); a Unified Data Management (UDM **2124**); and an Application Function (AF **2126**). In at least one embodiment, CN **2110** may also include other elements that are not shown, such as a Structured Data Storage network function (SDSF), an Unstructured Data Storage network function (UDSF), and variations thereof.

[0247] In at least one embodiment, UPF **2104** may act as an anchor point for intra-RAT and inter-RAT mobility, an external PDU session point of interconnect to DN **2106**, and a branching point to support multi-homed PDU session. In at least one embodiment, UPF **2104** may also perform packet routing and forwarding, packet inspection, enforce user plane part of policy rules, lawfully intercept packets (UP collection); traffic usage reporting, perform QoS handling for user plane (e.g. packet filtering, gating, UL/DL rate enforcement), perform Uplink Traffic verification (e.g., SDF to QoS flow mapping), transport level packet marking in uplink and downlink, and downlink packet buffering and downlink data notification triggering. In at least one embodiment, UPF **2104** may include an uplink classifier to support routing traffic flows to a data network. In at least one embodiment, DN **2106** may represent various network operator services, Internet access, or third party services.

[0248] In at least one embodiment, AUSF **2114** may store data for authentication of UE **2102** and handle authentication related functionality. In at least one embodiment, AUSF **2114** may facilitate a common authentication framework for various access types.

[0249] In at least one embodiment, AMF **2112** may be responsible for registration management (e.g., for register-

ing UE **2102**, etc.), connection management, reachability management, mobility management, and lawful interception of AMF-related events, and access authentication and authorization. In at least one embodiment, AMF **2112** may provide transport for SM messages for SMF **2118**, and act as a transparent proxy for routing SM messages. In at least one embodiment, AMF **2112** may also provide transport for short message service (SMS) messages between UE **2102** and an SMS function (SMSF) (not shown by FIG. **21**). In at least one embodiment, AMF **2112** may act as Security Anchor Function (SEA), which may include interaction with AUSF **2114** and UE **2102** and receipt of an intermediate key that was established as a result of UE **2102** authentication process. In at least one embodiment, where USIM based authentication is used, AMF **2112** may retrieve security material from AUSF **2114**. In at least one embodiment, AMF **2112** may also include a Security Context Management (SCM) function, which receives a key from SEA that it uses to derive access-network specific keys. In at least one embodiment, furthermore, AMF **2112** may be a termination point of RAN CP interface (N2 reference point), a termination point of NAS (NI) signaling, and perform NAS ciphering and integrity protection.

[0250] In at least one embodiment, AMF **2112** may also support NAS signaling with a UE **2102** over an N3 interworking-function (IWF) interface. In at least one embodiment, N3IWF may be used to provide access to untrusted entities. In at least one embodiment, N3IWF may be a termination point for N2 and N3 interfaces for control plane and user plane, respectively, and as such, may handle N2 signaling from SMF and AMF for PDU sessions and QoS, encapsulate/de-encapsulate packets for IPsec and N3 tunneling, mark N3 user-plane packets in uplink, and enforce QoS corresponding to N3 packet marking taking into account QoS requirements associated to such marking received over N2. In at least one embodiment, N3IWF may also relay uplink and downlink control-plane NAS (NI) signaling between UE **2102** and AMF **2112**, and relay uplink and downlink user-plane packets between UE **2102** and UPF **2104**. In at least one embodiment, N3IWF also provides mechanisms for IPsec tunnel establishment with UE **2102**.

[0251] In at least one embodiment, SMF **2118** may be responsible for session management (e.g., session establishment, modify and release, including tunnel maintain between UPF and AN node); UE IP address allocation & management (including optional Authorization); Selection and control of UP function; Configures traffic steering at UPF to route traffic to proper destination; termination of interfaces towards Policy control functions; control part of policy enforcement and QoS; lawful intercept (for SM events and interface to LI System); termination of SM parts of NAS messages; downlink Data Notification; initiator of AN specific SM information, sent via AMF over N2 to AN; determine SSC mode of a session. In at least one embodiment, SMF **2118** may include following roaming functionality: handle local enforcement to apply QoS SLAB (VPLMN); charging data collection and charging interface (VPLMN); lawful intercept (in VPLMN for SM events and interface to LI System); support for interaction with external DN for transport of signaling for PDU session authorization/authentication by external DN.

[0252] In at least one embodiment, NEF **2116** may provide means for securely exposing services and capabilities provided by 3GPP network functions for third party, internal

exposure/re-exposure, Application Functions (e.g., AF **2126**), edge computing or fog computing systems, etc. In at least one embodiment, NEF **2116** may authenticate, authorize, and/or throttle AFs. In at least one embodiment, NEF **2116** may also translate information exchanged with AF **2126** and information exchanged with internal network functions. In at least one embodiment, NEF **2116** may translate between an AF-Service-Identifier and an internal 5GC information. In at least one embodiment, NEF **2116** may also receive information from other network functions (NFs) based on exposed capabilities of other network functions. In at least one embodiment, this information may be stored at NEF **2116** as structured data, or at a data storage NF using a standardized interfaces. In at least one embodiment, stored information can then be re-exposed by NEF **2116** to other NFs and AFs, and/or used for other purposes such as analytics.

[0253] In at least one embodiment, NRF **2120** may support service discovery functions, receive NF Discovery Requests from NF instances, and provide information of discovered NF instances to NF instances. In at least one embodiment, NRF **2120** also maintains information of available NF instances and their supported services.

[0254] In at least one embodiment, PCF **2122** may provide policy rules to control plane function(s) to enforce them, and may also support unified policy framework to govern network behavior. In at least one embodiment, PCF **2122** may also implement a front end (FE) to access subscription information relevant for policy decisions in a UDR of UDM **2124**.

[0255] In at least one embodiment, UDM **2124** may handle subscription-related information to support a network entities' handling of communication sessions, and may store subscription data of UE **2102**. In at least one embodiment, UDM **2124** may include two parts, an application FE and a User Data Repository (UDR). In at least one embodiment, UDM may include a UDM FE, which is in charge of processing of credentials, location management, subscription management and so on. In at least one embodiment, several different front ends may serve a same user in different transactions. In at least one embodiment, UDM-FE accesses subscription information stored in an UDR and performs authentication credential processing; user identification handling; access authorization; registration/mobility management; and subscription management. In at least one embodiment, UDR may interact with PCF **2122**. In at least one embodiment, UDM **2124** may also support SMS management, wherein an SMS-FE implements a similar application logic as discussed previously.

[0256] In at least one embodiment, AF **2126** may provide application influence on traffic routing, access to a Network Capability Exposure (NCE), and interact with a policy framework for policy control. In at least one embodiment, NCE may be a mechanism that allows a 5GC and AF **2126** to provide information to each other via NEF **2116**, which may be used for edge computing implementations. In at least one embodiment, network operator and third party services may be hosted close to UE **2102** access point of attachment to achieve an efficient service delivery through a reduced end-to-end latency and load on a transport network. In at least one embodiment, for edge computing implementations, 5GC may select a UPF **2104** close to UE **2102** and execute traffic steering from UPF **2104** to DN **2106** via N6 interface. In at least one embodiment, this may be based on UE

subscription data, UE location, and information provided by AF **2126**. In at least one embodiment, AF **2126** may influence UPF (re)selection and traffic routing. In at least one embodiment, based on operator deployment, when AF **2126** is considered to be a trusted entity, a network operator may permit AF **2126** to interact directly with relevant NFs.

[**0257**] In at least one embodiment, CN **2110** may include an SMSF, which may be responsible for SMS subscription checking and verification, and relaying SM messages to/from UE **2102** to/from other entities, such as an SMS-GMSC/IWMS-C/SMS-router. In at least one embodiment, SMS may also interact with AMF **2112** and UDM **2124** for notification procedure that UE **2102** is available for SMS transfer (e.g., set a UE not reachable flag, and notifying UDM **2124** when UE **2102** is available for SMS).

[**0258**] In at least one embodiment, system **2100** may include following service-based interfaces: Namf: Service-based interface exhibited by AMF; Nsmf: Service-based interface exhibited by SMF; Nnef: Service-based interface exhibited by NEF; Npcf: Service-based interface exhibited by PCF; Nudm: Service-based interface exhibited by UDM; Naf: Service-based interface exhibited by AF; Nnrf: Service-based interface exhibited by NRF; and Nausf: Service-based interface exhibited by AUSF.

[**0259**] In at least one embodiment, system **2100** may include following reference points: N1: Reference point between UE and AMF; N2: Reference point between (R)AN and AMF; N3: Reference point between (R)AN and UPF; N4: Reference point between SMF and UPF; and N6: Reference point between UPF and a Data Network. In at least one embodiment, there may be many more reference points and/or service-based interfaces between a NF services in NFs, however, these interfaces and reference points have been omitted for clarity. In at least one embodiment, an NS reference point may be between a PCF and AF; an N7 reference point may be between PCF and SMF; an N11 reference point between AMF and SMF; etc. In at least one embodiment, CN **2110** may include an Nx interface, which is an inter-CN interface between MME and AMF **2112** in order to enable interworking between CN **2110** and CN **7221**.

[**0260**] In at least one embodiment, system **2100** may include multiple RAN nodes (such as (R)AN node **2108**) wherein an Xn interface is defined between two or more (R)AN node **2108** (e.g., gNBs) that connecting to 5GC **410**, between a (R)AN node **2108** (e.g., gNB) connecting to CN **2110** and an eNB (e.g., a macro RAN node), and/or between two eNBs connecting to CN **2110**.

[**0261**] In at least one embodiment, Xn interface may include an Xn user plane (Xn-U) interface and an Xn control plane (Xn-C) interface. In at least one embodiment, Xn-U may provide non-guaranteed delivery of user plane PDUs and support/provide data forwarding and flow control functionality. In at least one embodiment, Xn-C may provide management and error handling functionality, functionality to manage a Xn-C interface; mobility support for UE **2102** in a connected mode (e.g., CM-CONNECTED) including functionality to manage UE mobility for connected mode between one or more (R)AN node **2108**. In at least one embodiment, mobility support may include context transfer from an old (source) serving (R)AN node **2108** to new (target) serving (R)AN node **2108**; and control of user plane tunnels between old (source) serving (R)AN node **2108** to new (target) serving (R)AN node **2108**.

[**0262**] In at least one embodiment, a protocol stack of a Xn-U may include a transport network layer built on Internet Protocol (IP) transport layer, and a GTP-U layer on top of a UDP and/or IP layer(s) to carry user plane PDUs. In at least one embodiment, Xn-C protocol stack may include an application layer signaling protocol (referred to as Xn Application Protocol (Xn-AP)) and a transport network layer that is built on an SCTP layer. In at least one embodiment, SCTP layer may be on top of an IP layer. In at least one embodiment, SCTP layer provides a guaranteed delivery of application layer messages. In at least one embodiment, in a transport IP layer point-to-point transmission is used to deliver signaling PDUs. In at least one embodiment, Xn-U protocol stack and/or a Xn-C protocol stack may be same or similar to a user plane and/or control plane protocol stack(s) shown and described herein.

[**0263**] FIG. **22** is an illustration of a control plane protocol stack in accordance with some embodiments. In at least one embodiment, a control plane **2200** is shown as a communications protocol stack between UE **2002** (or alternatively, UE **2004**), RAN **2016**, and MME(s) **2028**.

[**0264**] In at least one embodiment, PHY layer **2202** may transmit or receive information used by MAC layer **2204** over one or more air interfaces. In at least one embodiment, PHY layer **2202** may further perform link adaptation or adaptive modulation and coding (AMC), power control, cell search (e.g., for initial synchronization and handover purposes), and other measurements used by higher layers, such as an RRC layer **2210**. In at least one embodiment, PHY layer **2202** may still further perform error detection on transport channels, forward error correction (FEC) coding/de-coding of transport channels, modulation/demodulation of physical channels, interleaving, rate matching, mapping onto physical channels, and Multiple Input Multiple Output (MIMO) antenna processing.

[**0265**] In at least one embodiment, MAC layer **2204** may perform mapping between logical channels and transport channels, multiplexing of MAC service data units (SDUs) from one or more logical channels onto transport blocks (TB) to be delivered to PHY via transport channels, demultiplexing MAC SDUs to one or more logical channels from transport blocks (TB) delivered from PHY via transport channels, multiplexing MAC SDUs onto TBs, scheduling information reporting, error correction through hybrid automatic repeat request (HARD), and logical channel prioritization.

[**0266**] In at least one embodiment, RLC layer **2206** may operate in a plurality of modes of operation, including: Transparent Mode™, Unacknowledged Mode (UM), and Acknowledged Mode (AM). In at least one embodiment, RLC layer **2206** may execute transfer of upper layer protocol data units (PDUs), error correction through automatic repeat request (ARQ) for AM data transfers, and concatenation, segmentation and reassembly of RLC SDUs for UM and AM data transfers. In at least one embodiment, RLC layer **2206** may also execute re-segmentation of RLC data PDUs for AM data transfers, reorder RLC data PDUs for UM and AM data transfers, detect duplicate data for UM and AM data transfers, discard RLC SDUs for UM and AM data transfers, detect protocol errors for AM data transfers, and perform RLC re-establishment.

[**0267**] In at least one embodiment, PDCP layer **2208** may execute header compression and decompression of IP data, maintain PDCP Sequence Numbers (SNs), perform in-se-

quence delivery of upper layer PDUs at re-establishment of lower layers, eliminate duplicates of lower layer SDUs at re-establishment of lower layers for radio bearers mapped on RLC AM, cipher and decipher control plane data, perform integrity protection and integrity verification of control plane data, control timer-based discard of data, and perform security operations (e.g., ciphering, deciphering, integrity protection, integrity verification, etc.).

[0268] In at least one embodiment, main services and functions of a RRC layer **2210** may include broadcast of system information (e.g., included in Master Information Blocks (MIBs) or System Information Blocks (SIBs) related to a non-access stratum (NAS)), broadcast of system information related to an access stratum (AS), paging, establishment, maintenance and release of an RRC connection between an UE and E-UTRAN (e.g., RRC connection paging, RRC connection establishment, RRC connection modification, and RRC connection release), establishment, configuration, maintenance and release of point-to-point radio bearers, security functions including key management, inter radio access technology (RAT) mobility, and measurement configuration for UE measurement reporting. In at least one embodiment, said MIBs and SIBs may comprise one or more information elements (IEs), which may each comprise individual data fields or data structures.

[0269] In at least one embodiment, UE **2002** and RAN **2016** may utilize a Uu interface (e.g., an LTE-Uu interface) to exchange control plane data via a protocol stack comprising PHY layer **2202**, MAC layer **2204**, RLC layer **2206**, PDCP layer **2208**, and RRC layer **2210**.

[0270] In at least one embodiment, non-access stratum (NAS) protocols (NAS protocols **2212**) form a highest stratum of a control plane between UE **2002** and MME(s) **2028**. In at least one embodiment, NAS protocols **2212** support mobility of UE **2002** and session management procedures to establish and maintain IP connectivity between UE **2002** and P-GW **2034**.

[0271] In at least one embodiment, Si Application Protocol (S1-AP) layer (Si-AP layer **2222**) may support functions of a Si interface and comprise Elementary Procedures (EPs). In at least one embodiment, an EP is a unit of interaction between RAN **2016** and CN **2028**. In at least one embodiment, S1-AP layer services may comprise two groups: UE-associated services and non UE-associated services. In at least one embodiment, these services perform functions including, but not limited to: E-UTRAN Radio Access Bearer (E-RAB) management, UE capability indication, mobility, NAS signaling transport, RAN Information Management (RIM), and configuration transfer.

[0272] In at least one embodiment, Stream Control Transmission Protocol (SCTP) layer (alternatively referred to as a stream control transmission protocol/internet protocol (SCTP/IP) layer) (SCTP layer **2220**) may ensure reliable delivery of signaling messages between RAN **2016** and MME(s) **2028** based, in part, on an IP protocol, supported by an IP layer **2218**. In at least one embodiment, L2 layer **2216** and an L1 layer **2214** may refer to communication links (e.g., wired or wireless) used by a RAN node and MME to exchange information.

[0273] In at least one embodiment, RAN **2016** and MME(s) **2028** may utilize an S1-MME interface to exchange control plane data via a protocol stack comprising a L1 layer **2214**, L2 layer **2216**, IP layer **2218**, SCTP layer **2220**, and Si-AP layer **2222**.

[0274] FIG. **23** is an illustration of a user plane protocol stack in accordance with at least one embodiment. In at least one embodiment, a user plane **2300** is shown as a communications protocol stack between a UE **2002**, RAN **2016**, S-GW **2030**, and P-GW **2034**. In at least one embodiment, user plane **2300** may utilize a same protocol layers as control plane **2200**. In at least one embodiment, UE **2002** and RAN **2016** may utilize a Uu interface (e.g., an LTE-Uu interface) to exchange user plane data via a protocol stack comprising PHY layer **2202**, MAC layer **2204**, RLC layer **2206**, PDCP layer **2208**.

[0275] In at least one embodiment, General Packet Radio Service (GPRS) Tunneling Protocol for a user plane (GTP-U) layer (GTP-U layer **2304**) may be used for carrying user data within a GPRS core network and between a radio access network and a core network. In at least one embodiment, user data transported can be packets in any of IPv4, IPv6, or PPP formats. In at least one embodiment, UDP and IP security (UDP/IP) layer (UDP/IP layer **2302**) may provide checksums for data integrity, port numbers for addressing different functions at a source and destination, and encryption and authentication on selected data flows. In at least one embodiment, RAN **2016** and S-GW **2030** may utilize an S1-U interface to exchange user plane data via a protocol stack comprising L1 layer **2214**, L2 layer **2216**, UDP/IP layer **2302**, and GTP-U layer **2304**. In at least one embodiment, S-GW **2030** and P-GW **2034** may utilize an S5/S8a interface to exchange user plane data via a protocol stack comprising L1 layer **2214**, L2 layer **2216**, UDP/IP layer **2302**, and GTP-U layer **2304**. In at least one embodiment, as discussed above with respect to FIG. **22**, NAS protocols support a mobility of UE **2002** and session management procedures to establish and maintain IP connectivity between UE **2002** and P-GW **2034**.

[0276] FIG. **24** illustrates components **2400** of a core network in accordance with at least one embodiment. In at least one embodiment, components of CN **2038** may be implemented in one physical node or separate physical nodes including components to read and execute instructions from a machine-readable or computer-readable medium (e.g., a non-transitory machine-readable storage medium). In at least one embodiment, Network Functions Virtualization (NFV) is utilized to virtualize any or all of above described network node functions via executable instructions stored in one or more computer readable storage mediums (described in further detail below). In at least one embodiment, a logical instantiation of CN **2038** may be referred to as a network slice **2402** (e.g., network slice **2402** is shown to include HSS **2032**, MME(s) **2028**, and S-GW **2030**). In at least one embodiment, a logical instantiation of a portion of CN **2038** may be referred to as a network sub-slice **2404** (e.g., network sub-slice **2404** is shown to include P-GW **2034** and PCRF **2036**).

[0277] In at least one embodiment, NFV architectures and infrastructures may be used to virtualize one or more network functions, alternatively performed by proprietary hardware, onto physical resources comprising a combination of industry-standard server hardware, storage hardware, or switches. In at least one embodiment, NFV systems can be used to execute virtual or reconfigurable implementations of one or more EPC components/functions.

[0278] FIG. **25** is a block diagram illustrating components, according to at least one embodiment, of a system **2500** to support network function virtualization (NFV). In at least

one embodiment, system **2500** is illustrated as including a virtualized infrastructure manager (shown as **VIM 2502**), a network function virtualization infrastructure (shown as **NFVI 2504**), a VNF manager (shown as **VNFM 2506**), virtualized network functions (shown as **VNF 2508**), an element manager (shown as **EM 2510**), an NFV Orchestrator (shown as **NFVO 2512**), and a network manager (shown as **NM 2514**).

[0279] In at least one embodiment, **VIM 2502** manages resources of **NFVI 2504**. In at least one embodiment, **NFVI 2504** can include physical or virtual resources and applications (including hypervisors) used to execute system **2500**. In at least one embodiment, **VIM 2502** may manage a life cycle of virtual resources with **NFVI 2504** (e.g., creation, maintenance, and tear down of virtual machines (VMs) associated with one or more physical resources), track VM instances, track performance, fault and security of VM instances and associated physical resources, and expose VM instances and associated physical resources to other management systems.

[0280] In at least one embodiment, **VNFM 2506** may manage **VNF 2508**. In at least one embodiment, **VNF 2508** may be used to execute EPC components/functions. In at least one embodiment, **VNFM 2506** may manage a life cycle of **VNF 2508** and track performance, fault and security of virtual aspects of **VNF 2508**. In at least one embodiment, **EM 2510** may track performance, fault and security of functional aspects of **VNF 2508**. In at least one embodiment, tracking data from **VNFM 2506** and **EM 2510** may comprise, in at least one embodiment, performance measurement (PM) data used by **VIM 2502** or **NFVI 2504**. In at least one embodiment, both **VNFM 2506** and **EM 2510** can scale up/down a quantity of VNFs of system **2500**.

[0281] In at least one embodiment, **NFVO 2512** may coordinate, authorize, release and engage resources of **NFVI 2504** in order to provide a requested service (e.g., to execute an EPC function, component, or slice). In at least one embodiment, **NM 2514** may provide a package of end-user functions with responsibility for a management of a network, which may include network elements with VNFs, non-virtualized network functions, or both (management of VNFs may occur via an **EM 2510**).

Computer-Based Systems

[0282] The following figures set forth, without limitation, exemplary computer-based systems that can be used to implement at least one embodiment.

[0283] FIG. 26 illustrates a processing system **2600**, in accordance with at least one embodiment. In at least one embodiment, processing system **2600** includes one or more processors **2602** and one or more graphics processors **2608**, and may be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **2602** or processor cores **2607**. In at least one embodiment, processing system **2600** is a processing platform incorporated within a system-on-a-chip (“SoC”) integrated circuit for use in mobile, handheld, or embedded devices.

[0284] In at least one embodiment, processing system **2600** can include, or be incorporated within a server-based gaming platform, a game console, a media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, processing system **2600** is a mobile phone, smart phone, tablet com-

puting device or mobile Internet device. In at least one embodiment, processing system **2600** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart eyewear device, augmented reality device, or virtual reality device. In at least one embodiment, processing system **2600** is a television or set top box device having one or more processors **2602** and a graphical interface generated by one or more graphics processors **2608**.

[0285] In at least one embodiment, one or more processors **2602** each include one or more processor cores **2607** to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, each of one or more processor cores **2607** is configured to process a specific instruction set **2609**. In at least one embodiment, instruction set **2609** may facilitate Complex Instruction Set Computing (“CISC”), Reduced Instruction Set Computing (“RISC”), or computing via a Very Long Instruction Word (“VLIW”). In at least one embodiment, processor cores **2607** may each process a different instruction set **2609**, which may include instructions to facilitate emulation of other instruction sets. In at least one embodiment, processor core **2607** may also include other processing devices, such as a digital signal processor (“DSP”).

[0286] In at least one embodiment, processor **2602** includes cache memory (“cache”) **2604**. In at least one embodiment, processor **2602** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory is shared among various components of processor **2602**. In at least one embodiment, processor **2602** also uses an external cache (e.g., a Level 3 (“L3”) cache or Last Level Cache (“LLC”)) (not shown), which may be shared among processor cores **2607** using known cache coherency techniques. In at least one embodiment, register file **2606** is additionally included in processor **2602** which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **2606** may include general-purpose registers or other registers.

[0287] In at least one embodiment, one or more processor(s) **2602** are coupled with one or more interface bus(es) **2610** to transmit communication signals such as address, data, or control signals between processor **2602** and other components in processing system **2600**. In at least one embodiment interface bus **2610**, in one embodiment, can be a processor bus, such as a version of a Direct Media Interface (“DMI”) bus. In at least one embodiment, interface bus **2610** is not limited to a DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., “PCI,” PCI Express (“PCIe”)), memory buses, or other types of interface buses. In at least one embodiment processor(s) **2602** include an integrated memory controller **2616** and a platform controller hub **2630**. In at least one embodiment, memory controller **2616** facilitates communication between a memory device and other components of processing system **2600**, while platform controller hub (“PCH”) **2630** provides connections to Input/Output (“I/O”) devices via a local I/O bus.

[0288] In at least one embodiment, memory device **2620** can be a dynamic random access memory (“DRAM”) device, a static random access memory (“SRAM”) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to

serve as processor memory. In at least one embodiment memory device **2620** can operate as system memory for processing system **2600**, to store data **2622** and instructions **2621** for use when one or more processors **2602** executes an application or process. In at least one embodiment, memory controller **2616** also couples with an optional external graphics processor **2612**, which may communicate with one or more graphics processors **2608** in processors **2602** to perform graphics and media operations. In at least one embodiment, a display device **2611** can connect to processor(s) **2602**. In at least one embodiment display device **2611** can include one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device **2611** can include a head mounted display (“HMD”) such as a stereoscopic display device for use in virtual reality (“VR”) applications or augmented reality (“AR”) applications.

[0289] In at least one embodiment, platform controller hub **2630** enables peripherals to connect to memory device **2620** and processor **2602** via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, an audio controller **2646**, a network controller **2634**, a firmware interface **2628**, a wireless transceiver **2626**, touch sensors **2625**, a data storage device **2624** (e.g., hard disk drive, flash memory, etc.). In at least one embodiment, data storage device **2624** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as PCI, or PCIe. In at least one embodiment, touch sensors **2625** can include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver **2626** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (“LTE”) transceiver. In at least one embodiment, firmware interface **2628** enables communication with system firmware, and can be, in at least one embodiment, a unified extensible firmware interface (“UEFI”). In at least one embodiment, network controller **2634** can enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not shown) couples with interface bus **2610**. In at least one embodiment, audio controller **2646** is a multi-channel high definition audio controller. In at least one embodiment, processing system **2600** includes an optional legacy I/O controller **2640** for coupling legacy (e.g., Personal System 2 (“PS/2”)) devices to processing system **2600**. In at least one embodiment, platform controller hub **2630** can also connect to one or more Universal Serial Bus (“USB”) controllers **2642** connect input devices, such as keyboard and mouse **2643** combinations, a camera **2644**, or other USB input devices.

[0290] In at least one embodiment, an instance of memory controller **2616** and platform controller hub **2630** may be integrated into a discreet external graphics processor, such as external graphics processor **2612**. In at least one embodiment, platform controller hub **2630** and/or memory controller **2616** may be external to one or more processor(s) **2602**. In at least one embodiment, processing system **2600** can include an external memory controller **2616** and platform controller hub **2630**, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with processor(s) **2602**.

[0291] FIG. 27 illustrates a computer system **2700**, in accordance with at least one embodiment. In at least one embodiment, computer system **2700** may be a system with

interconnected devices and components, an SOC, or some combination. In at least one embodiment, computer system **2700** is formed with a processor **2702** that may include execution units to execute an instruction. In at least one embodiment, computer system **2700** may include, without limitation, a component, such as processor **2702** to employ execution units including logic to perform algorithms for processing data. In at least one embodiment, computer system **2700** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, Calif., although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **2700** may execute a version of WINDOWS® operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux in at least one embodiment), embedded software, and/or graphical user interfaces, may also be used.

[0292] In at least one embodiment, computer system **2700** may be used in other devices such as handheld devices and embedded applications. Some ones of the at least one embodiments of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants (“PDAs”), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (DSP), an SoC, network computers (“NetPCs”), set-top boxes, network hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions.

[0293] In at least one embodiment, computer system **2700** may include, without limitation, processor **2702** that may include, without limitation, one or more execution units **2708** that may be configured to execute a Compute Unified Device Architecture (“CUDA”) (CUDA® is developed by NVIDIA Corporation of Santa Clara, Calif.) program. In at least one embodiment, a CUDA program is at least a portion of a software application written in a CUDA programming language. In at least one embodiment, computer system **2700** is a single processor desktop or server system. In at least one embodiment, computer system **2700** may be a multiprocessor system. In at least one embodiment, processor **2702** may include, without limitation, a CISC microprocessor, a RISC microprocessor, a VLIW microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, in at least one embodiment. In at least one embodiment, processor **2702** may be coupled to a processor bus **2710** that may transmit data signals between processor **2702** and other components in computer system **2700**.

[0294] In at least one embodiment, processor **2702** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **2704**. In at least one embodiment, processor **2702** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **2702**. In at least one embodiment, processor **2702** may also include a combination of both internal and external caches. In at least one embodiment, a register file **2706** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

[0295] In at least one embodiment, execution unit **2708**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **2702**. Processor **2702** may also include a microcode (“ucode”)

read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **2708** may include logic to handle a packed instruction set **2709**. In at least one embodiment, by including packed instruction set **2709** in an instruction set of a general-purpose processor **2702**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **2702**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor’s data bus for performing operations on packed data, which may eliminate a need to transfer smaller units of data across a processor’s data bus to perform one or more operations one data element at a time.

[0296] In at least one embodiment, execution unit **2708** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **2700** may include, without limitation, a memory **2720**. In at least one embodiment, memory **2720** may be implemented as a DRAM device, an SRAM device, flash memory device, or other memory device. Memory **2720** may store instruction (s) **2719** and/or data **2721** represented by data signals that may be executed by processor **2702**.

[0297] In at least one embodiment, a system logic chip may be coupled to processor bus **2710** and memory **2720**. In at least one embodiment, a system logic chip may include, without limitation, a memory controller hub (“MCH”) **2716**, and processor **2702** may communicate with MCH **2716** via processor bus **2710**. In at least one embodiment, MCH **2716** may provide a high bandwidth memory path **2718** to memory **2720** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **2716** may direct data signals between processor **2702**, memory **2720**, and other components in computer system **2700** and to bridge data signals between processor bus **2710**, memory **2720**, and a system I/O **2722**. In at least one embodiment, system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **2716** may be coupled to memory **2720** through high bandwidth memory path **2718** and graphics/video card **2712** may be coupled to MCH **2716** through an Accelerated Graphics Port (“AGP”) interconnect **2714**.

[0298] In at least one embodiment, computer system **2700** may use system I/O **2722** that is a proprietary hub interface bus to couple MCH **2716** to I/O controller hub (“ICH”) **2730**. In at least one embodiment, ICH **2730** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **2720**, a chipset, and processor **2702**. Examples may include, without limitation, an audio controller **2729**, a firmware hub (“flash BIOS”) **2728**, a wireless transceiver **2726**, a data storage **2724**, a legacy I/O controller **2723** containing a user input interface **2725** and a keyboard interface, a serial expansion port **2777**, such as a USB, and a network controller **2734**. Data storage **2724** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

[0299] In at least one embodiment, FIG. 27 illustrates a system, which includes interconnected hardware devices or “chips.” In at least one embodiment, FIG. 27 may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. 27 may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe),

or some combination thereof. In at least one embodiment, one or more components of system **2700** are interconnected using compute express link (“CXL”) interconnects.

[0300] FIG. 28 illustrates a system **2800**, in accordance with at least one embodiment. In at least one embodiment, system **2800** is an electronic device that utilizes a processor **2810**. In at least one embodiment, system **2800** may be, in at least one embodiment and without limitation, a notebook, a tower server, a rack server, a blade server, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, or any other suitable electronic device.

[0301] In at least one embodiment, system **2800** may include, without limitation, processor **2810** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **2810** is coupled using a bus or interface, such as an I2C bus, a System Management Bus (“SMBus”), a Low Pin Count (“LPC”) bus, a Serial Peripheral Interface (“SPI”), a High Definition Audio (“HDA”) bus, a Serial Advance Technology Attachment (“SATA”) bus, a USB (versions 1, 2, 3), or a Universal Asynchronous Receiver/Transmitter (“UART”) bus. In at least one embodiment, FIG. 28 illustrates a system which includes interconnected hardware devices or “chips.” In at least one embodiment, FIG. 28 may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. 28 may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of FIG. 28 are interconnected using CXL interconnects.

[0302] In at least one embodiment, FIG. 28 may include a display **2824**, a touch screen **2825**, a touch pad **2830**, a Near Field Communications unit (“NFC”) **2845**, a sensor hub **2840**, a thermal sensor **2846**, an Express Chipset (“EC”) **2835**, a Trusted Platform Module (“TPM”) **2838**, BIOS/firmware/flash memory (“BIOS, FW Flash”) **2822**, a DSP **2860**, a Solid State Disk (“SSD”) or Hard Disk Drive (“HDD”) **2820**, a wireless local area network unit (“WLAN”) **2850**, a Bluetooth unit **2852**, a Wireless Wide Area Network unit (“WWAN”) **2856**, a Global Positioning System (“GPS”) **2855**, a camera (“USB 3.0 camera”) **2854** such as a USB 3.0 camera, or a Low Power Double Data Rate (“LPDDR”) memory unit (“LPDDR3”) **2815** implemented, in at least one embodiment, LPDDR3 standard. These components may each be implemented in any suitable manner.

[0303] In at least one embodiment, other components may be communicatively coupled to processor **2810** through components discussed above. In at least one embodiment, an accelerometer **2841**, an Ambient Light Sensor (“ALS”) **2842**, a compass **2843**, and a gyroscope **2844** may be communicatively coupled to sensor hub **2840**. In at least one embodiment, a thermal sensor **2839**, a fan **2837**, a keyboard **2846**, and a touch pad **2830** may be communicatively coupled to EC **2835**. In at least one embodiment, a speaker **2863**, a headphones **2864**, and a microphone (“mic”) **2865** may be communicatively coupled to an audio unit (“audio codec and class d amp”) **2864**, which may in turn be communicatively coupled to DSP **2860**. In at least one embodiment, audio unit **2864** may include, without limitation, an audio coder/decoder (“codec”) and a class D amplifier. In at least one embodiment, a SIM card (“SIM”) **2857** may be communicatively coupled to WWAN unit **2856**. In at least one embodiment, components such as WLAN unit **2850** and Bluetooth unit **2852**, as well as WWAN unit **2856** may be implemented in a Next Generation Form Factor (“NGFF”).

[0304] FIG. 29 illustrates an exemplary integrated circuit 2900, in accordance with at least one embodiment. In at least one embodiment, exemplary integrated circuit 2900 is an SoC that may be fabricated using one or more IP cores. In at least one embodiment, integrated circuit 2900 includes one or more application processor(s) 2905 (e.g., CPUs), at least one graphics processor 2910, and may additionally include an image processor 2915 and/or a video processor 2920, any of which may be a modular IP core. In at least one embodiment, integrated circuit 2900 includes peripheral or bus logic including a USB controller 2925, a UART controller 2930, an SPI/SDIO controller 2935, and an I2S/I2C controller 2940. In at least one embodiment, integrated circuit 2900 can include a display device 2945 coupled to one or more of a high-definition multimedia interface (“HDMI”) controller 2950 and a mobile industry processor interface (“MIPI”) display interface 2955. In at least one embodiment, storage may be provided by a flash memory subsystem 2960 including flash memory and a flash memory controller. In at least one embodiment, a memory interface may be provided via a memory controller 2965 for access to SDRAM or SRAM memory devices. In at least one embodiment, some integrated circuits additionally include an embedded security engine 2970.

[0305] FIG. 30 illustrates a computing system 3000, according to at least one embodiment. In at least one embodiment, computing system 3000 includes a processing subsystem 3001 having one or more processor(s) 3002 and a system memory 3004 communicating via an interconnection path that may include a memory hub 3005. In at least one embodiment, memory hub 3005 may be a separate component within a chipset component or may be integrated within one or more processor(s) 3002. In at least one embodiment, memory hub 3005 couples with an I/O subsystem 3011 via a communication link 3006. In at least one embodiment, I/O subsystem 3011 includes an I/O hub 3007 that can enable computing system 3000 to receive input from one or more input device(s) 3008. In at least one embodiment, I/O hub 3007 can enable a display controller, which may be included in one or more processor(s) 3002, to provide outputs to one or more display device(s) 3010A. In at least one embodiment, one or more display device(s) 3010A coupled with I/O hub 3007 can include a local, internal, or embedded display device.

[0306] In at least one embodiment, processing subsystem 3001 includes one or more parallel processor(s) 3012 coupled to memory hub 3005 via a bus or other communication link 3013. In at least one embodiment, communication link 3013 may be one of any number of standards based communication link technologies or protocols, such as, but not limited to PCIe, or may be a vendor specific communications interface or communications fabric. In at least one embodiment, one or more parallel processor(s) 3012 form a computationally focused parallel or vector processing system that can include a large number of processing cores and/or processing clusters, such as a many integrated core processor. In at least one embodiment, one or more parallel processor(s) 3012 form a graphics processing subsystem that can output pixels to one of one or more display device(s) 3010A coupled via I/O Hub 3007. In at least one embodiment, one or more parallel processor(s) 3012 can also include a display controller and display interface (not shown) to enable a direct connection to one or more display device(s) 3010B.

[0307] In at least one embodiment, a system storage unit 3014 can connect to I/O hub 3007 to provide a storage mechanism for computing system 3000. In at least one

embodiment, an I/O switch 3016 can be used to provide an interface mechanism to enable connections between I/O hub 3007 and other components, such as a network adapter 3018 and/or wireless network adapter 3019 that may be integrated into a platform, and various other devices that can be added via one or more add-in device(s) 3020. In at least one embodiment, network adapter 3018 can be an Ethernet adapter or another wired network adapter. In at least one embodiment, wireless network adapter 3019 can include one or more of a Wi-Fi, Bluetooth, NFC, or other network device that includes one or more wireless radios.

[0308] In at least one embodiment, computing system 3000 can include other components not explicitly shown, including USB or other port connections, optical storage drives, video capture devices, and/or variations thereof, that may also be connected to I/O hub 3007. In at least one embodiment, communication paths interconnecting various components in FIG. 30 may be implemented using any suitable protocols, such as PCI based protocols (e.g., PCIe), or other bus or point-to-point communication interfaces and/or protocol(s), such as NVLink high-speed interconnect, or interconnect protocols.

[0309] In at least one embodiment, one or more parallel processor(s) 3012 incorporate circuitry optimized for graphics and video processing, including, in at least one embodiment, video output circuitry, and constitutes a graphics processing unit (“GPU”). In at least one embodiment, one or more parallel processor(s) 3012 incorporate circuitry optimized for general purpose processing. In at least one embodiment, components of computing system 3000 may be integrated with one or more other system elements on a single integrated circuit. In at least one embodiment, one or more parallel processor(s) 3012, memory hub 3005, processor(s) 3002, and I/O hub 3007 can be integrated into a SoC integrated circuit. In at least one embodiment, components of computing system 3000 can be integrated into a single package to form a system in package (“SIP”) configuration. In at least one embodiment, at least a portion of components of computing system 3000 can be integrated into a multi-chip module (“MCM”), which can be interconnected with other multi-chip modules into a modular computing system. In at least one embodiment, I/O subsystem 3011 and display devices 3010B are omitted from computing system 3000.

Processing Systems

[0310] The following figures set forth, without limitation, exemplary processing systems that can be used to implement at least one embodiment.

[0311] FIG. 31 illustrates an accelerated processing unit (“APU”) 3100, in accordance with at least one embodiment. In at least one embodiment, APU 3100 is developed by AMD Corporation of Santa Clara, Calif. In at least one embodiment, APU 3100 can be configured to execute an application program, such as a CUDA program. In at least one embodiment, APU 3100 includes, without limitation, a core complex 3110, a graphics complex 3140, fabric 3160, I/O interfaces 3170, memory controllers 3180, a display controller 3192, and a multimedia engine 3194. In at least one embodiment, APU 3100 may include, without limitation, any number of core complexes 3110, any number of graphics complexes 3150, any number of display controllers 3192, and any number of multimedia engines 3194 in any combination. For explanatory purposes, multiple instances of like objects are denoted herein with reference numbers identifying an object and parenthetical numbers identifying an instance where needed.

[0312] In at least one embodiment, core complex 3110 is a CPU, graphics complex 3140 is a GPU, and APU 3100 is a processing unit that integrates, without limitation, 3110 and 3140 onto a single chip. In at least one embodiment, some tasks may be assigned to core complex 3110 and other tasks may be assigned to graphics complex 3140. In at least one embodiment, core complex 3110 is configured to execute main control software associated with APU 3100, such as an operating system. In at least one embodiment, core complex 3110 is a master processor of APU 3100, controlling and coordinating operations of other processors. In at least one embodiment, core complex 3110 issues commands that control an operation of graphics complex 3140. In at least one embodiment, core complex 3110 can be configured to execute host executable code derived from CUDA source code, and graphics complex 3140 can be configured to execute device executable code derived from CUDA source code.

[0313] In at least one embodiment, core complex 3110 includes, without limitation, cores 3120(1)-3120(4) and an L3 cache 3130. In at least one embodiment, core complex 3110 may include, without limitation, any number of cores 3120 and any number and type of caches in any combination. In at least one embodiment, cores 3120 are configured to execute instructions of a particular instruction set architecture (“ISA”). In at least one embodiment, each core 3120 is a CPU core.

[0314] In at least one embodiment, each core 3120 includes, without limitation, a fetch/decode unit 3122, an integer execution engine 3124, a floating point execution engine 3126, and an L2 cache 3128. In at least one embodiment, fetch/decode unit 3122 fetches instructions, decodes such instructions, generates micro-operations, and dispatches separate micro-instructions to integer execution engine 3124 and floating point execution engine 3126. In at least one embodiment, fetch/decode unit 3122 can concurrently dispatch one micro-instruction to integer execution engine 3124 and another micro-instruction to floating point execution engine 3126. In at least one embodiment, integer execution engine 3124 executes, without limitation, integer and memory operations. In at least one embodiment, floating point engine 3126 executes, without limitation, floating point and vector operations. In at least one embodiment, fetch-decode unit 3122 dispatches micro-instructions to a single execution engine that replaces both integer execution engine 3124 and floating point execution engine 3126.

[0315] In at least one embodiment, each core 3120(*i*), where *i* is an integer representing a particular instance of core 3120, may access L2 cache 3128(*i*) included in core 3120(*i*). In at least one embodiment, each core 3120 included in core complex 3110(*j*), where *j* is an integer representing a particular instance of core complex 3110, is connected to other cores 3120 included in core complex 3110(*j*) via L3 cache 3130(*j*) included in core complex 3110(*j*). In at least one embodiment, cores 3120 included in core complex 3110(*j*), where *j* is an integer representing a particular instance of core complex 3110, can access all of L3 cache 3130(*j*) included in core complex 3110(*j*). In at least one embodiment, L3 cache 3130 may include, without limitation, any number of slices.

[0316] In at least one embodiment, graphics complex 3140 can be configured to perform compute operations in a highly-parallel fashion. In at least one embodiment, graphics complex 3140 is configured to execute graphics pipeline operations such as draw commands, pixel operations, geometric computations, and other operations associated with rendering an image to a display. In at least one embodiment,

graphics complex 3140 is configured to execute operations unrelated to graphics. In at least one embodiment, graphics complex 3140 is configured to execute both operations related to graphics and operations unrelated to graphics.

[0317] In at least one embodiment, graphics complex 3140 includes, without limitation, any number of compute units 3150 and an L2 cache 3142. In at least one embodiment, compute units 3150 share L2 cache 3142. In at least one embodiment, L2 cache 3142 is partitioned. In at least one embodiment, graphics complex 3140 includes, without limitation, any number of compute units 3150 and any number (including zero) and type of caches. In at least one embodiment, graphics complex 3140 includes, without limitation, any amount of dedicated graphics hardware.

[0318] In at least one embodiment, each compute unit 3150 includes, without limitation, any number of SIMD units 3152 and a shared memory 3154. In at least one embodiment, each SIMD unit 3152 implements a SIMD architecture and is configured to perform operations in parallel. In at least one embodiment, each compute unit 3150 may execute any number of thread blocks, but each thread block executes on a single compute unit 3150. In at least one embodiment, a thread block includes, without limitation, any number of threads of execution. In at least one embodiment, a workgroup is a thread block. In at least one embodiment, each SIMD unit 3152 executes a different warp. In at least one embodiment, a warp is a group of threads (e.g., 16 threads), where each thread in a warp belongs to a single thread block and is configured to process a different set of data based on a single set of instructions. In at least one embodiment, predication can be used to disable one or more threads in a warp. In at least one embodiment, a lane is a thread. In at least one embodiment, a work item is a thread. In at least one embodiment, a wavefront is a warp. In at least one embodiment, different wavefronts in a thread block may synchronize together and communicate via shared memory 3154.

[0319] In at least one embodiment, fabric 3160 is a system interconnect that facilitates data and control transmissions across core complex 3110, graphics complex 3140, I/O interfaces 3170, memory controllers 3180, display controller 3192, and multimedia engine 3194. In at least one embodiment, APU 3100 may include, without limitation, any amount and type of system interconnect in addition to or instead of fabric 3160 that facilitates data and control transmissions across any number and type of directly or indirectly linked components that may be internal or external to APU 3100. In at least one embodiment, I/O interfaces 3170 are representative of any number and type of I/O interfaces (e.g., PCI, PCI-Extended (“PCI-X”), PCIe, gigabit Ethernet (“GBE”), USB, etc.). In at least one embodiment, various types of peripheral devices are coupled to I/O interfaces 3170. In at least one embodiment, peripheral devices that are coupled to I/O interfaces 3170 may include, without limitation, keyboards, mice, printers, scanners, joysticks or other types of game controllers, media recording devices, external storage devices, network interface cards, and so forth.

[0320] In at least one embodiment, display controller AMD92 displays images on one or more display device(s), such as a liquid crystal display (“LCD”) device. In at least one embodiment, multimedia engine 240 includes, without limitation, any amount and type of circuitry that is related to multimedia, such as a video decoder, a video encoder, an image signal processor, etc. In at least one embodiment, memory controllers 3180 facilitate data transfers between APU 3100 and a unified system memory 3190. In at least

one embodiment, core complex **3110** and graphics complex **3140** share unified system memory **3190**.

[0321] In at least one embodiment, APU **3100** implements a memory subsystem that includes, without limitation, any amount and type of memory controllers **3180** and memory devices (e.g., shared memory **3154**) that may be dedicated to one component or shared among multiple components. In at least one embodiment, APU **3100** implements a cache subsystem that includes, without limitation, one or more cache memories (e.g., L2 caches **2728**, L3 cache **3130**, and L2 cache **3142**) that may each be private to or shared between any number of components (e.g., cores **3120**, core complex **3110**, SIMD units **3152**, compute units **3150**, and graphics complex **3140**).

[0322] FIG. 32 illustrates a CPU **3200**, in accordance with at least one embodiment. In at least one embodiment, CPU **3200** is developed by AMD Corporation of Santa Clara, Calif. In at least one embodiment, CPU **3200** can be configured to execute an application program. In at least one embodiment, CPU **3200** is configured to execute main control software, such as an operating system. In at least one embodiment, CPU **3200** issues commands that control an operation of an external GPU (not shown). In at least one embodiment, CPU **3200** can be configured to execute host executable code derived from CUDA source code, and an external GPU can be configured to execute device executable code derived from such CUDA source code. In at least one embodiment, CPU **3200** includes, without limitation, any number of core complexes **3210**, fabric **3260**, I/O interfaces **3270**, and memory controllers **3280**.

[0323] In at least one embodiment, core complex **3210** includes, without limitation, cores **3220(1)-3220(4)** and an L3 cache **3230**. In at least one embodiment, core complex **3210** may include, without limitation, any number of cores **3220** and any number and type of caches in any combination. In at least one embodiment, cores **3220** are configured to execute instructions of a particular ISA. In at least one embodiment, each core **3220** is a CPU core.

[0324] In at least one embodiment, each core **3220** includes, without limitation, a fetch/decode unit **3222**, an integer execution engine **3224**, a floating point execution engine **3226**, and an L2 cache **3228**. In at least one embodiment, fetch/decode unit **3222** fetches instructions, decodes such instructions, generates micro-operations, and dispatches separate micro-instructions to integer execution engine **3224** and floating point execution engine **3226**. In at least one embodiment, fetch/decode unit **3222** can concurrently dispatch one micro-instruction to integer execution engine **3224** and another micro-instruction to floating point execution engine **3226**. In at least one embodiment, integer execution engine **3224** executes, without limitation, integer and memory operations. In at least one embodiment, floating point engine **3226** executes, without limitation, floating point and vector operations. In at least one embodiment, fetch-decode unit **3222** dispatches micro-instructions to a single execution engine that replaces both integer execution engine **3224** and floating point execution engine **3226**.

[0325] In at least one embodiment, each core **3220(i)**, where *i* is an integer representing a particular instance of core **3220**, may access L2 cache **3228(i)** included in core **3220(i)**. In at least one embodiment, each core **3220** included in core complex **3210(j)**, where *j* is an integer representing a particular instance of core complex **3210**, is connected to other cores **3220** in core complex **3210(j)** via L3 cache **3230(j)** included in core complex **3210(j)**. In at least one embodiment, cores **3220** included in core complex **3210(j)**, where *j* is an integer representing a particular

instance of core complex **3210**, can access all of L3 cache **3230(j)** included in core complex **3210(j)**. In at least one embodiment, L3 cache **3230** may include, without limitation, any number of slices.

[0326] In at least one embodiment, fabric **3260** is a system interconnect that facilitates data and control transmissions across core complexes **3210(1)-3210(N)** (where *N* is an integer greater than zero), I/O interfaces **3270**, and memory controllers **3280**. In at least one embodiment, CPU **3200** may include, without limitation, any amount and type of system interconnect in addition to or instead of fabric **3260** that facilitates data and control transmissions across any number and type of directly or indirectly linked components that may be internal or external to CPU **3200**. In at least one embodiment, I/O interfaces **3270** are representative of any number and type of I/O interfaces (e.g., PCI, PCI-X, PCIe, GBE, USB, etc.). In at least one embodiment, various types of peripheral devices are coupled to I/O interfaces **3270**. In at least one embodiment, peripheral devices that are coupled to I/O interfaces **3270** may include, without limitation, displays, keyboards, mice, printers, scanners, joysticks or other types of game controllers, media recording devices, external storage devices, network interface cards, and so forth.

[0327] In at least one embodiment, memory controllers **3280** facilitate data transfers between CPU **3200** and a system memory **3290**. In at least one embodiment, core complex **3210** and graphics complex **3240** share system memory **3290**. In at least one embodiment, CPU **3200** implements a memory subsystem that includes, without limitation, any amount and type of memory controllers **3280** and memory devices that may be dedicated to one component or shared among multiple components. In at least one embodiment, CPU **3200** implements a cache subsystem that includes, without limitation, one or more cache memories (e.g., L2 caches **3228** and L3 caches **3230**) that may each be private to or shared between any number of components (e.g., cores **3220** and core complexes **3210**).

[0328] FIG. 33 illustrates an exemplary accelerator integration slice **3390**, in accordance with at least one embodiment. As used herein, a “slice” comprises a specified portion of processing resources of an accelerator integration circuit. In at least one embodiment, an accelerator integration circuit provides cache management, memory access, context management, and interrupt management services on behalf of multiple graphics processing engines included in a graphics acceleration module. Graphics processing engines may each comprise a separate GPU. Alternatively, graphics processing engines may comprise different types of graphics processing engines within a GPU such as graphics execution units, media processing engines (e.g., video encoders/decoders), samplers, and blit engines. In at least one embodiment, a graphics acceleration module may be a GPU with multiple graphics processing engines. In at least one embodiment, graphics processing engines may be individual GPUs integrated on a common package, line card, or chip.

[0329] An application effective address space **3382** within system memory **3314** stores process elements **3383**. In one embodiment, process elements **3383** are stored in response to GPU invocations **3381** from applications **3380** executed on processor **3307**. A process element **3383** contains process state for corresponding application **3380**. A work descriptor (“WD”) **3384** contained in process element **3383** can be a single job requested by an application or may contain a pointer to a queue of jobs. In at least one embodiment, WD **3384** is a pointer to a job request queue in application effective address space **3382**.

[0330] Graphics acceleration module 3346 and/or individual graphics processing engines can be shared by all or a subset of processes in a system. In at least one embodiment, an infrastructure for setting up process state and sending WD 3384 to graphics acceleration module 3346 to start a job in a virtualized environment may be included.

[0331] In at least one embodiment, a dedicated-process programming model is implementation-specific. In this model, a single process owns graphics acceleration module 3346 or an individual graphics processing engine. Because graphics acceleration module 3346 is owned by a single process, a hypervisor initializes an accelerator integration circuit for an owning partition and an operating system initializes accelerator integration circuit for an owning process when graphics acceleration module 3346 is assigned.

[0332] In operation, a WD fetch unit 3391 in accelerator integration slice 3390 fetches next WD 3384 which includes an indication of work to be done by one or more graphics processing engines of graphics acceleration module 3346. Data from WD 3384 may be stored in registers 3345 and used by a memory management unit (“MMU”) 3339, interrupt management circuit 3347 and/or context management circuit 3348 as illustrated. In at least one embodiment of MMU 3339 includes segment/page walk circuitry for accessing segment/page tables 3386 within OS virtual address space 3385. Interrupt management circuit 3347 may process interrupt events (“INT”) 3392 received from graphics acceleration module 3346. When performing graphics operations, an effective address 3393 generated by a graphics processing engine is translated to a real address by MMU 3339.

[0333] In one embodiment, a same set of registers 3345 are duplicated for each graphics processing engine and/or graphics acceleration module 3346 and may be initialized by a hypervisor or operating system. Each of these duplicated registers may be included in accelerator integration slice 3390. Exemplary registers that may be initialized by a hypervisor are shown in Table 1.

TABLE 1

Hypervisor Initialized Registers	
1	Slice Control Register
2	Real Address (RA) Scheduled Processes Area Pointer
3	Authority Mask Override Register
4	Interrupt Vector Table Entry Offset
5	Interrupt Vector Table Entry Limit
6	State Register
7	Logical Partition ID
8	Real address (RA) Hypervisor Accelerator Utilization Record Pointer
9	Storage Description Register

[0334] Exemplary registers that may be initialized by an operating system are shown in Table 2.

TABLE 2

Operating System Initialized Registers	
1	Process and Thread Identification
2	Effective Address (EA) Context Save/Restore Pointer
3	Virtual Address (VA) Accelerator Utilization Record Pointer
4	Virtual Address (VA) Storage Segment Table Pointer

TABLE 2-continued

Operating System Initialized Registers	
5	Authority Mask
6	Work descriptor

[0335] In one embodiment, each WD 3384 is specific to a particular graphics acceleration module 3346 and/or a particular graphics processing engine. It contains all information required by a graphics processing engine to do work or it can be a pointer to a memory location where an application has set up a command queue of work to be completed.

[0336] FIGS. 34A-34B illustrate exemplary graphics processors, in accordance with at least one embodiment. In at least one embodiment, any of the exemplary graphics processors may be fabricated using one or more IP cores. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores. In at least one embodiment, the exemplary graphics processors are for use within an SoC.

[0337] FIG. 34A illustrates an exemplary graphics processor 3410 of an SoC integrated circuit that may be fabricated using one or more IP cores, in accordance with at least one embodiment. FIG. 34B illustrates an additional exemplary graphics processor 3440 of an SoC integrated circuit that may be fabricated using one or more IP cores, in accordance with at least one embodiment. In at least one embodiment, graphics processor 3410 of FIG. 34A is a low power graphics processor core. In at least one embodiment, graphics processor 3440 of FIG. 34B is a higher performance graphics processor core. In at least one embodiment, each of graphics processors 3410, 3440 can be variants of graphics processor 510 of FIG. 5.

[0338] In at least one embodiment, graphics processor 3410 includes a vertex processor 3405 and one or more fragment processor(s) 3415A-3415N (e.g., 3415A, 3415B, 3415C, 3415D, through 3415N-1, and 3415N). In at least one embodiment, graphics processor 3410 can execute different shader programs via separate logic, such that vertex processor 3405 is optimized to execute operations for vertex shader programs, while one or more fragment processor(s) 3415A-3415N execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. In at least one embodiment, vertex processor 3405 performs a vertex processing stage of a 3D graphics pipeline and generates primitives and vertex data. In at least one embodiment, fragment processor(s) 3415A-3415N use primitive and vertex data generated by vertex processor 3405 to produce a framebuffer that is displayed on a display device. In at least one embodiment, fragment processor(s) 3415A-3415N are optimized to execute fragment shader programs as provided for in an OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in a Direct 3D API.

[0339] In at least one embodiment, graphics processor 3410 additionally includes one or more MMU(s) 3420A-3420B, cache(s) 3425A-3425B, and circuit interconnect(s) 3430A-3430B. In at least one embodiment, one or more MMU(s) 3420A-3420B provide for virtual to physical address mapping for graphics processor 3410, including for vertex processor 3405 and/or fragment processor(s) 3415A-3415N, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in one or more cache(s) 3425A-3425B. In at least one embodiment, one or more MMU(s) 3420A-3420B may be

synchronized with other MMUs within a system, including one or more MMUs associated with one or more application processor(s) 505, image processors 515, and/or video processors 520 of FIG. 5, such that each processor 505-520 can participate in a shared or unified virtual memory system. In at least one embodiment, one or more circuit interconnect(s) 3430A-3430B enable graphics processor 3410 to interface with other IP cores within an SoC, either via an internal bus of an SoC or via a direct connection.

[0340] In at least one embodiment, graphics processor 3440 includes one or more MMU(s) 3420A-3420B, caches 3425A-3425B, and circuit interconnects 3430A-3430B of graphics processor 3410 of FIG. 34A. In at least one embodiment, graphics processor 3440 includes one or more shader core(s) 3455A-3455N (e.g., 3455A, 3455B, 3455C, 3455D, 3455E, 3455F, through 3455N-1, and 3455N), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. In at least one embodiment, a number of shader cores can vary. In at least one embodiment, graphics processor 3440 includes an inter-core task manager 3445, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores 3455A-3455N and a tiling unit 3458 to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, in at least one embodiment to exploit local spatial coherence within a scene or to optimize use of internal caches.

[0341] FIG. 35A illustrates a graphics core 3500, in accordance with at least one embodiment. In at least one embodiment, graphics core 3500 may be included within graphics processor 2410 of FIG. 24. In at least one embodiment, graphics core 3500 may be a unified shader core 3455A-3455N as in FIG. 34B. In at least one embodiment, graphics core 3500 includes a shared instruction cache 3502, a texture unit 3518, and a cache/shared memory 3520 that are common to execution resources within graphics core 3500. In at least one embodiment, graphics core 3500 can include multiple slices 3501A-3501N or partition for each core, and a graphics processor can include multiple instances of graphics core 3500. Slices 3501A-3501N can include support logic including a local instruction cache 3504A-3504N, a thread scheduler 3506A-3506N, a thread dispatcher 3508A-3508N, and a set of registers 3510A-3510N. In at least one embodiment, slices 3501A-3501N can include a set of additional function units (“AFUs”) 3512A-3512N, floating-point units (“FPUs”) 3514A-3514N, integer arithmetic logic units (“ALUs”) 3516-3516N, address computational units (“ACUs”) 3513A-3513N, double-precision floating-point units (“DPFPUs”) 3515A-3515N, and matrix processing units (“MPUs”) 3517A-3517N.

[0342] In at least one embodiment, FPUs 3514A-3514N can perform single-precision (32-bit) and half-precision (16-bit) floating point operations, while DPFPUs 3515A-3515N perform double precision (64-bit) floating point operations. In at least one embodiment, ALUs 3516A-3516N can perform variable precision integer operations at 8-bit, 16-bit, and 32-bit precision, and can be configured for mixed precision operations. In at least one embodiment, MPUs 3517A-3517N can also be configured for mixed precision matrix operations, including half-precision floating point and 8-bit integer operations. In at least one embodiment, MPUs 3517-3517N can perform a variety of matrix operations to accelerate CUDA programs, including enabling support for accelerated general matrix to matrix

multiplication (“GEMM”). In at least one embodiment, AFUs 3512A-3512N can perform additional logic operations not supported by floating-point or integer units, including trigonometric operations (e.g., Sine, Cosine, etc.).

[0343] FIG. 35B illustrates a general-purpose graphics processing unit (“GPGPU”) 3530, in accordance with at least one embodiment. In at least one embodiment, GPGPU 3530 is highly-parallel and suitable for deployment on a multi-chip module. In at least one embodiment, GPGPU 3530 can be configured to enable highly-parallel compute operations to be performed by an array of GPUs. In at least one embodiment, GPGPU 3530 can be linked directly to other instances of GPGPU 3530 to create a multi-GPU cluster to improve execution time for CUDA programs. In at least one embodiment, GPGPU 3530 includes a host interface 3532 to enable a connection with a host processor. In at least one embodiment, host interface 3532 is a PCIe interface. In at least one embodiment, host interface 3532 can be a vendor specific communications interface or communications fabric. In at least one embodiment, GPGPU 3530 receives commands from a host processor and uses a global scheduler 3534 to distribute execution threads associated with those commands to a set of compute clusters 3536A-3536H. In at least one embodiment, compute clusters 3536A-3536H share a cache memory 3538. In at least one embodiment, cache memory 3538 can serve as a higher-level cache for cache memories within compute clusters 3536A-3536H.

[0344] In at least one embodiment, GPGPU 3530 includes memory 3544A-3544B coupled with compute clusters 3536A-3536H via a set of memory controllers 3542A-3542B. In at least one embodiment, memory 3544A-3544B can include various types of memory devices including DRAM or graphics random access memory, such as synchronous graphics random access memory (“SGRAM”), including graphics double data rate (“GDDR”) memory.

[0345] In at least one embodiment, compute clusters 3536A-3536H each include a set of graphics cores, such as graphics core 3500 of FIG. 35A, which can include multiple types of integer and floating point logic units that can perform computational operations at a range of precisions including suited for computations associated with CUDA programs. In at least one embodiment, at least a subset of floating point units in each of compute clusters 3536A-3536H can be configured to perform 16-bit or 32-bit floating point operations, while a different subset of floating point units can be configured to perform 64-bit floating point operations.

[0346] In at least one embodiment, multiple instances of GPGPU 3530 can be configured to operate as a compute cluster. In at least one embodiment, compute clusters 3536A-3536H may implement any technically feasible communication techniques for synchronization and data exchange. In at least one embodiment, multiple instances of GPGPU 3530 communicate over host interface 3532. In at least one embodiment, GPGPU 3530 includes an I/O hub 3539 that couples GPGPU 3530 with a GPU link 3540 that enables a direct connection to other instances of GPGPU 3530. In at least one embodiment, GPU link 3540 is coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of GPGPU 3530. In at least one embodiment GPU link 3540 couples with a high speed interconnect to transmit and receive data to other GPGPUs 3530 or parallel processors. In at least one embodiment, multiple instances of GPGPU 3530 are located in separate data processing systems and communicate via a network device that is accessible via host

interface 3532. In at least one embodiment GPU link 3540 can be configured to enable a connection to a host processor in addition to or as an alternative to host interface 3532. In at least one embodiment, GPGPU 3530 can be configured to execute a CUDA program.

[0347] FIG. 36A illustrates a parallel processor 3600, in accordance with at least one embodiment. In at least one embodiment, various components of parallel processor 3600 may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (“ASICs”), or FPGAs.

[0348] In at least one embodiment, parallel processor 3600 includes a parallel processing unit 3602. In at least one embodiment, parallel processing unit 3602 includes an I/O unit 3604 that enables communication with other devices, including other instances of parallel processing unit 3602. In at least one embodiment, I/O unit 3604 may be directly connected to other devices. In at least one embodiment, I/O unit 3604 connects with other devices via use of a hub or switch interface, such as memory hub 605. In at least one embodiment, connections between memory hub 605 and I/O unit 3604 form a communication link. In at least one embodiment, I/O unit 3604 connects with a host interface 3606 and a memory crossbar 3616, where host interface 3606 receives commands directed to performing processing operations and memory crossbar 3616 receives commands directed to performing memory operations.

[0349] In at least one embodiment, when host interface 3606 receives a command buffer via I/O unit 3604, host interface 3606 can direct work operations to perform those commands to a front end 3608. In at least one embodiment, front end 3608 couples with a scheduler 3610, which is configured to distribute commands or other work items to a processing array 3612. In at least one embodiment, scheduler 3610 ensures that processing array 3612 is properly configured and in a valid state before tasks are distributed to processing array 3612. In at least one embodiment, scheduler 3610 is implemented via firmware logic executing on a microcontroller. In at least one embodiment, microcontroller implemented scheduler 3610 is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, enabling rapid preemption and context switching of threads executing on processing array 3612. In at least one embodiment, host software can provide workloads for scheduling on processing array 3612 via one of multiple graphics processing doorbells. In at least one embodiment, workloads can then be automatically distributed across processing array 3612 by scheduler 3610 logic within a microcontroller including scheduler 3610.

[0350] In at least one embodiment, processing array 3612 can include up to “N” clusters (e.g., cluster 3614A, cluster 3614B, through cluster 3614N). In at least one embodiment, each cluster 3614A-3614N of processing array 3612 can execute a large number of concurrent threads. In at least one embodiment, scheduler 3610 can allocate work to clusters 3614A-3614N of processing array 3612 using various scheduling and/or work distribution algorithms, which may vary depending on a workload arising for each type of program or computation. In at least one embodiment, scheduling can be handled dynamically by scheduler 3610, or can be assisted in part by compiler logic during compilation of program logic configured for execution by processing array 3612. In at least one embodiment, different clusters 3614A-3614N of processing array 3612 can be allocated for processing different types of programs or for performing different types of computations.

[0351] In at least one embodiment, processing array 3612 can be configured to perform various types of parallel processing operations. In at least one embodiment, processing array 3612 is configured to perform general-purpose parallel compute operations. In at least one embodiment, processing array 3612 can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

[0352] In at least one embodiment, processing array 3612 is configured to perform parallel graphics processing operations. In at least one embodiment, processing array 3612 can include additional logic to support execution of such graphics processing operations, including, but not limited to texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. In at least one embodiment, processing array 3612 can be configured to execute graphics processing related shader programs such as, but not limited to vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. In at least one embodiment, parallel processing unit 3602 can transfer data from system memory via I/O unit 3604 for processing. In at least one embodiment, during processing, transferred data can be stored to on-chip memory (e.g., a parallel processor memory 3622) during processing, then written back to system memory.

[0353] In at least one embodiment, when parallel processing unit 3602 is used to perform graphics processing, scheduler 3610 can be configured to divide a processing workload into approximately equal sized tasks, to better enable distribution of graphics processing operations to multiple clusters 3614A-3614N of processing array 3612. In at least one embodiment, portions of processing array 3612 can be configured to perform different types of processing. In at least one embodiment, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. In at least one embodiment, intermediate data produced by one or more of clusters 3614A-3614N may be stored in buffers to allow intermediate data to be transmitted between clusters 3614A-3614N for further processing.

[0354] In at least one embodiment, processing array 3612 can receive processing tasks to be executed via scheduler 3610, which receives commands defining processing tasks from front end 3608. In at least one embodiment, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and commands defining how data is to be processed (e.g., what program is to be executed). In at least one embodiment, scheduler 3610 may be configured to fetch indices corresponding to tasks or may receive indices from front end 3608. In at least one embodiment, front end 3608 can be configured to ensure processing array 3612 is configured to a valid state before a workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

[0355] In at least one embodiment, each of one or more instances of parallel processing unit 3602 can couple with parallel processor memory 3622. In at least one embodiment, parallel processor memory 3622 can be accessed via memory crossbar 3616, which can receive memory requests from processing array 3612 as well as I/O unit 3604. In at least one embodiment, memory crossbar 3616 can access parallel processor memory 3622 via a memory interface

3618. In at least one embodiment, memory interface **3618** can include multiple partition units (e.g., a partition unit **3620A**, partition unit **3620B**, through partition unit **3620N**) that can each couple to a portion (e.g., memory unit) of parallel processor memory **3622**. In at least one embodiment, a number of partition units **3620A-3620N** is configured to be equal to a number of memory units, such that a first partition unit **3620A** has a corresponding first memory unit **3624A**, a second partition unit **3620B** has a corresponding memory unit **3624B**, and an Nth partition unit **3620N** has a corresponding Nth memory unit **3624N**. In at least one embodiment, a number of partition units **3620A-3620N** may not be equal to a number of memory devices.

[0356] In at least one embodiment, memory units **3624A-3624N** can include various types of memory devices, including DRAM or graphics random access memory, such as SGRAM, including GDDR memory. In at least one embodiment, memory units **3624A-3624N** may also include 3D stacked memory, including but not limited to high bandwidth memory (“HBM”). In at least one embodiment, render targets, such as frame buffers or texture maps may be stored across memory units **3624A-3624N**, allowing partition units **3620A-3620N** to write portions of each render target in parallel to efficiently use available bandwidth of parallel processor memory **3622**. In at least one embodiment, a local instance of parallel processor memory **3622** may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

[0357] In at least one embodiment, any one of clusters **3614A-3614N** of processing array **3612** can process data that will be written to any of memory units **3624A-3624N** within parallel processor memory **3622**. In at least one embodiment, memory crossbar **3616** can be configured to transfer an output of each cluster **3614A-3614N** to any partition unit **3620A-3620N** or to another cluster **3614A-3614N**, which can perform additional processing operations on an output. In at least one embodiment, each cluster **3614A-3614N** can communicate with memory interface **3618** through memory crossbar **3616** to read from or write to various external memory devices. In at least one embodiment, memory crossbar **3616** has a connection to memory interface **3618** to communicate with I/O unit **3604**, as well as a connection to a local instance of parallel processor memory **3622**, enabling processing units within different clusters **3614A-3614N** to communicate with system memory or other memory that is not local to parallel processing unit **3602**. In at least one embodiment, memory crossbar **3616** can use virtual channels to separate traffic streams between clusters **3614A-3614N** and partition units **3620A-3620N**.

[0358] In at least one embodiment, multiple instances of parallel processing unit **3602** can be provided on a single add-in card, or multiple add-in cards can be interconnected. In at least one embodiment, different instances of parallel processing unit **3602** can be configured to interoperate even if different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. In at least one embodiment, some instances of parallel processing unit **3602** can include higher precision floating point units relative to other instances. In at least one embodiment, systems incorporating one or more instances of parallel processing unit **3602** or parallel processor **3600** can be implemented in a variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems.

[0359] FIG. **36B** illustrates a processing cluster **3694**, in accordance with at least one embodiment. In at least one embodiment, processing cluster **3694** is included within a parallel processing unit. In at least one embodiment, processing cluster **3694** is one of processing clusters **3614A-3614N** of FIG. **36**. In at least one embodiment, processing cluster **3694** can be configured to execute many threads in parallel, where the term “thread” refers to an instance of a particular program executing on a particular set of input data. In at least one embodiment, single instruction, multiple data (“SIMD”) instruction issue techniques are used to support parallel execution of a large number of threads without providing multiple independent instruction units. In at least one embodiment, single instruction, multiple thread (“SIMT”) techniques are used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of processing engines within each processing cluster **3694**.

[0360] In at least one embodiment, operation of processing cluster **3694** can be controlled via a pipeline manager **3632** that distributes processing tasks to SIMT parallel processors. In at least one embodiment, pipeline manager **3632** receives instructions from scheduler **3610** of FIG. **36** and manages execution of those instructions via a graphics multiprocessor **3634** and/or a texture unit **3636**. In at least one embodiment, graphics multiprocessor **3634** is an exemplary instance of a SIMT parallel processor. However, in at least one embodiment, various types of SIMT parallel processors of differing architectures may be included within processing cluster **3694**. In at least one embodiment, one or more instances of graphics multiprocessor **3634** can be included within processing cluster **3694**. In at least one embodiment, graphics multiprocessor **3634** can process data and a data crossbar **3640** can be used to distribute processed data to one of multiple possible destinations, including other shader units. In at least one embodiment, pipeline manager **3632** can facilitate distribution of processed data by specifying destinations for processed data to be distributed via data crossbar **3640**.

[0361] In at least one embodiment, each graphics multiprocessor **3634** within processing cluster **3694** can include an identical set of functional execution logic (e.g., arithmetic logic units, load/store units (“LSUs”), etc.). In at least one embodiment, functional execution logic can be configured in a pipelined manner in which new instructions can be issued before previous instructions are complete. In at least one embodiment, functional execution logic supports a variety of operations including integer and floating point arithmetic, comparison operations, Boolean operations, bit-shifting, and computation of various algebraic functions. In at least one embodiment, same functional-unit hardware can be leveraged to perform different operations and any combination of functional units may be present.

[0362] In at least one embodiment, instructions transmitted to processing cluster **3694** constitute a thread. In at least one embodiment, a set of threads executing across a set of parallel processing engines is a thread group. In at least one embodiment, a thread group executes a program on different input data. In at least one embodiment, each thread within a thread group can be assigned to a different processing engine within graphics multiprocessor **3634**. In at least one embodiment, a thread group may include fewer threads than a number of processing engines within graphics multiprocessor **3634**. In at least one embodiment, when a thread group includes fewer threads than a number of processing engines, one or more of processing engines may be idle during cycles

in which that thread group is being processed. In at least one embodiment, a thread group may also include more threads than a number of processing engines within graphics multiprocessor 3634. In at least one embodiment, when a thread group includes more threads than a number of processing engines within graphics multiprocessor 3634, processing can be performed over consecutive clock cycles. In at least one embodiment, multiple thread groups can be executed concurrently on graphics multiprocessor 3634.

[0363] In at least one embodiment, graphics multiprocessor 3634 includes an internal cache memory to perform load and store operations. In at least one embodiment, graphics multiprocessor 3634 can forego an internal cache and use a cache memory (e.g., L1 cache 3648) within processing cluster 3694. In at least one embodiment, each graphics multiprocessor 3634 also has access to Level 2 (“L2”) caches within partition units (e.g., partition units 3620A-3620N of FIG. 36A) that are shared among all processing clusters 3694 and may be used to transfer data between threads. In at least one embodiment, graphics multiprocessor 3634 may also access off-chip global memory, which can include one or more of local parallel processor memory and/or system memory. In at least one embodiment, any memory external to parallel processing unit 3602 may be used as global memory. In at least one embodiment, processing cluster 3694 includes multiple instances of graphics multiprocessor 3634 that can share common instructions and data, which may be stored in L1 cache 3648.

[0364] In at least one embodiment, each processing cluster 3694 may include an MMU 3645 that is configured to map virtual addresses into physical addresses. In at least one embodiment, one or more instances of MMU 3645 may reside within memory interface 3618 of FIG. 36. In at least one embodiment, MMU 3645 includes a set of page table entries (“PTEs”) used to map a virtual address to a physical address of a tile and optionally a cache line index. In at least one embodiment, MMU 3645 may include address translation lookaside buffers (“TLBs”) or caches that may reside within graphics multiprocessor 3634 or L1 cache 3648 or processing cluster 3694. In at least one embodiment, a physical address is processed to distribute surface data access locality to allow efficient request interleaving among partition units. In at least one embodiment, a cache line index may be used to determine whether a request for a cache line is a hit or miss.

[0365] In at least one embodiment, processing cluster 3694 may be configured such that each graphics multiprocessor 3634 is coupled to a texture unit 3636 for performing texture mapping operations, e.g., determining texture sample positions, reading texture data, and filtering texture data. In at least one embodiment, texture data is read from an internal texture L1 cache (not shown) or from an L1 cache within graphics multiprocessor 3634 and is fetched from an L2 cache, local parallel processor memory, or system memory, as needed. In at least one embodiment, each graphics multiprocessor 3634 outputs a processed task to data crossbar 3640 to provide a processed task to another processing cluster 3694 for further processing or to store a processed task in an L2 cache, a local parallel processor memory, or a system memory via memory crossbar 3616. In at least one embodiment, a pre-raster operations unit (“preROP”) 3642 is configured to receive data from graphics multiprocessor 3634, direct data to ROP units, which may be located with partition units as described herein (e.g., partition units 3620A-3620N of FIG. 36). In at least one embodi-

ment, PreROP 3642 can perform optimizations for color blending, organize pixel color data, and perform address translations.

[0366] FIG. 36C illustrates a graphics multiprocessor 3696, in accordance with at least one embodiment. In at least one embodiment, graphics multiprocessor 3696 is graphics multiprocessor 3634 of FIG. 36B. In at least one embodiment, graphics multiprocessor 3696 couples with pipeline manager 3632 of processing cluster 3694. In at least one embodiment, graphics multiprocessor 3696 has an execution pipeline including but not limited to an instruction cache 3652, an instruction unit 3654, an address mapping unit 3656, a register file 3658, one or more GPGPU cores 3662, and one or more LSUs 3666. GPGPU cores 3662 and LSUs 3666 are coupled with cache memory 3672 and shared memory 3670 via a memory and cache interconnect 3668.

[0367] In at least one embodiment, instruction cache 3652 receives a stream of instructions to execute from pipeline manager 3632. In at least one embodiment, instructions are cached in instruction cache 3652 and dispatched for execution by instruction unit 3654. In at least one embodiment, instruction unit 3654 can dispatch instructions as thread groups (e.g., warps), with each thread of a thread group assigned to a different execution unit within GPGPU core 3662. In at least one embodiment, an instruction can access any of a local, shared, or global address space by specifying an address within a unified address space. In at least one embodiment, address mapping unit 3656 can be used to translate addresses in a unified address space into a distinct memory address that can be accessed by LSUs 3666.

[0368] In at least one embodiment, register file 3658 provides a set of registers for functional units of graphics multiprocessor 3696. In at least one embodiment, register file 3658 provides temporary storage for operands connected to data paths of functional units (e.g., GPGPU cores 3662, LSUs 3666) of graphics multiprocessor 3696. In at least one embodiment, register file 3658 is divided between each of functional units such that each functional unit is allocated a dedicated portion of register file 3658. In at least one embodiment, register file 3658 is divided between different thread groups being executed by graphics multiprocessor 3696.

[0369] In at least one embodiment, GPGPU cores 3662 can each include FPUs and/or integer ALUs that are used to execute instructions of graphics multiprocessor 3696. GPGPU cores 3662 can be similar in architecture or can differ in architecture. In at least one embodiment, a first portion of GPGPU cores 3662 include a single precision FPU and an integer ALU while a second portion of GPGPU cores 3662 include a double precision FPU. In at least one embodiment, FPUs can implement IEEE 754-2008 standard for floating point arithmetic or enable variable precision floating point arithmetic. In at least one embodiment, graphics multiprocessor 3696 can additionally include one or more fixed function or special function units to perform specific functions such as copy rectangle or pixel blending operations. In at least one embodiment one or more of GPGPU cores 3662 can also include fixed or special function logic.

[0370] In at least one embodiment, GPGPU cores 3662 include SIMD logic capable of performing a single instruction on multiple sets of data. In at least one embodiment GPGPU cores 3662 can physically execute SIMD4, SIMD8, and SIMD16 instructions and logically execute SIMD1, SIMD2, and SIMD32 instructions. In at least one embodiment, SIMD instructions for GPGPU cores 3662 can be generated at compile time by a shader compiler or automati-

cally generated when executing programs written and compiled for single program multiple data (“SPMD”) or SIMT architectures. In at least one embodiment, multiple threads of a program configured for an SIMT execution model can be executed via a single SIMD instruction. In at least one embodiment, eight SIMT threads that perform the same or similar operations can be executed in parallel via a single SIMD8 logic unit.

[0371] In at least one embodiment, memory and cache interconnect 3668 is an interconnect network that connects each functional unit of graphics multiprocessor 3696 to register file 3658 and to shared memory 3670. In at least one embodiment, memory and cache interconnect 3668 is a crossbar interconnect that allows LSU 3666 to implement load and store operations between shared memory 3670 and register file 3658. In at least one embodiment, register file 3658 can operate at a same frequency as GPGPU cores 3662, thus data transfer between GPGPU cores 3662 and register file 3658 is very low latency. In at least one embodiment, shared memory 3670 can be used to enable communication between threads that execute on functional units within graphics multiprocessor 3696. In at least one embodiment, cache memory 3672 can be used as a data cache in at least one embodiment, to cache texture data communicated between functional units and texture unit 3636. In at least one embodiment, shared memory 3670 can also be used as a program managed cache. In at least one embodiment, threads executing on GPGPU cores 3662 can programmatically store data within shared memory in addition to automatically cached data that is stored within cache memory 3672.

[0372] In at least one embodiment, a parallel processor or GPGPU as described herein is communicatively coupled to host/processor cores to accelerate graphics operations, machine-learning operations, pattern analysis operations, and various general purpose GPU (GPGPU) functions. In at least one embodiment, a GPU may be communicatively coupled to host processor/cores over a bus or other interconnect (e.g., a high speed interconnect such as PCIe or NVLink). In at least one embodiment, a GPU may be integrated on a same package or chip as cores and communicatively coupled to cores over a processor bus/interconnect that is internal to a package or a chip. In at least one embodiment, regardless of a manner in which a GPU is connected, processor cores may allocate work to a GPU in a form of sequences of commands/instructions contained in a WD. In at least one embodiment, a GPU then uses dedicated circuitry/logic for efficiently processing these commands/instructions.

General Computing

[0373] The following figures set forth, without limitation, exemplary software constructs within general computing that can be used to implement at least one embodiment.

[0374] FIG. 37 illustrates a software stack of a programming platform, in accordance with at least one embodiment. In at least one embodiment, a programming platform is a platform for leveraging hardware on a computing system to accelerate computational tasks. A programming platform may be accessible to software developers through libraries, compiler directives, and/or extensions to programming languages, in at least one embodiment. In at least one embodiment, a programming platform may be, but is not limited to, CUDA, Radeon Open Compute Platform (“ROCm”), OpenCL (OpenCL™ is developed by Khronos group), SYCL, or Intel One API.

[0375] In at least one embodiment, a software stack 3700 of a programming platform provides an execution environment for an application 3701. In at least one embodiment, application 3701 may include any computer software capable of being launched on software stack 3700. In at least one embodiment, application 3701 may include, but is not limited to, an artificial intelligence (“AI”)/machine learning (“ML”) application, a high performance computing (“HPC”) application, a virtual desktop infrastructure (“VDI”), or a datacenter workload.

[0376] In at least one embodiment, application 3701 and software stack 3700 run on hardware 3707. Hardware 3707 may include one or more GPUs, CPUs, FPGAs, AI engines, and/or other types of compute devices that support a programming platform, in at least one embodiment. In at least one embodiment, such as with CUDA, software stack 3700 may be vendor specific and compatible with only devices from particular vendor(s). In at least one embodiment, such as in with OpenCL, software stack 3700 may be used with devices from different vendors. In at least one embodiment, hardware 3707 includes a host connected to one more devices that can be accessed to perform computational tasks via application programming interface (“API”) calls. A device within hardware 3707 may include, but is not limited to, a GPU, FPGA, AI engine, or other compute device (but may also include a CPU) and its memory, as opposed to a host within hardware 3707 that may include, but is not limited to, a CPU (but may also include a compute device) and its memory, in at least one embodiment.

[0377] In at least one embodiment, software stack 3700 of a programming platform includes, without limitation, a number of libraries 3703, a runtime 3705, and a device kernel driver 3706. Each of libraries 3703 may include data and programming code that can be used by computer programs and leveraged during software development, in at least one embodiment. In at least one embodiment, libraries 3703 may include, but are not limited to, pre-written code and subroutines, classes, values, type specifications, configuration data, documentation, help data, and/or message templates. In at least one embodiment, libraries 3703 include functions that are optimized for execution on one or more types of devices. In at least one embodiment, libraries 3703 may include, but are not limited to, functions for performing mathematical, deep learning, and/or other types of operations on devices. In at least one embodiment, libraries 3803 are associated with corresponding APIs 3802, which may include one or more APIs, that expose functions implemented in libraries 3803.

[0378] In at least one embodiment, application 3701 is written as source code that is compiled into executable code, as discussed in greater detail below in conjunction with FIG. 42. Executable code of application 3701 may run, at least in part, on an execution environment provided by software stack 3700, in at least one embodiment. In at least one embodiment, during execution of application 3701, code may be reached that needs to run on a device, as opposed to a host. In such a case, runtime 3705 may be called to load and launch requisite code on a device, in at least one embodiment. In at least one embodiment, runtime 3705 may include any technically feasible runtime system that is able to support execution of application S01.

[0379] In at least one embodiment, runtime 3705 is implemented as one or more runtime libraries associated with corresponding APIs, which are shown as API(s) 3704. One or more of such runtime libraries may include, without limitation, functions for memory management, execution control, device management, error handling, and/or synchro-

nization, among other things, in at least one embodiment. In at least one embodiment, memory management functions may include, but are not limited to, functions to allocate, deallocate, and copy device memory, as well as transfer data between host memory and device memory. In at least one embodiment, execution control functions may include, but are not limited to, functions to launch a function (sometimes referred to as a “kernel” when a function is a global function callable from a host) on a device and set attribute values in a buffer maintained by a runtime library for a given function to be executed on a device.

[0380] Runtime libraries and corresponding API(s) **3704** may be implemented in any technically feasible manner, in at least one embodiment. In at least one embodiment, one (or any number of) API may expose a low-level set of functions for fine-grained control of a device, while another (or any number of) API may expose a higher-level set of such functions. In at least one embodiment, a high-level runtime API may be built on top of a low-level API. In at least one embodiment, one or more of runtime APIs may be language-specific APIs that are layered on top of a language-independent runtime API.

[0381] In at least one embodiment, device kernel driver **3706** is configured to facilitate communication with an underlying device. In at least one embodiment, device kernel driver **3706** may provide low-level functionalities upon which APIs, such as API(s) **3704**, and/or other software relies. In at least one embodiment, device kernel driver **3706** may be configured to compile intermediate representation (“IR”) code into binary code at runtime. For CUDA, device kernel driver **3706** may compile Parallel Thread Execution (“PTX”) IR code that is not hardware specific into binary code for a specific target device at runtime (with caching of compiled binary code), which is also sometimes referred to as “finalizing” code, in at least one embodiment. Doing so may permit finalized code to run on a target device, which may not have existed when source code was originally compiled into PTX code, in at least one embodiment. Alternatively, in at least one embodiment, device source code may be compiled into binary code offline, without requiring device kernel driver **3706** to compile IR code at runtime.

[0382] FIG. **38** illustrates a CUDA implementation of software stack **3700** of FIG. **37**, in accordance with at least one embodiment. In at least one embodiment, a CUDA software stack **3800**, on which an application **3801** may be launched, includes CUDA libraries **3803**, a CUDA runtime **3805**, a CUDA driver **3807**, and a device kernel driver **3808**. In at least one embodiment, CUDA software stack **3800** executes on hardware **3809**, which may include a GPU that supports CUDA and is developed by NVIDIA Corporation of Santa Clara, Calif.

[0383] In at least one embodiment, application **3801**, CUDA runtime **3805**, and device kernel driver **3808** may perform similar functionalities as application **3701**, runtime **3705**, and device kernel driver **3706**, respectively, which are described above in conjunction with FIG. **37**. In at least one embodiment, CUDA driver **3807** includes a library (libcudart) that implements a CUDA driver API **3806**. Similar to a CUDA runtime API **3804** implemented by a CUDA runtime library (cudart), CUDA driver API **3806** may, without limitation, expose functions for memory management, execution control, device management, error handling, synchronization, and/or graphics interoperability, among other things, in at least one embodiment. In at least one embodiment, CUDA driver API **3806** differs from CUDA runtime API **3804** in that CUDA runtime API **3804** simplifies device code man-

agement by providing implicit initialization, context (analogous to a process) management, and module (analogous to dynamically loaded libraries) management. In contrast to high-level CUDA runtime API **3804**, CUDA driver API **3806** is a low-level API providing more fine-grained control of a device, particularly with respect to contexts and module loading, in at least one embodiment. In at least one embodiment, CUDA driver API **3806** may expose functions for context management that are not exposed by CUDA runtime API **3804**. In at least one embodiment, CUDA driver API **3806** is also language-independent and supports, e.g., OpenCL in addition to CUDA runtime API **3804**. Further, in at least one embodiment, development libraries, including CUDA runtime **3805**, may be considered as separate from driver components, including user-mode CUDA driver **3807** and kernel-mode device driver **3808** (also sometimes referred to as a “display” driver).

[0384] In at least one embodiment, CUDA libraries **3803** may include, but are not limited to, mathematical libraries, deep learning libraries, parallel algorithm libraries, and/or signal/image/video processing libraries, which parallel computing applications such as application **3801** may utilize. In at least one embodiment, CUDA libraries **3803** may include mathematical libraries such as a cuBLAS library that is an implementation of Basic Linear Algebra Subprograms (“BLAS”) for performing linear algebra operations, a cuFFT library for computing fast Fourier transforms (“FFTs”), and a cuRAND library for generating random numbers, among others. In at least one embodiment, CUDA libraries **3803** may include deep learning libraries such as a cuDNN library of primitives for deep neural networks and a TensorRT platform for high-performance deep learning inference, among others.

[0385] FIG. **39** illustrates a ROCm implementation of software stack **3700** of FIG. **37**, in accordance with at least one embodiment. In at least one embodiment, a ROCm software stack **3900**, on which an application **3901** may be launched, includes a language runtime **3903**, a system runtime **3905**, a thunk **3907**, a ROCm kernel driver **3908**, and a device kernel driver **3909**. In at least one embodiment, ROCm software stack **3900** executes on hardware **3910**, which may include a GPU that supports ROCm and is developed by AMD Corporation of Santa Clara, Calif.

[0386] In at least one embodiment, application **3901** may perform similar functionalities as application **3701** discussed above in conjunction with FIG. **37**. In addition, language runtime **3903** and system runtime **3905** may perform similar functionalities as runtime **3705** discussed above in conjunction with FIG. **37**, in at least one embodiment. In at least one embodiment, language runtime **3903** and system runtime **3905** differ in that system runtime **3905** is a language-independent runtime that implements a ROCr system runtime API **3904** and makes use of a Heterogeneous System Architecture (“HAS”) Runtime API. HAS runtime API is a thin, user-mode API that exposes interfaces to access and interact with an AMD GPU, including functions for memory management, execution control via architected dispatch of kernels, error handling, system and agent information, and runtime initialization and shutdown, among other things, in at least one embodiment. In contrast to system runtime **3905**, language runtime **3903** is an implementation of a language-specific runtime API **3902** layered on top of ROCr system runtime API **3904**, in at least one embodiment. In at least one embodiment, language runtime API may include, but is not limited to, a Heterogeneous compute Interface for Portability (“HIP”) language runtime API, a Heterogeneous Compute Compiler (“HCC”) lan-

guage runtime API, or an OpenCL API, among others. HIP language in particular is an extension of C++ programming language with functionally similar versions of CUDA mechanisms, and, in at least one embodiment, a HIP language runtime API includes functions that are similar to those of CUDA runtime API **3804** discussed above in conjunction with FIG. **38**, such as functions for memory management, execution control, device management, error handling, and synchronization, among other things.

[0387] In at least one embodiment, thunk (ROCT) **3907** is an interface that can be used to interact with underlying ROCm driver **3908**. In at least one embodiment, ROCm driver **3908** is a ROCK driver, which is a combination of an AMDGPU driver and a HAS kernel driver (amdkfd). In at least one embodiment, AMDGPU driver is a device kernel driver for GPUs developed by AMD that performs similar functionalities as device kernel driver **3706** discussed above in conjunction with FIG. **37**. In at least one embodiment, HAS kernel driver is a driver permitting different types of processors to share system resources more effectively via hardware features.

[0388] In at least one embodiment, various libraries (not shown) may be included in ROCm software stack **3900** above language runtime **3903** and provide functionality similarity to CUDA libraries **3803**, discussed above in conjunction with FIG. **38**. In at least one embodiment, various libraries may include, but are not limited to, mathematical, deep learning, and/or other libraries such as a hipBLAS library that implements functions similar to those of CUDA cuBLAS, a rocFFT library for computing FFTs that is similar to CUDA cuFFT, among others.

[0389] FIG. **40** illustrates an OpenCL implementation of software stack **3700** of FIG. **37**, in accordance with at least one embodiment. In at least one embodiment, an OpenCL software stack **4000**, on which an application **4001** may be launched, includes an OpenCL framework **4005**, an OpenCL runtime **4006**, and a driver **4007**. In at least one embodiment, OpenCL software stack **4000** executes on hardware **3809** that is not vendor-specific. As OpenCL is supported by devices developed by different vendors, specific OpenCL drivers may be required to interoperate with hardware from such vendors, in at least one embodiment.

[0390] In at least one embodiment, application **4001**, OpenCL runtime **4006**, device kernel driver **4007**, and hardware **4008** may perform similar functionalities as application **3701**, runtime **3705**, device kernel driver **3706**, and hardware **3707**, respectively, that are discussed above in conjunction with FIG. **37**. In at least one embodiment, application **4001** further includes an OpenCL kernel **4002** with code that is to be executed on a device.

[0391] In at least one embodiment, OpenCL defines a “platform” that allows a host to control devices connected to a host. In at least one embodiment, an OpenCL framework provides a platform layer API and a runtime API, shown as platform API **4003** and runtime API **4005**. In at least one embodiment, runtime API **4005** uses contexts to manage execution of kernels on devices. In at least one embodiment, each identified device may be associated with a respective context, which runtime API **4005** may use to manage command queues, program objects, and kernel objects, share memory objects, among other things, for that device. In at least one embodiment, platform API **4003** exposes functions that permit device contexts to be used to select and initialize devices, submit work to devices via command queues, and enable data transfer to and from devices, among other things. In addition, OpenCL framework provides various built-in functions (not shown), including math functions,

relational functions, and image processing functions, among others, in at least one embodiment.

[0392] In at least one embodiment, a compiler **4004** is also included in OpenCL frame-work **4005**. Source code may be compiled offline prior to executing an application or online during execution of an application, in at least one embodiment. In contrast to CUDA and ROCm, OpenCL applications in at least one embodiment may be compiled online by compiler **4004**, which is included to be representative of any number of compilers that may be used to compile source code and/or IR code, such as Standard Portable Intermediate Representation (“SPIR-V”) code, into binary code. Alternatively, in at least one embodiment, OpenCL applications may be compiled offline, prior to execution of such applications.

[0393] FIG. **41** illustrates software that is supported by a programming platform, in accordance with at least one embodiment. In at least one embodiment, a programming platform **4104** is configured to support various programming models **4103**, middlewares and/or libraries **4102**, and frameworks **4101** that an application **4100** may rely upon. In at least one embodiment, application **4100** may be an AI/ML application implemented using, in at least one embodiment, a deep learning framework such as MXNet, PyTorch, or TensorFlow, which may rely on libraries such as cuDNN, NVIDIA Collective Communications Library (“NCCL”), and/or NVIDIA Developer Data Loading Library (“DALI”) CUDA libraries to provide accelerated computing on underlying hardware.

[0394] In at least one embodiment, programming platform **4104** may be one of a CUDA, ROCm, or OpenCL platform described above in conjunction with FIG. **33**, FIG. **34**, and FIG. **40**, respectively. In at least one embodiment, programming platform **4104** supports multiple programming models **4103**, which are abstractions of an underlying computing system permitting expressions of algorithms and data structures. Programming models **4103** may expose features of underlying hardware in order to improve performance, in at least one embodiment. In at least one embodiment, programming models **4103** may include, but are not limited to, CUDA, HIP, OpenCL, C++ Accelerated Massive Parallelism (“C++ AMP”), Open Multi-Processing (“OpenMP”), Open Accelerators (“OpenACC”), and/or Vulkan Compute.

[0395] In at least one embodiment, libraries and/or middlewares **4102** provide implementations of abstractions of programming models **4104**. In at least one embodiment, such libraries include data and programming code that may be used by computer programs and leveraged during software development. In at least one embodiment, such middlewares include software that provides services to applications beyond those available from programming platform **4104**. In at least one embodiment, libraries and/or middlewares **4102** may include, but are not limited to, cuBLAS, cuFFT, cuRAND, and other CUDA libraries, or rocBLAS, rocFFT, rocRAND, and other ROCm libraries. In addition, in at least one embodiment, libraries and/or middlewares **4102** may include NCCL and ROCm Communication Collectives Library (“RCCL”) libraries providing communication routines for GPUs, a MIOpen library for deep learning acceleration, and/or an Eigen library for linear algebra, matrix and vector operations, geometrical transformations, numerical solvers, and related algorithms.

[0396] In at least one embodiment, application frameworks **4101** depend on libraries and/or middlewares **4102**. In at least one embodiment, each of application frameworks **4101** is a software framework used to implement a standard structure of application software. An AI/ML application may

be implemented using a framework such as Caffe, Caffe2, TensorFlow, Keras, PyTorch, or MxNet deep learning frameworks, in at least one embodiment.

[0397] FIG. 42 illustrates compiling code to execute on one of programming platforms of FIGS. 37-40, in accordance with at least one embodiment. In at least one embodiment, a compiler 4201 receives source code 4200 that includes both host code as well as device code. In at least one embodiment, compiler 4201 is configured to convert source code 4200 into host executable code 4202 for execution on a host and device executable code 4203 for execution on a device. In at least one embodiment, source code 4200 may either be compiled offline prior to execution of an application, or online during execution of an application.

[0398] In at least one embodiment, source code 4200 may include code in any programming language supported by compiler 4201, such as C++, C, Fortran, etc. In at least one embodiment, source code 4200 may be included in a single-source file having a mixture of host code and device code, with locations of device code being indicated therein. In at least one embodiment, a single-source file may be a .cu file that includes CUDA code or a .hip.cpp file that includes HIP code. Alternatively, in at least one embodiment, source code 4200 may include multiple source code files, rather than a single-source file, into which host code and device code are separated.

[0399] In at least one embodiment, compiler 4201 is configured to compile source code 4200 into host executable code 4202 for execution on a host and device executable code 4203 for execution on a device. In at least one embodiment, compiler 4201 performs operations including parsing source code 4200 into an abstract system tree (AST), performing optimizations, and generating executable code. In at least one embodiment in which source code 4200 includes a single-source file, compiler 4201 may separate device code from host code in such a single-source file, compile device code and host code into device executable code 4203 and host executable code 4202, respectively, and link device executable code 4203 and host executable code 4202 together in a single file, as discussed in greater detail below with respect to FIG. 26.

[0400] In at least one embodiment, host executable code 4202 and device executable code 4203 may be in any suitable format, such as binary code and/or IR code. In a case of CUDA, host executable code 4202 may include native object code and device executable code 4203 may include code in PTX intermediate representation, in at least one embodiment. In a case of ROCm, both host executable code 4202 and device executable code 4203 may include target binary code, in at least one embodiment.

[0401] Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit disclosure to specific form or forms disclosed, but on contrary, intention is to cover all modifications, alternative constructions, and equivalents falling within spirit and scope of disclosure, as defined in appended claims.

[0402] Use of terms “a” and “an” and “the” and similar referents in context of describing disclosed embodiments (especially in context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-

ended terms (meaning “including, but not limited to,”) unless otherwise noted. term “connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within range, unless otherwise indicated herein and each separate value is incorporated into specification as if it were individually recited herein. In at least one embodiment, use of term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, term “subset” of a corresponding set does not necessarily denote a proper subset of corresponding set, but subset and corresponding set may be equal.

[0403] Conjunctive language, such as phrases of form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of set of A and B and C. In at least one embodiment of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). In at least one embodiment, a number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, phrase “based on” means “based at least in part on” and not “based solely on.”

[0404] Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium. In at least one embodiment, in form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause computer system to perform operations described herein. A set of non-transitory computer-readable storage media, in at least one embodiment, comprises multiple

non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of code while multiple non-transitory computer-readable storage media collectively store all of code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors—in at least one embodiment, a non-transitory computer-readable storage medium store instructions and a main central processing unit (“CPU”) executes some of instructions while a graphics processing unit (“GPU”) executes other instructions. In at least one embodiment, different components of a computer system have separate processors and different processors execute different subsets of instructions.

[0405] Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

[0406] Use of any and all of the at least one embodiments, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of disclosure and does not pose a limitation on scope of disclosure unless otherwise claimed. No language in specification should be construed as indicating any non-claimed element as essential to practice of disclosure.

[0407] All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

[0408] In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may be not intended as synonyms for each other. Rather, in ones of at least one embodiments, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

[0409] Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system’s registers and/or memories into other data similarly represented as physical quantities within computing system’s memories, registers or other such information storage, transmission or display devices.

[0410] In a similar manner, term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting ones of the at least one embodiments, “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more proces-

sors. As used herein, “software” processes may include, in at least one embodiment, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. Terms “system” and “method” are used herein interchangeably insofar as system may embody one or more methods and methods may be considered a system.

[0411] In present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In some implementations, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In another implementation, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. References may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various ones of the at least one embodiments, process of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or interprocess communication mechanism.

[0412] Although discussion above sets forth ones of the at least one embodiments having implementations of described techniques, other architectures may be used to implement described functionality, and are intended to be within scope of this disclosure. Furthermore, although specific distributions of responsibilities are defined above for purposes of discussion, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

[0413] Furthermore, although subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

What is claimed is:

1. A datacenter cooling system, comprising:

a cold plate comprising microchannels and a heat pipe, the microchannels to support a first fluid in an active mode of operation of the cold plate, the first fluid to flow through the microchannels by action of at least one flow controller, and the heat pipe to support a second fluid in a passive mode of operation of the cold plate, the second fluid to circulate in the heat pipe by differential temperature and capillary actions between an evaporator section and a condenser section of the heat pipe.

2. The datacenter cooling system of claim 1, further comprising:

at least one processor to determine a temperature associated with a computing device or the cold plate and to cause the at least one flow controller to pass the first fluid through the microchannels to enable the active mode for the cold plate.

3. The datacenter cooling system of claim 1, further comprising:

at least one processor to determine a temperature or power associated with a computing device and to cause one or more of the active mode or a combination of the active mode and the passive mode for the cold plate.

4. The datacenter cooling system of claim 1, further comprising:

the cold plate associated with a computing device and having ports to support entry and egress of the first fluid through the microchannels, wherein the first fluid is maintained distinctly from the second fluid of the heat pipe.

5. The datacenter cooling system of claim 1, further comprising:

at least one processor to receive sensor inputs from sensors associated with a computing device or the cold plate, the at least one processor to determine a first cooling requirement associated with the passive mode and a second cooling requirement associated with a combination of the active mode and the passive mode, based in part on the sensor inputs.

6. The datacenter cooling system of claim 5, further comprising:

one or more neural networks to receive the sensor inputs and to infer the first cooling requirement and the second cooling requirement.

7. The datacenter cooling system of claim 1, further comprising:

at least one processor to receive sensor inputs from temperature sensors and from power sensors associated with the computing device, the at least one processor to determine a first cooling requirement and a first power requirement associated with the passive mode, and a second cooling requirement and a second power requirement associated with a combination of the active mode and the passive mode, based in part on the sensor inputs.

8. The datacenter cooling system of claim 1, wherein the second fluid comprises a phase-changeable fluid operative in a vapor phase in the evaporator section and a liquid or partly liquid phase in the condenser section.

9. The datacenter cooling system of claim 1, wherein the heat pipe comprises a wicking material to enable passive movement of the second fluid from the condenser section to the evaporator section.

10. The datacenter cooling system of claim 1, further comprising:

a liquid-to-air heat exchanger associated with the heat pipe to enable dissipation of heat or cooling of the second fluid at the condenser section, the condenser section located remotely from the cold plate.

11. A processor comprising one or more circuits and associated with a cold plate comprising microchannels and a heat pipe, the processor to cause an active mode of operation for the cold plate so that microchannels support a first fluid through the microchannels by action of at least one flow controller, the active mode to supplement a passive mode of operation of the cold plate, the passive mode to enable the heat pipe to circulate a second fluid therein by differential temperature and capillary actions between an evaporator section and a condenser section of the heat pipe.

12. The processor of claim 11, further comprising:

an output of the one or more circuits to provide signals for the at least one flow controller to enable flow of the first fluid through the microchannels.

13. The processor of claim 11, further comprising:

an input of the one or more circuits to receive sensor inputs from sensors associated with at least one computing device or the second fluid, the processor to determine a cooling requirement based in part on the sensor inputs.

14. The processor of claim 13, further comprising:

one or more neural networks to receive the sensor inputs and to infer the cooling requirement.

15. The processor of claim 11, further comprising:

one or more neural networks to infer a power requirement threshold for at least one computing device and the one or more circuits to cause the active mode for the cold plate.

16. A processor comprising one or more circuits, the one or more circuits to train one or more neural networks to infer a cooling requirement from sensor inputs of sensors associated with a computing device or a cold plate, the processor to cause an active mode of operation for the cold plate so that microchannels therein support flow of a first fluid by an action of at least one flow controller, the active mode to supplement a passive mode of operation of the cold plate, the passive mode to enable circulation of a second fluid through a heat pipe of the cold plate by differential temperature and capillary actions between an evaporator section and a condenser section of the heat pipe.

17. The processor of claim 16, further comprising:

an output of the one or more circuits to provide signals for the at least one flow controller to enable flow of the first fluid through the microchannels.

18. The processor of claim 16, further comprising:

the one or more neural networks to receive the sensor inputs from the sensors and to be trained to infer the cooling requirement as part of an analysis of prior sensor inputs and prior cooling requirements.

19. The processor of claim 16, further comprising:

an output of the one or more circuits to provide signals to cause the at least one flow controller to activate or deactivate the active mode.

20. The processor of claim 16, further comprising:

an input of the one or more circuits to receive the sensor inputs associated with a temperature from at least one computing device or the second fluid, the one or more neural networks trained to infer a change in coolant state has occurred based in part on the temperature and on prior temperatures.

21. A processor comprising one or more circuits, the one or more circuits to comprise one or more neural networks to infer a cooling requirement from sensor inputs of sensors associated with a computing device or a cold plate, the processor to cause an active mode of operation for the cold plate so that microchannels therein support flow of a first fluid by an action of at least one flow controller, the active mode to supplement a passive mode of operation of the cold plate, the passive mode to enable circulation of a second fluid through a heat pipe of the cold plate by differential temperature and capillary actions between an evaporator section and a condenser section of the heat pipe.

22. The processor of claim 21, further comprising:

an output of the one or more circuits to provide signals for the at least one flow controller to enable flow of the first fluid through the microchannels.

23. The processor of claim 21, further comprising:

the one or more neural networks to receive the sensor inputs from the sensors and to infer the cooling requirement as part of an analysis of prior sensor inputs and prior cooling requirements.

- 24.** The processor of claim **21**, further comprising:
an output of the one or more circuits to provide signals to cause the at least one flow controller to activate or deactivate the active mode.
- 25.** The processor of claim **21**, further comprising:
an input of the one or more circuits to receive the sensor inputs associated with a temperature from at least one computing device or the second fluid, the one or more neural networks to infer a change in coolant state has occurred based in part on the temperature and on prior temperatures.
- 26.** A method for datacenter cooling system, comprising:
providing a cold plate comprising microchannels and a heat pipe;
determining at least one cooling requirement for at least one computing device;
enabling the microchannels to support a first fluid in an active mode of operation of the cold plate, the first fluid forced through the microchannels by at least one flow controller; and
enabling the heat pipe to support a second fluid in a passive mode of operation of the cold plate, the second fluid forced to circulate in the heat pipe by differential temperature and capillary actions between an evaporator section and a condenser section of the heat pipe.
- 27.** The method of claim **26**, further comprising:
determining, using at least one processor, a temperature associated with a computing device or the cold plate; and

causing the at least one flow controller to pass the first fluid through the microchannels to enable the active mode for the cold plate.

- 28.** The method of claim **26**, further comprising:
determining, using at least one processor, a temperature or power associated with a computing device; and
causing one or more of the active mode or a combination of the active mode and the passive mode for the cold plate.

- 29.** The method of claim **26**, further comprising:
receiving, in at least one processor, sensor inputs from sensors associated with a computing device or the cold plate; and
determining, by the at least one processor, a first cooling requirement associated with the passive mode and a second cooling requirement associated with a combination of the active mode and the passive mode, based in part on the sensor inputs.

- 30.** The method of claim **26**, further comprising:
receiving, in at least one processor, sensor inputs from temperature sensors and from power sensors associated with the computing device; and
determining, using the at least one processor a first cooling requirement and a first power requirement associated with the passive mode, a second cooling requirement and a second power requirement associated with a combination of the active mode and the passive mode, based in part on the sensor inputs.

* * * * *